



COMPUTER SCIENCE

INDIANA UNIVERSITY

School of Informatics and Computing
Bloomington

Hidden variable models



Reminders/Comments

- Not requiring the assignment to be written-up in tex, but if writing it by hand, it needs to be legible
- Small comment: can equivalently use
 - diagonal matrix and standard matrix multiplication: $C v$
 - element-wise multiplication (Hadamard product) of vectors: $c \circ v$
- Naive Bayes question asks about adding a column of ones; I may have provided a chunk of code that is robust to this, but answer the question assuming I had not done so
- Questions about accuracy of the algorithms in the assignment



Accuracy of learned models

- The datasets from UCI are somewhat notoriously simple
 - “Very Simple Classification Rules Perform Well on Most Commonly Used Datasets”, Holte, 1993
- They have gotten more interesting, but still some issues
- For many machine learning algorithms, the differences are only evident on some datasets (not any dataset)
- Here the goal is to implement the algorithms and try to ensure their correctness
- Question 3: adding regularizers *can* outperform the base logistic regression; if it is not, try to see why and explain
 - look at the (final) weights as a debugging tool
 - print out the function values that are obtained along the way, ensure they steadily improving



Question 3 and regularizers

- Adding regularizers *can* outperform the base logistic regression
- If it is not, try to see why and explain
 - look at the (final) weights as a debugging tool
 - print out the function values that are obtained along the way, ensure they steadily improving
 - why should they be steadily improving?
 - think about your range of regularizers and what it *should* be; for example, how did your choice of l2 regularizer affect the solution in linear regression?



Neural networks

- Using neural networks *can* (and will if tuned well) outperform the base logistic regression
- If it is not, try to understand why
 - again, look at the (final) weights as a debugging tool
 - in this case, should your objective value be steadily decreasing? Is this true for batch gradient descent or stochastic gradient descent?
- For all your algorithms, consider comparing to python's library as a sanity check
 - if their learned models are significantly outperforming your learned models, then you might have a bug
 - if their models perform similarly poorly, then this might simply be a hard problem for that algorithm and/or tuning is difficult
 - if your model out-does python, feel proud and don't be too surprised; a capable implementer can often outperform packages



Student example for neural net

- “Strange” behavior in neural network
- I ran it with stepsize = .001 with the following iterations and accuracies:
 - 2: 50%
 - 3: 63%
 - 4: 68%
 - 5: 70% <--peak
 - 6: 60%
 - 7: 40%
 - 8: 48%
 - 9: 67% <--another peak
 - 10: 47%
 - 11+ ~50%
- I then ran it 5 times with stepsize = .00005 with 100 iterations and got the following accuracies: 80%, 45%, 63%, 73%, 73%.




Hidden variables

- Different from missing variables, in the sense that we *could* have observed the missing information
 - e.g., if the person had just filled in the box on the form
- Hidden variables are never observed; rather they are useful for model description
 - e.g., hidden, latent representation
 - e.g., hidden state that drives dynamics
- Hidden variables make specification of distribution simpler
 - $p(x \mid D) = \int p(x \mid D, h) p(h)$
 - $p(x \mid D, h)$ is often much simpler to specify



Intuitive example

- Underlying “state” influencing what we observe; partial observability makes what we observe difficult to interpret
- Image we can never see that a kitten is present; but it clearly helps to explain the data





Hidden variable models

- Probabilistic PCA and factor analysis
 - common in psychology
- Mixture models
- Hidden Markov Models
 - commonly used for NLP and modeling dynamical systems



Probabilistic PCA

- In PCA, we learned $p(\mathbf{x} | \mathbf{D}, \mathbf{h})$
 - What were the assumptions on $p(\mathbf{x} | \mathbf{D}, \mathbf{h})$?
- For Probabilistic PCA, we learn $p(\mathbf{x} | \mathbf{D})$
- Given some prior $p(\mathbf{h})$, we have

$$p(\mathbf{x} | \mathbf{D}) = \int_{\mathcal{H}} p(\mathbf{x} | \mathbf{D}, \mathbf{h}) p(\mathbf{h}) d\mathbf{h}$$



Modified goal

- The interpretation of the hidden factors as a new representation is still reasonable in this setting
- Now our goal is to obtain a distribution over x , only given the dictionary and not the representation
 - Why do we care about having distributions over x ? Why isn't $p(x | D, h)$ “good” enough?
 - What can we do with $p(x | D)$ that we could not do with $p(x | D, h)$, assuming we have learned D ?




Resulting differences


- Resulting solution for D is actually very similar, in the case of probabilistic PCA
 - In PCA, $D = U \Sigma$
 - In probabilistic PCA, $D = U (\Sigma - \sigma^2 I)$
- The model now is generative
 - can think of the previous one as discriminative, since need h to obtain the distribution over x
 - parallels $p(y | x)$ versus $p(x, y)$
- Solution approach is different
- Probabilistic PCA extends to more generally to other probabilistic models (e.g. factor analysis) that does not have such a similar solution

Generating data

- Sample \mathbf{h} from $p(\mathbf{h})$, then sample \mathbf{x} from $p(\mathbf{x} | \mathbf{D}, \mathbf{h})$
 - both of these distributions are Gaussian and so simple to sample



(a)



(b)

Figure : Factor Analysis: 1000 points generated from the model. **(a):** 1000 latent two-dimensional points \mathbf{h}^n sampled from $\mathcal{N}(\mathbf{h} | \mathbf{0}, \mathbf{I})$. These are transformed to a point on the three-dimensional plane by $\mathbf{x}_0^n = \mathbf{c} + \mathbf{F}\mathbf{h}^n$. The covariance of \mathbf{x}_0 is degenerate, with covariance matrix $\mathbf{F}\mathbf{F}^\top$. **(b):** For each point \mathbf{x}_0^n on the plane a random noise vector is drawn from $\mathcal{N}(\boldsymbol{\epsilon} | \mathbf{0}, \boldsymbol{\Psi})$ and added to the in-plane vector to form a sample \mathbf{x}^n , plotted in red. The distribution of points forms a ‘pancake’ in space. Points ‘underneath’ the plane are not shown.



Other hidden variable models

- Probabilistic PCA and factor analysis
 - common in psychology
- Mixture models
- Hidden Markov Models
 - commonly used for NLP and modeling dynamical systems



Gaussian mixture model


A D dimensional Gaussian distribution for a continuous variable \mathbf{x} is

$$p(\mathbf{x}|\mathbf{m}, \mathbf{S}) = \frac{1}{\sqrt{\det(2\pi\mathbf{S})}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mathbf{m})^\top \mathbf{S}^{-1} (\mathbf{x} - \mathbf{m}) \right\}$$

where \mathbf{m} is the mean and \mathbf{S} is the covariance matrix. A mixture of Gaussians is then


$$p(\mathbf{x}) = \sum_{i=1}^H p(\mathbf{x}|\mathbf{m}_i, \mathbf{S}_i) p(i)$$

where $p(i)$ is the mixture weight for component i .





Hidden Markov Model



The observation are x_1, x_2, x_3
Temporally related
Dynamics driven by hidden state



Closed-form solutions

- For some hidden variable models, have a closed form solution
 - probabilistic PCA
 - factor analysis
- Probabilistic PCA solution:

$$\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$$

$$\mathbf{D} = \mathbf{U}_k(\mathbf{\Sigma}_k^2 - \sigma^2\mathbf{I}_k)^{1/2}$$

$$\sigma^2 = \frac{1}{d-k} \sum_{i=k+1}^d \sigma_i^2$$



Expectation-maximization

- We can use an expectation-maximization approach instead to incrementally compute the solution (rather than a closed form)
- Similar to alternating descent approach taken for RFMs
 - For PCA, instead of computing a closed-form solution to D and H , we could have simply used gradient descent with our objective
- What is the advantage to using the incremental EM approach, when we already have a closed form?
 - other than as an educational example of EM



Whiteboard

- Closed form solution for probabilistic PCA
- Expectation-maximization for probabilistic PCA