



COMPUTER SCIENCE

INDIANA UNIVERSITY

School of Informatics and Computing  
Bloomington

# Performance measures



# Reminders/comments

- Assignment 4
  - For the first question, there \*may not\* be a solution
  - For the kernels question, if you are getting terrible performance, consider the choices for your kernel algorithm.



# Experimental set-up

- Cross-validation? External or Internal?
- Repeated subsampling?
- Careful statistical work done on executing empirical studies; pros and cons to each (and some are also just wrong)
- My personal preference:
  - Repeated subsampling to generate training/validation splits (can choose as many as desired, say  $m$ )
  - $k$ -fold cross-validation on training set to select parameters





# k-fold CV for training set $i$

Randomly and evenly split into 4 non-overlapping partitions


$D$

20 data points




Partition 1.

Data points: 1, 3, 5, 15, 16




Partition 2.

Data points: 6, 10, 11, 14, 17



Partition 3.

Data points: 4, 9, 12, 19, 20



Partition 4.



Data points: 2, 7, 8, 13, 17



# k-fold CV for training set i


- For parameter setting  $j$ , learn model on  $k-1$  folds and test on the hold-out fold (done  $k$  times); average  $k$  error estimates
- Select parameter setting  $j$  that gives lowest average error

20 data points



Partition 1.

Data points: 1, 3, 5, 15, 16




Partition 2.

Data points: 6, 10, 11, 14, 17



Partition 3.

Data points: 4, 9, 12, 19, 20



Partition 4.

Data points: 2, 7, 8, 13, 17



# Training set $i$

- Once “best” parameters obtained for the training set, learn on the entire training set  $i$  with those parameters, and validate on validation set  $i$
- This gives you an estimate of error for training/validation split  $i$
- Over all  $m$  splits, get  $m$  estimates of error (different from  $k$ )



# Experiments

- “What if I cannot find any difference between the algorithms?”
  - If you ran a fair experiment, with lots of repetitions (random splits of training and test) to get a large enough number of samples of the error, then that is a fine result
  - Remember that algorithms have many parameters that can strongly affect their performance
- “How do I select parameter ranges?”
  - the best is to provide a large enough range; this can be slow
- “Do I have to sweep all parameters in the CV?”
  - No, but remember that any choice of parameters affects your conclusion  
—> it is much less interesting to conclude that linear regression with regularization weight = 0.1 is outperformed by a NN with



# Why CV?

- CV a reasonable strategy to allow data to pick your parameters: measures on folds give idea of generalization
- Also simulates how someone (or you) might use the model in practice:
  - given a dataset set, you might do 10-fold CV to select across your meta-parameters
  - learn one model with those meta-parameters and now can use for future test data (which we do not have access to)





# What to avoid

- Data snooping: looking at test data and iteratively modifying your model
  - BAD: do CV with set of parameters on training data, then see if model did well on test; if did not do well, go back and change set of meta-parameters given to CV
  - GOOD: do CV with set of parameters on training data, see average errors on k-folds, refine set of meta-parameters and then finally complete the CV (all only using the training data)
- Data snooping will typically give you an overestimate of the performance of your model, since you “overfit” to your test
- Leave-one-out CV: higher variance, not usually a good choice



# What if I have another approach to selecting meta-parameters?

- Imagine that you are comparing two algorithms, with proposed approaches for automatically selecting their meta-parameters
- These are typically called “meta-parameter free”
- How does this change the experiment?
- What are you trying to test?




# Regression objectives

- We have looked at  $l_2$  error for estimating parameters (i.e., as an objective) and to measure performance
- Other options:
  - $l_1$  error — can be difficult to optimize, still a useful measure of error
  - smooth  $l_1$  — smooth and convex, easier to optimize, not usually used as a measure of error (unless reporting accuracy of optimizer)
  - R-squared — coefficient of determination
  - Variance unexplained
  - Percentage error — rescale by magnitude of values



# R-squared measure

- Also called “coefficient of determination”



- The sum of squares of residuals, also called the **residual sum of squares**:

$$SS_{\text{res}} = \sum_i (y_i - f_i)^2$$

- The **total sum of squares** (proportional to the **variance** of the data):

$$SS_{\text{tot}} = \sum_i (y_i - \bar{y})^2,$$



# R-squared is monotone in number of features

- As add more features, the R-squared measure cannot decrease. Why?

$$R^2 \equiv 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}, \quad SS_{\text{res}} = \sum_i (y_i - f_i)^2 \quad SS_{\text{tot}} = \sum_i (y_i - \bar{y})^2,$$

- Is this an issue?
- Alternative: adjusted R-squared — penalize the number of explanatory variables (features)



# Variance explained

- Variance explained = variance of predictor to mean of y
- R-squared = Variance explained / Total variance
- FractionVarianceUnexplained(predictor)
  - =  $\text{MSE}[\text{predictor}] / \text{Var}[y]$
  - = 1 - R-squared



# Percentage error

- If use error  $\| \text{val1} - \text{val2} \|$ , and get 0.1, is this good?
- One option: mean percentage error (issues?)
- Another option: mean absolute percentage error (MAPE)

$$M = \frac{1}{n} \sum_{t=1}^n \left| \frac{A_t - F_t}{A_t} \right|,$$

- Another option: symmetric MAPE

$$\text{SMAPE} = \frac{1}{n} \sum_{t=1}^n \frac{|F_t - A_t|}{(|A_t| + |F_t|)/2}$$



# Classification terminology

- True positives — samples predicted by classifier to be positive that have true label positive
- False positives — samples predicted by classifier to be positive that have true label negative
- True negatives — samples predicted by classifier to be negative that have true label negative
- False negatives — samples predicted by classifier to be negative that have true label positive





# Classification measures

Name	Symbol	Definition
Classification error	$error$	$error = \frac{fp+fn}{tp+fp+tn+fn}$
Classification accuracy	$accuracy$	$accuracy = 1 - error$
True positive rate	$tpr$	$tpr = \frac{tp}{tp+fn}$
False negative rate	$fnr$	$fnr = \frac{fn}{tp+fn}$
True negative rate	$tnr$	$tnr = \frac{tn}{tn+fp}$
False positive rate	$fpr$	$fpr = \frac{fp}{tn+fp}$
Precision	$pr$	$pr = \frac{tp}{tp+fp}$
Recall	$rc$	$rc = \frac{tp}{tp+fn}$



# Why these specific values?

- These measures exist for multiple reasons
- Separate the importance of false positives and false negatives
  - In some cases, much more hazardous to have a false positive than a false negative (or vice versa) e.g.
- Avoid issues with imbalanced datasets



# Confusion Matrix for binary classification

Predicted class	True class	
	0	1
0	$N_{00}$ 😊	$N_{01}$ 😞
1	$N_{10}$ 😞	$N_{11}$ 😊

$$Accuracy = \frac{N_{00} + N_{11}}{N_{00} + N_{10} + N_{01} + N_{11}}$$

Number of data points whose true class was 0 but predicted class was 1.

$$Error = 1 - Accuracy$$



# Classification accuracy and error

$$Accuracy = \frac{N_{correct}}{N}$$

$N_{correct}$ : number of correctly classified data points

$N$ : total number of data points

$$Error = 1 - \frac{N_{correct}}{N}$$

Another naming convention:

$N_{00}$  = number of true negatives

$N_{01}$  = number of false negatives

$N_{10}$  = number of false positives

$N_{11}$  = number of true positives




# Example of importance of measures: imbalanced datasets

16 data points have class 0 (majority class)

4 data points have class 1 (minority class)

---



---

**Trivial classifier:** always predict majority class

Accuracy of a trivial classifier is:  $16/20 = 80\%$

---

**Random classifier:** predict class 0 with probability 0.8 and class 1 with probability 0.2

Accuracy of the random classifier: 68%

$(0.8^2 + 0.2^2 = 0.68)$



# More on accuracy

$$sn = \frac{N_{11}}{N_{01} + N_{11}} \quad \longleftarrow \quad \text{Sensitivity or accuracy on data points whose class is 1. Also called true positive rate.}$$

$$sp = \frac{N_{00}}{N_{10} + N_{00}} \quad \longleftarrow \quad \text{Specificity or accuracy on data points whose class is 0. Also called true negative rate.}$$

$$1 - sn = 1 - \frac{N_{11}}{N_{01} + N_{11}} \quad \longleftarrow \quad \text{False negative rate.}$$

$$1 - sp = 1 - \frac{N_{00}}{N_{10} + N_{00}} \quad \longleftarrow \quad \text{False positive rate.}$$

$$Accuracy_B = \frac{sn + sp}{2}$$

$\longleftarrow$  **Balanced-sample accuracy**



# F-measure

$$F_{\beta} = \frac{(1 + \beta^2) \cdot \text{true positive}}{(1 + \beta^2) \cdot \text{true positive} + \beta^2 \cdot \text{false negative} + \text{false positive}}.$$

- Related measure, called F1 score has beta=1
  - best value at 1, worst at 0
- Weighted values of true positives and false negatives
- Focus on positives can be switched to negatives just by considering negatives as positive



# More on accuracy

$$rc = \frac{N_{11}}{N_{01} + N_{11}} \quad \longleftarrow$$

**Recall is accuracy on data points whose class is 1. Same as sensitivity or true positive rate.**

$$pr = \frac{N_{11}}{N_{10} + N_{11}} \quad \longleftarrow$$

**Precision is accuracy on data points that were predicted as 1. Also called positive predictive value.**

$$1 - rc = 1 - \frac{N_{11}}{N_{01} + N_{11}} \quad \longleftarrow$$

**False negative rate.**

$$1 - pr = 1 - \frac{N_{11}}{N_{10} + N_{11}} \quad \longleftarrow$$

**False discovery rate. Important: it is very different from the false positive rate!!!**

$$F_{\beta} = (1 + \beta^2) \cdot \frac{pr \cdot rc}{\beta^2 \cdot pr + rc}$$




**F-measure.**





# Precision and recall

- Example: when a search engine returns 30 pages only 20 of which were relevant while failing to return 40 additional relevant pages, its precision is  $20/30 = 2/3$  while its recall is  $20/60 = 1/3$ .



How many selected items are relevant?


$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$




# ROC Curve






# ROC space





# ROC space


A			B			C		
TP=63	FP=28	91	TP=77	FP=77	154	TP=24	FP=88	112
FN=37	TN=72	109	FN=23	TN=23	46	FN=76	TN=12	88
100	100	200	100	100	200	100	100	200
TPR = 0.63			TPR = 0.77			TPR = 0.24		
FPR = 0.28			FPR = 0.77			FPR = 0.88		
PPV = 0.69			PPV = 0.50			PPV = 0.21		
F1 = 0.66			F1 = 0.61			F1 = 0.22		
ACC = 0.68			ACC = 0.50			ACC = 0.18		






# ROC Curve example

e.g., diseased people, healthy people  
blood protein levels normally distributed  
Parameter that changes: threshold



TP	FP
FN	TN
1	1






# Area under the curve

- AUC or AUCROC gives the area under the ROC curve
- AUC is equal to the probability that a classifier will rank a randomly chosen positive instance higher than a randomly chosen negative one
- Some issues in using AUC to compare classifiers (see “Measuring classifier performance: a coherent alternative to the area under the ROC curve”, Hand, JMLR, 2009)
  - can give unequal importance to a FPR or TPR for different classifiers



# Classification objectives





# Clipped objectives




Figure 1: Comparing standard losses (dashed) with corresponding “clipped” losses (solid),  $\rho$ -relaxed losses (dotted), and non-convex robust losses (dash-dotted). **Left:** squared loss (dashed), clipped (solid),  $1/3$ -relaxed (dotted), robust Geman and McClure loss [2] (dash-dotted). **Center:** SVM hinge loss (dashed), clipped [27, 30] (solid),  $1/2$ -relaxed (upper dotted), robust  $1 - \tanh(y\hat{y})$  loss [19] (dash-dotted). **Right:** Adaboost exponential loss (dashed), clipped (solid),  $1/2$ -relaxed (upper dotted), robust  $1 - \tanh(y\hat{y})$  loss [19] (dash-dotted).

see: “Relaxed Clipping: A Global Training Method for Robust Regression and Classification”, Yu et al, 2010





# Miscellaneous

- Understand the expected properties of your estimator
- Example: give NN the same training data, in the same order, with the same starting point —> should it produce the same final set of weights?
- Example: with enough samples, should I get close to the true parameters? for any small random subset, how different might my learned parameters be?