

B555: Machine Learning Notes

Predrag Radivojac and Martha White

November 18, 2015

Table of Contents

1	Introduction to Probabilistic Modeling	4
1.1	Probability Theory	7
1.1.1	Axioms of probability	7
1.1.2	Probability mass functions	9
1.1.3	Probability density functions	13
1.1.4	Multidimensional distributions	16
1.1.5	Elementary conditional probabilities	18
1.1.6	Independence of events	18
1.1.7	Interpretation of probability	20
1.2	Random Variables	20
1.2.1	Formal definition of random variable	23
1.2.2	Joint and marginal distributions	25
1.2.3	Conditional distributions	27
1.2.4	Independence of random variables	28
1.2.5	Expectations and moments	30
1.2.6	Mixtures of distributions	34
1.2.7	Graphical representation of probability distributions	36
2	Basic Principles of Parameter Estimation	41
2.1	Maximum a posteriori and maximum likelihood estimation	42
3	Introduction to Prediction Problems	50
3.1	Problem Statement	50
3.2	Useful notation	52
3.3	Optimal classification and regression models	52
3.4	Bayes Optimal Models	57
4	Linear Regression and Its Applications	60
4.1	Maximum likelihood formulation	60
4.2	Ordinary Least-Squares (OLS) Regression	63
4.2.1	Weighted error function	66
4.2.2	Predicting multiple outputs simultaneously	67
4.2.3	Expectation and variance for the solution vector	67
4.3	An Algebraic Perspective	68

4.3.1	The four fundamental subspaces	69
4.3.2	Minimizing $\ \mathbf{Ax} - \mathbf{b}\ _2$	71
4.4	Linear regression for non-linear problems	74
4.4.1	Polynomial curve fitting	74
4.5	Practical considerations	77
4.5.1	Stability and sensitivity of the solutions	77
4.5.2	Regularization	78
4.5.3	Handling big data sets	79
5	Generalized Linear Models	81
5.1	Loglinear link and Poisson distribution	82
5.2	Exponential family of distributions	83
5.3	Formalizing generalized linear models	86
5.4	Logistic regression	87
6	Linear Classifiers	88
6.1	Logistic regression	89
6.1.1	Predicting class labels	89
6.1.2	Maximum conditional likelihood estimation	90
6.1.3	Stochastic optimization for logistic regression	93
6.1.4	Issues with minimizing Euclidean distance	94
6.2	Naive Bayes Classifier	96
6.2.1	Binary features and linear classification	97
6.2.2	Continuous naive Bayes	99
6.3	Multinomial logistic regression	100
7	Representations for machine learning	103
7.1	Radial basis function networks and kernel representations	103
7.2	Learning representations	105
7.2.1	Neural networks	105
7.2.2	Unsupervised learning and matrix factorization	111
8	Empirical Evaluation	116
8.1	Comparison of Learning Algorithms	116
8.2	Performance of Classification Models	117
9	Advanced topics	119
9.1	Bayesian estimation	119
9.2	Parameter estimation for mixtures of distributions	121
9.3	Boosting	128

Bibliography	129
A Optimization background	130
A.1 Second order optimization: Newton-Raphson method	130

Chapter 1

Introduction to Probabilistic Modeling

Modeling the world around us and making predictions about the occurrence of events is a multidisciplinary endeavor standing on the solid foundations of probability theory, statistics, and computer science. Although intertwined in the process of modeling, these fields have relatively discernible roles and can be, to a degree, studied individually. Probability theory brings the mathematical infrastructure, firmly grounded in its axioms, for manipulating probabilities and equips us with a broad range of models with well-understood theoretical properties. Statistics contributes frameworks to formulate inference and the process of narrowing down the model space based on the observed data and our experience in order to find, and then analyze, solutions. Computer science provides us with theories, algorithms, and software to manage the data, compute the solutions, and study the relationship between solutions and available resources (time, space, computer architecture, etc.). As such, these three disciplines form the core quantitative framework for all of empirical science and beyond.

Probability theory and statistics have a relatively long history; the formal origins of both can be traced to the 17th century. Probability theory was developed out of efforts to understand games of chance and gambling. The correspondence between Blaise Pascal and Pierre de Fermat in 1654 serves as the oldest record of modern probability theory. Statistics, on the other hand, originated from data collection initiatives and attempts to understand trends in the society (e.g. manufacturing, mortality causes, value of land) and political affairs (e.g. public revenues, taxation, armies). The two disciplines started to merge in the 18th century with the use of data for inferential purposes in astronomy, geography, and social sciences. The increased complexity of models and availability of data in the 19th century emphasized the importance of computing machines. This contributed to establishing the foundations of the field of computer science in the 20th century, which is generally attributed to the introduction of the von Neumann architecture and formalization of the concept of an algorithm. The convergence of the three

disciplines has reached the status of a principled theory of probabilistic inference with widespread applications in science, business, medicine, military, political campaigns, etc. Interestingly, various other disciplines have also contributed to the core of probabilistic modeling. Concepts such as a Boltzmann distribution, a genetic algorithm, or a neural network illustrate the influence of disciplines such as physics, biology, psychology, and engineering.

We will refer to the process of modeling, inference, and decision making based on probabilistic models as *reasoning under uncertainty*. Some form of reasoning under uncertainty is a necessary component of everyday life. When driving, for example, we often make decisions based on our expectations about which way would be best to take. While these situations do not usually involve an explicit use of probabilities and probabilistic models, an intelligent driverless car such as Google Chauffeur must make use of them. And so must a spam detection software in an email client, a credit card fraud detection system, or an algorithm that infers whether a particular genetic mutation will result in disease. Therefore, we first need to understand the concept of probability and then introduce a formal theory to incorporate evidence (e.g. data collected from instruments) in order to make good decisions in a range of situations. At a basic level, probabilities are used to quantify the chance of the occurrence of events. As Jacob Bernoulli brilliantly put it in his work *The Art of Conjecturing* (1713), “[p]robability, [...] is the degree of certainty, and it differs from the latter as a part differs from the whole”. He later adds, “To make a conjecture [prediction] about something is the same as to measure its probability. Therefore, we define the art of conjecturing [science of prediction] or stochastics, as the art of measuring probabilities of things as accurately as possible, to the end that, in judgements and actions, we may always choose or follow that which has been found to be better, more satisfactory, safer, or more carefully considered.” The techniques of probabilistic modeling formalize many intuitive concepts. In a nutshell, they provide toolkits for rigorous mathematical analysis and inference, often in the presence of evidence, about events influenced by factors that we either do not fully understand or have no control of.

To provide a quick insight into the concept of uncertainty and modeling, consider rolling a fair six-sided die. We could accurately predict, or so we think, the outcome of a roll if we carefully incorporated the initial position, force, friction, shape defects, and other physical factors and then executed the experiment. But the physical laws may not be known, they can be difficult to incorporate or such actions may not even be allowed by the rules of the experiment. Thus, it is practically useful to simply assume that each outcome is equally likely; in fact, if we rolled the die many times,




Figure 1.1: A histogram of recordings of the commute time (in minutes) to work. The data set contains 340 measurements collected over one year, for a distance of roughly 3.1 miles. The data was modeled using a gamma family of probability distributions, with the particular location and scale parameters estimated from the raw data. The values of the gamma distribution are shown as dark circles.

we would indeed observe that each number is observed roughly equally. Assigning an equal chance (probability) to each outcome of the roll of a die provides an efficient and elegant way of modeling uncertainties inherent to the experiment.

Another, more realistic example in which collecting data provides a basis for simple probabilistic modeling is a situation of driving to work every day and predicting how long it will take us to reach the destination tomorrow. If we recorded the “time to work” for a few months we would observe that trips generally took different times depending on many internal (e.g. preferred speed for the day) and also external factors (e.g., weather, road works, encountering a slow driver). While these events, if known, could be used to predict the exact duration of the commute, it is unrealistic to expect to have full information. Therefore, it is useful to provide ways of aggregating external factors via collecting data over a period of time and providing the distribution of the commute time. Such a distribution, in the absence of any other information, would then facilitate reasoning about events such as making it on time to an important meeting at 9 am. One way of using the recorded data is to create histograms and calculate percentiles. Another would be to estimate the parameters of some mathematical function that fits the data well. Both approaches are illustrated in Figure 1.1.

As the examples above suggest, the techniques of probabilistic modeling provide a formalism for dealing with repetitive “experiments” influenced by a number of external factors over which we have little control or knowledge. However, we shall see later that probabilities need not be assigned only to

events that repeat, but to any event in general. As long as they are assigned according to the formal axioms of probability, we can make inferences because the mathematical formalism does not depend on how the probabilities are assigned. This provides us an opportunity to incorporate our assumptions and existing knowledge into modeling, including the subjective assessments (beliefs) about occurrence of non-repetitive events. But let us start from the beginning.

1.1 Probability Theory

Probability theory can be seen as a branch of mathematics that deals with set functions. At the heart of probability theory is the concept of an *experiment*. An experiment can be the process of tossing a coin, rolling a die, checking the temperature tomorrow or figuring out the location of one's keys. When carried out, each experiment has an *outcome*, which is an element “drawn” from a set of predefined options, potentially infinite in size. The outcome of a roll of a die is a number between one and six; the temperature tomorrow is typically an integer or sometimes a real number; the outcome of the location of one's keys can be a discrete set of places such as a kitchen table, under a couch, in office etc. In many ways, the main goal of probabilistic modeling is to formulate a particular question or a hypothesis pertaining to the physical world as an experiment, collect the data, and then construct a model. Once a model is created, we can compute quantitative measures of sets of outcomes we are interested in and assess the confidence we should have in these measures.

1.1.1 Axioms of probability

We start by introducing the *axioms of probability*. Let the *sample space* (Ω) be a non-empty set of outcomes of the experiment and the *event space* (\mathcal{F}) be a non-empty set of subsets of Ω such that

1. $A \in \mathcal{F} \Rightarrow A^c \in \mathcal{F}$
2. $A_1, A_2, \dots \in \mathcal{F} \Rightarrow \bigcup_{i=1}^{\infty} A_i \in \mathcal{F}$

where A and all A_i 's are *events*, and A^c is the complement of A ; i.e. $A^c = \Omega - A$. If both conditions hold, \mathcal{F} is called a sigma field, or sigma algebra, and is a set of so-called measurable events.¹ The tuple (Ω, \mathcal{F}) is then called

¹This terminology is due to historic reasons; so though it sounds complex, do not be frightened by this name. Generally, one can simply think of it as the set of events.

a *measurable space*.

It is important to emphasize that the definition of sigma field requires that \mathcal{F} be closed under both finite and countably infinite number of basic set operations (union, intersection, complementation and set difference). The operations union and complementation are in the definition. For intersection, we can use De Morgan's laws: $\cup A_i = (\cap A_i^c)^c$ and $\cap A_i = (\cup A_i^c)^c$. Any intersection of sets in \mathcal{F} must again be in \mathcal{F} because \mathcal{F} is closed under union and complementation. Therefore, a sigma field is also closed under intersection. Similarly for set difference, we can write $A_1 \setminus A_2 = (A_1 \cap A_2)^c \cap A_1$ implies $A_1 \setminus A_2 \in \mathcal{F}$ because $A_1 \cap A_2 \in \mathcal{F} \implies (A_1 \cap A_2)^c \in \mathcal{F} \implies (A_1 \cap A_2)^c \cap A_1 \in \mathcal{F}$. Observe that all the above conditions imply that $\Omega \in \mathcal{F}$ and $\emptyset \in \mathcal{F}$, where \emptyset is the empty set.

Let (Ω, \mathcal{F}) be a measurable space. Any function $P : \mathcal{F} \rightarrow [0, 1]$ such that

1. (unit measure) $P(\Omega) = 1$
2. (σ -additivity) Any countable sequence of disjoint events $A_1, A_2, \dots \in \mathcal{F}$ (i.e., $A_i \cap A_j = \emptyset \forall i, j$) satisfies $P(\cup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} P(A_i)$

is called a *probability measure* or a *probability distribution* and is said to satisfy the axioms of probability.² The tuple (Ω, \mathcal{F}, P) is called the *probability space*.

The beauty of these axioms lies in their compactness and elegance. Many useful expressions can be derived from the axioms of probability. For example, it is obvious that $P(\emptyset) = 0$ or $P(A^c) = 1 - P(A)$. Similarly, closure under infinite unions of disjoint sets (σ -additivity) implies finite closure (additivity), because the remaining sets can be set to the empty set \emptyset : $\forall A_1, A_2 \in \mathcal{F}$ with $A_1 \cap A_2 = \emptyset$, set $A_i = \emptyset$ for $i > 2$ to get $P(A_1 \cup A_2) = P(\cup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} P(A_i) = P(A_1) + P(A_2)$. Another formula that is particularly important can be derived by considering a *partition* of the sample space; i.e. a set of k non-overlapping sets $\{B_i\}_{i=1}^k$ such that $\Omega = \cup_{i=1}^k B_i$. That is, if A is any set in Ω and if $\{B_i\}_{i=1}^k$ is a partition of Ω

²It seems intuitive that the second condition could be replaced with a union of finite sets (the simpler requirement of additivity rather than σ -additivity). However, for σ -fields, closure under *finite* unions may not result in closure under *infinite* unions.

it follows that

$$\begin{aligned}
P(A) &= P(A \cap \Omega) \\
&= P\left(A \cap \left(\bigcup_{i=1}^k B_i\right)\right) \\
&= P\left(\bigcup_{i=1}^k (A \cap B_i)\right) \\
&= \sum_{i=1}^k P(A \cap B_i),
\end{aligned} \tag{1.1}$$

where the last line followed from the axioms of probability. We will refer to this expression as the *sum rule*. Another important expression, shown here without a derivation, is that $P(A \cup B) = P(A) + P(B) - P(A \cap B)$.

It is convenient to separately consider discrete (countable) and continuous (uncountable) sample spaces. A roll of a die draws numbers from a finite space $\Omega = \{1, 2, 3, 4, 5, 6\}$. For finite and other countable sample spaces (e.g. the set of integers \mathbb{Z}), \mathcal{F} is usually the power set $\mathcal{P}(\Omega)$. An example of continuous sample space is the set of real numbers \mathbb{R} . As we shall see later, for uncountable spaces, \mathcal{F} must be a proper subset of $\mathcal{P}(\Omega)$, i.e. $\mathcal{F} \subset \mathcal{P}(\Omega)$, because there exist sets over which one cannot integrate. Technically, sample spaces can also be mixed; e.g. $\Omega = [0, 1] \cup \{2\}$, but we will not consider such spaces here. Discrete and continuous sample spaces give rise to discrete and continuous probability distributions, respectively.

Owing to many constraints in defining the distribution function P , it is clear that it cannot be chosen arbitrarily. For example, if $\Omega = [0, 1]$ and $P([0, \frac{1}{2})) = \frac{1}{2}$, we cannot arbitrarily assign $P([\frac{1}{2}, 1]) = \frac{1}{3}$ because probabilities of complement sets must sum to one. It turns out, in practice it is easier to define P indirectly, by selecting a probability mass function or a probability density function. These functions are defined directly on the sample space where we have fewer restrictions to be concerned with compared to the event space. We address these two ways of defining probability distributions next.

1.1.2 Probability mass functions

Let Ω be a discrete (finite or countably infinite) sample space and $\mathcal{F} = \mathcal{P}(\Omega)$. A function $p : \Omega \rightarrow [0, 1]$ is called a *probability mass function* (pmf) if

$$\sum_{\omega \in \Omega} p(\omega) = 1.$$

The probability of any event $A \in \mathcal{F}$ is defined as

$$P(A) = \sum_{\omega \in A} p(\omega).$$

It is straightforward to verify that P satisfies the axioms of probability and, thus, is a probability distribution.

Example 1: Consider a roll of a fair six-sided die, i.e. $\Omega = \{1, 2, 3, 4, 5, 6\}$, and the event space $\mathcal{F} = \mathcal{P}(\Omega)$. What is the probability that the outcome is a number greater than 4?

First, because the die is fair, we know that $p(\omega) = \frac{1}{6}$ for $\forall \omega \in \Omega$. Now, let A be an event in \mathcal{F} that the outcome is greater than 4, i.e. $A = \{5, 6\}$. Thus,

$$P(A) = \sum_{\omega \in A} p(\omega) = \frac{1}{3}.$$

It is important to note that P is defined on the elements of \mathcal{F} , whereas p is defined on the elements of Ω . That is, $P(\{1\}) = p(1)$, $P(\{2\}) = p(2)$, $P(\{1, 2\}) = p(1) + p(2)$, etc. \square

In discrete cases, $P(\{\omega\}) = p(\omega)$ for every $\omega \in \Omega$, and the probability of a set is always equal to the sum of probabilities of individual elements. We can define a discrete probability space by providing the tuple (Ω, \mathcal{F}, P) , but it is often much simpler to define P indirectly by assuming that $\mathcal{F} = \mathcal{P}(\Omega)$ and providing a probability mass function p . In this case we say that the probability measure P is induced by a pmf p . In fact, we rarely define (Ω, \mathcal{F}, P) directly.

A few useful pmfs

Let us now look at some families of functions that are often used as discrete probability distributions. This is by no means a comprehensive review of the subject; we shall simply focus on a few basic concepts and will later introduce other distributions as needed. For simplicity, we will often refer to both pmfs and probability distributions they induce as distribution functions. Several of the following distributions are shown in Figure 1.2.

The *Bernoulli distribution* derives from the concept of a Bernoulli trial, an experiment that has two possible outcomes: success and failure. In a Bernoulli trial, a success occurs with probability α and, thus, failure occurs with probability $1 - \alpha$. A toss of a coin (heads/tails), a basketball game (win/loss), or a roll of a die (even/odd) can all be seen as Bernoulli trials.

We model this distribution by setting a sample space to two elements and defining the probability of one of them as α . More specifically, $\Omega = \{S, F\}$ and

$$p(\omega) = \begin{cases} \alpha & \omega = S \\ 1 - \alpha & \omega = F \end{cases}$$

where $\alpha \in (0, 1)$ is a parameter. If $\Omega = \{0, 1\}$, we can compactly write the Bernoulli distribution as $p(k) = \alpha^k \cdot (1 - \alpha)^{1-k}$ for $k \in \Omega$. Here we replaced ω with k , which is a more common notation when the sample space is comprised of integers.

To be precise, the Bernoulli distribution as presented above is actually a family of discrete probability distributions, one for each α . We shall refer to each such distribution as $\text{Bernoulli}(\alpha)$. However, we do not need to concern ourselves with semantics because the correct interpretation of a family vs. individual distributions will usually be clear from the context. The Bernoulli distribution is sometimes referred to as the binary distribution.

The *Binomial distribution* is used to describe a sequence of n independent and identically distributed (i.i.d.) Bernoulli trials. At each value k in the sample space the distribution gives the probability that the success happened exactly k times out of n trials, where of course $0 \leq k \leq n$. More formally, given $\Omega = \{0, 1, \dots, n\}$, for $\forall k \in \Omega$ the binomial pmf is defined as

$$p(k) = \binom{n}{k} \alpha^k (1 - \alpha)^{n-k},$$

where $\alpha \in (0, 1)$, as before, is the parameter indicating the probability of success in a single trial. Here, the binomial coefficient

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

enumerates all ways in which one can pick k elements from a list of n elements (e.g. there are 3 different ways in which one can pick 2 elements from a group of 3 elements). We will refer to a binomial distribution with parameters n and α as $\text{Binomial}(n, \alpha)$. The experiment leading to a binomial distribution can be generalized to a situation with more than two possible outcomes. This experiment results in a multidimensional probability mass function (one dimension per possible outcome) called the *multinomial distribution*.

The *Poisson distribution* can be derived as a limit of the binomial distribution as $n \rightarrow \infty$ with a fixed expected number of successes (λ). Here,

$\Omega = \{0, 1, \dots\}$ and for $\forall k \in \Omega$

$$p(k) = \frac{\lambda^k e^{-\lambda}}{k!},$$

where $\lambda \in (0, \infty)$ is a parameter (the relationship with the binomial distribution can be obtained by taking $\alpha = \lambda/n$). The Poisson distribution is often used to model counts of events occurring sequentially and independently but with a fixed average (λ) in a particular time interval. Unlike the previous two distributions, $\text{Poisson}(\lambda)$ is defined over an infinite sample space, but still countable.

The *geometric distribution* is also used to model a sequence of independent Bernoulli trials with the probability of success α . At each point $k \in \Omega$, it gives the probability that the first success occurs exactly in the k -th trial. Here, $\Omega = \{1, 2, \dots\}$ and for $\forall k \in \Omega$

$$p(k) = (1 - \alpha)^{k-1} \alpha,$$

where $\alpha \in (0, 1)$ is a parameter. The geometric distribution, $\text{Geometric}(\alpha)$, is defined over an infinite sample space, i.e. $\Omega = \mathbb{N}$.

For the *hypergeometric distribution*, consider a finite population of N elements of two types (e.g. success and failure), K of which are of one type (e.g. success). The experiment consists of drawing n elements, without replacement, from this population such that the elements remaining in the population are equiprobable in terms of being selected in the next draw. The probability of drawing k successes out of n trials can be described as

$$p(k) = \frac{\binom{K}{k} \cdot \binom{N-K}{n-k}}{\binom{N}{n}},$$

where $0 \leq n \leq N$ and $k \leq n$. The hypergeometric distribution is intimately related to the binomial distribution where the elements are drawn with replacement ($\alpha = K/N$). There, the probability of drawing a success does not change in subsequent trials. We will refer to the hypergeometric distribution as $\text{Hypergeometric}(n, N, K)$.

The *uniform distribution* for discrete sample spaces is defined over a finite set of outcomes each of which is equally likely to occur. Here we can set $\Omega = \{1, \dots, n\}$; then for $\forall k \in \Omega$

$$p(k) = \frac{1}{n}.$$

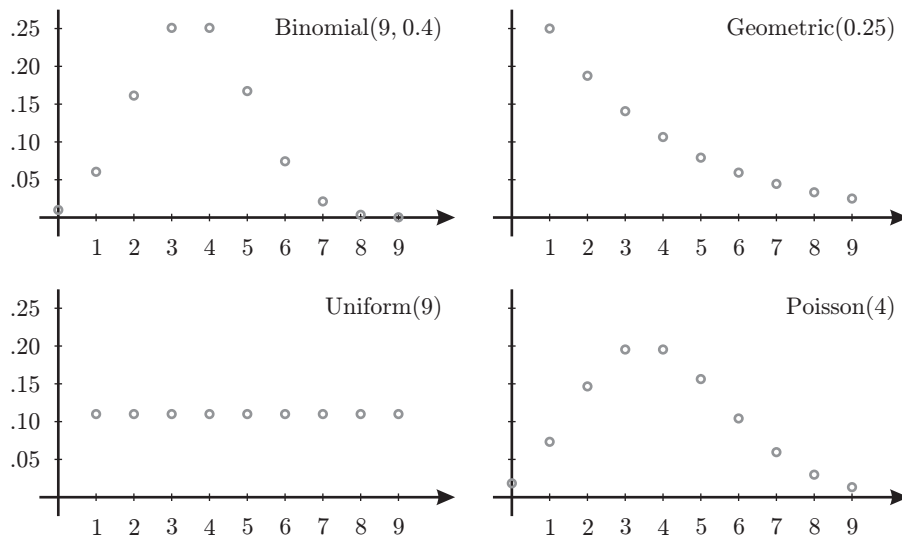


Figure 1.2: Four discrete probability mass functions.

The uniform distribution does not contain parameters; it is defined by the size of the sample space. We refer to this distribution as $\text{Uniform}(n)$. We will see later that the uniform distribution can also be defined over finite intervals in continuous spaces.

All of the functions above satisfy the definition of a probability mass function, which we can verify by summing over all possible outcomes in the sample space. Four examples are shown in Figure 1.2. In general, for discrete spaces, we can assign probabilities to outcomes quite freely. For example, one could have a table of 365 values between zero and one for the probability of a birthday falling on each day, as long as the probabilities sum to 1. We will see that for continuous spaces it is more difficult to define consistent probabilities and we will often restrict to a limited set of known distributions.

1.1.3 Probability density functions

We shall see soon that the treatment of continuous probability spaces is analogous to that of discrete spaces, with probability density functions replacing probability mass functions and integrals replacing sums. Mathematically, however, there are fundamental differences between the two situations which we should keep in mind whenever working with continuous spaces.

The main obstacle in generalizing the theory to uncountable sample spaces lies in addressing mathematical nuances involving infinitesimal calculus, the countable nature of a sigma field, and limitations of the definition of integrals. For example, there exist sets over which we cannot integrate and thus the set of events \mathcal{F} cannot be the power set of any uncountable set (e.g. \mathbb{R}). It is therefore necessary to define an adequate event space which would be applicable to a vast majority of practically important situations.

To operate with continuous sample spaces we proceed by taking $\Omega = \mathbb{R}$ and defining the Borel field. The Borel field on \mathbb{R} , denoted by $\mathcal{B}(\mathbb{R})$, is a set that contains all points in \mathbb{R} , all open, semi-open and closed intervals in \mathbb{R} , as well as sets that can be obtained by a countable number of basic set operations on them. By definition, $\mathcal{B}(\mathbb{R})$ is a sigma field; $\mathcal{B}(\mathbb{R})$ is an uncountably infinite set, but still smaller than $\mathcal{P}(\mathbb{R})$. The construction of subsets of \mathbb{R} that are not in $\mathcal{B}(\mathbb{R})$ is difficult and only of theoretical importance (e.g., Vitali sets), but nevertheless, the use of $\mathcal{P}(\mathbb{R})$ as the event space would lead to a flawed theory. Therefore, when discussing probability distributions over continuous sample spaces, we will usually take $\Omega = \mathbb{R}$ to be the sample space and implicitly consider $\mathcal{F} = \mathcal{B}(\mathbb{R})$ to be the event space.

Let now Ω be a continuous sample space and $\mathcal{F} = \mathcal{B}(\Omega)$. A function $p : \Omega \rightarrow [0, \infty)$ is called a *probability density function* (pdf) if

$$\int_{\Omega} p(\omega) d\omega = 1.$$

The probability of an event $A \in \mathcal{B}(\Omega)$ is defined as

$$P(A) = \int_A p(\omega) d\omega.$$

There are a few mathematical nuances associated with this definition. First, interestingly, the standard Riemann integration does not work for some sets in the Borel field (e.g. how would you integrate over the set of rational or irrational numbers within $[0, 1]$ for any pdf?). For that reason, probability density functions are formally defined using Lebesgue integration which allows us to integrate over all sets in $\mathcal{B}(\Omega)$. Luckily, Riemann integration, when possible, provides identical results as Lebesgue's integrals; thus, it will suffice to use Riemann integration in all situations of our interest.

Second, we mentioned before for pmfs that the probability of a singleton event $\{\omega\}$ is the value of the pmf at the sample point ω , i.e. $P(\{\omega\}) = p(\omega)$. In contrast, the value of a pdf at point ω is not a probability; it can be greater than 1. The probability at any single point, but also over any finite

or countably infinite set is 0 (i.e. they constitute a set of measure zero). One way to think about the probabilities in continuous spaces is to look at small intervals $A = [x, x + \Delta x]$ as

$$P(A) = \int_x^{x+\Delta x} p(\omega) d\omega \approx p(x) \Delta x.$$

Here, a potentially large value of the density function is compensated by the small interval Δx to result in a number between 0 and 1.

A few useful pdfs

Some important probability density functions are reviewed below. As before, the sample space will be defined for each distribution and the Borel field will be implicitly assumed as the event space.

The *uniform distribution* is defined by an equal value of a probability density function over a finite interval in \mathbb{R} . Thus, for $\Omega = [a, b]$ the uniform probability density function $\forall \omega \in [a, b]$ is defined as

$$p(\omega) = \frac{1}{b - a}.$$

Note that one can also define $\text{Uniform}(a, b)$ by taking $\Omega = \mathbb{R}$ and setting $p(\omega) = 0$ whenever ω is outside of $[a, b]$. This form is convenient because $\Omega = \mathbb{R}$ can then be used consistently for all one-dimensional probability distributions.

The *exponential distribution* is defined over a set of non-negative numbers, i.e. $\Omega = [0, \infty)$. Its probability density function is

$$p(\omega) = \lambda e^{-\lambda \omega},$$

where $\lambda > 0$ is a parameter. As before, the sample space can be extended to all real numbers, in which case we would set $p(\omega) = 0$ for $\omega < 0$.

The *Gaussian distribution* or normal distribution is one of the most frequently used probability distributions. It is defined over $\Omega = \mathbb{R}$ as

$$p(\omega) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(\omega-\mu)^2}$$

with two parameters, $\mu \in \mathbb{R}$ and $\sigma > 0$. We will refer to this distribution as $\text{Gaussian}(\mu, \sigma^2)$ or $\mathcal{N}(\mu, \sigma^2)$. Both Gaussian and exponential distribution are




Figure 1.3: Selection of a random number (x) from the unit interval $[0, 1]$.

members of a broader family of distributions called the exponential family. We will see a general definition of this family later.

Example 2: Consider selecting a number (x) between 0 and 1 uniformly randomly (Figure 1.3). What is the probability that the number is greater than $\frac{3}{4}$ or lower than $\frac{1}{4}$?

We know that $\Omega = [0, 1]$. We define an event of interest as $A = [0, \frac{1}{4}) \cup (\frac{3}{4}, 1]$ and calculate its probability as

$$\begin{aligned} P(A) &= \int_0^{1/4} d\omega + \int_{3/4}^1 d\omega &> p(\omega) = \frac{1}{b-a} = 1 \\ &= \frac{1}{2}. \end{aligned}$$

Because the probability of any individual event in a continuous case is 0, there is no difference in integration if we consider open or closed intervals. \square

1.1.4 Multidimensional distributions

It is often convenient to think of the sample space as a multidimensional space. In the discrete case, one can think of the sample space Ω as a multidimensional array or as a k -dimensional tensor (note a matrix is a 2-d tensor). That is, $\Omega = \Omega_1 \times \Omega_2 \times \dots \times \Omega_k$, where Ω_i can be seen as the sample space along dimension i . Then, any function $p : \Omega_1 \times \Omega_2 \times \dots \times \Omega_k \rightarrow [0, 1]$ is called a multidimensional probability mass function if

$$\sum_{\omega_1 \in \Omega_1} \dots \sum_{\omega_k \in \Omega_k} p(\omega_1, \omega_2, \dots, \omega_k) = 1.$$

One example of the multidimensional pmf is the multinomial distribution, which generalizes the binomial distribution to the case when the number of outcomes in any trial is a positive integer $k \geq 2$.

The *multinomial distribution* is used to model a sequence of n independent and identically distributed (i.i.d.) trials with k outcomes. At each point (n_1, n_2, \dots, n_k) in the sample space, the multinomial pmf provides the probability that the outcome 1 occurred n_1 times, outcome 2 occurred n_2 times, etc. Of course, $0 \leq n_i \leq n$ for $\forall i$ and $\sum_{i=1}^k n_i = n$. More formally, given the sample space $\Omega = \{0, 1, \dots, n\}^k$, the multinomial pmf is defined as

$$p(n_1, n_2, \dots, n_k) = \begin{cases} \binom{n}{n_1, n_2, \dots, n_k} \alpha_1^{n_1} \alpha_2^{n_2} \dots \alpha_k^{n_k} & n_1 + n_2 + \dots + n_k = n \\ 0 & \text{otherwise} \end{cases}$$

where α_i 's are positive coefficients such that $\sum_{i=1}^k \alpha_i = 1$. That is, each coefficient α_i gives the probability of outcome i in any trial. The multinomial coefficient

$$\binom{n}{n_1, n_2, \dots, n_k} = \frac{n!}{n_1! n_2! \dots n_k!}$$

generalizes the binomial coefficient by enumerating all ways in which one can distribute n balls into k boxes such that the first box contains n_1 balls, the second box n_2 balls, etc. An experiment consisting of n tosses of a fair six-sided die and counting the number of occurrences of each number can be described by a multinomial distribution. Clearly, in this case $\alpha_i = 1/6$, for each $i \in \{1, 2, 3, 4, 5, 6\}$.

In the continuous case, we can think of the sample space as the k -dimensional Euclidean space; i.e. $\Omega = \mathbb{R}^k$ and an event space as $\mathcal{F} = \mathcal{B}(\mathbb{R})^k$. Then, the k -dimensional probability density function can be defined as any function $p : \mathbb{R}^k \rightarrow [0, \infty)$ such that

$$\int_{-\infty}^{\infty} \dots \int_{-\infty}^{\infty} p(\omega_1, \omega_2, \dots, \omega_k) d\omega_1 \dots d\omega_k = 1.$$

The *multivariate Gaussian distribution* is a generalization of the Gaussian or normal distribution to the k -dimensional case, with $\Omega = \mathbb{R}^k$. It is defined as

$$p(\boldsymbol{\omega}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp \left(-\frac{1}{2} (\boldsymbol{\omega} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\boldsymbol{\omega} - \boldsymbol{\mu}) \right),$$

with parameters $\boldsymbol{\mu} \in \mathbb{R}^k$ and a positive definite k -by- k matrix $\boldsymbol{\Sigma}$ (note that $|\boldsymbol{\Sigma}|$ is the determinant of $\boldsymbol{\Sigma}$). We will refer to this distribution as $\text{Gaussian}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ or $\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$.

1.1.5 Elementary conditional probabilities

Let (Ω, \mathcal{F}, P) be a probability space and B an event that already occurred. We are interested in the probability that event A also occurred, i.e. $P(A|B)$. The elementary conditional probability is defined as

$$P(A|B) = \frac{P(A \cap B)}{P(B)}, \quad (1.2)$$

where $P(B) > 0$. From this expression, which is sometimes referred to as *product rule*, we can now derive two important formulas. The first one is *Bayes' rule*

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}.$$

The second formula, referred to as the *chain rule*, applies to a collection of k events $\{A_i\}_{i=1}^k$ and can be derived by recursively applying the product rule. Then,

$$P(A_1 \cap A_2 \dots \cap A_k) = P(A_1)P(A_2|A_1) \dots P(A_k|A_1 \cap A_2 \dots \cap A_{k-1}).$$

In some situations we refer to the probability $P(A)$ as *prior probability* because it quantifies the likelihood of occurrence of event A in absence of any other information or evidence. The probability $P(A|B)$ is referred to as *posterior probability* because it quantifies the likelihood about A in the presence of additional information (event B). The product rule from Eq. (1.2) has long history; it was derived by Abraham de Moivre in 1718.

One way to think about conditional probabilities is to consider that the experiment has already been conducted, but that we do not know the outcome yet. For example, a fair die has been rolled and we are interested in an event that the outcome was 4; i.e. $A = \{4\}$. The prior probability of event A is $P(A) = \frac{1}{6}$. But imagine that someone had observed the experiment and told us that the number was even ($B = \{2, 4, 6\}$). The probability after hearing this news becomes $P(A|B) = \frac{1}{3}$. Proper estimation of posterior probabilities from data is central to statistical inference.

1.1.6 Independence of events

Let (Ω, \mathcal{F}, P) be a probability space. Two events A and B from \mathcal{F} are defined as *independent* if

$$P(A \cap B) = P(A) \cdot P(B)$$




Figure 1.4: Visualization of dependent and independent events. Events A and B on the left are dependent because the occurrence of one excludes the occurrence of the other one. Events C and D on the right are independent. Each event occupies $1/4$ of the sample space Ω , while their intersection occupies $1/16$ of the sample space.

or, alternatively, if $P(A|B) = P(A)$ or $P(B|A) = P(B)$. More broadly, two or more events are (mutually or jointly) independent, if the probability of intersection of any group of events (of size two, three, etc.) can be expressed as the product of probabilities of individual events. For k events, there are $2^k - k - 1$ independence tests, one for each subset excluding the empty set and singletons.

It is important to distinguish between mutually exclusive events and independent events. Mutually exclusive events are in fact never independent because the knowledge that the outcome of the experiment belongs to event A excludes the possibility that it is in B (Figure 1.4). It is often difficult, and quite non-intuitive, to simply look at events and conclude whether they are independent or not. One should (almost) always calculate $P(A \cap B)$ and $P(A) \cdot P(B)$ and numerically verify independence. Sometimes there may exist deep physical reasons why particular events are independent or assumed to be independent. In other occasions it may just be a numerical coincidence.

Let (Ω, \mathcal{F}, P) be a probability space and A , B , and C some events from \mathcal{F} . Events A and B are defined as *conditionally independent* given C if

$$P(A \cap B|C) = P(A|C) \cdot P(B|C)$$

or, alternatively, if $P(A|B \cap C) = P(A|C)$. Independence between events does not imply conditional independence and, likewise, conditional independence between events does not imply their independence. We shall see an example later.

1.1.7 Interpretation of probability

There are two different interpretations of probability: a *frequentist* (*objectivist*) and a *subjectivist* one. In the frequentist view, to determine probabilities one needs repetitive experiments, e.g. Bernoulli trials. For example, in rolling a die 100 times, one could obtain a good approximation of $p(6)$ and even a better one if they rolled it 100,000 times. As the number of experiments increases, the fraction of favorable outcomes converges to the probability of an event. Thus, probability is defined through a limit in an infinite series of trials.

However, consider the following question: What is the probability of war between England and France next year? They certainly have been in many wars before (e.g. the Hundred Years' War) so perhaps it is not zero. Such an event cannot be tested multiple times, as there is only one next year. Yet we may have a pretty good sense of probability of war and can come up with a reasonable number (a subjectivist view). This number is our *degree of belief* or *conviction* derived from particular parameters of the situation. In the frequentist view, this probability would be undefined, or we would have to resort to *ad hoc* methods; for example, the probability of war is ϵ , where ϵ is the fraction of war days since formation of England and France (this is not a pure frequentist view). Even further, we could take probabilities of war between any two neighboring countries during their existence and apply this knowledge to England and France.

Interestingly, it can be argued that both frequentist and subjectivist interpretations of probability originated around the same time, perhaps in the same work. In *The Art of Conjecturing* Bernoulli proved the weak law of large numbers (undoubtedly a frequentist's concept) but also introduced a subjective notion of probability as something that depends on a person's knowledge. The good news is, the philosophical debate has almost no bearing on the use of probability theory in practice. No matter how probabilities are assigned, the mechanics of probabilistic manipulations are the same and valid, as long as the assignments adhere to the axioms of probability.

1.2 Random Variables

Until now we operated on relatively simple sample spaces and produced measure functions over sets of outcomes. In many situations, however, we would like to use probabilistic modeling on sets (e.g. a group of people) where elements can be associated with various descriptors. For example, a person may be associated with his/her age, height, citizenship, IQ, or marital status

and we may be interested in events related to such descriptors. In other situations, we may be interested in transformations of sample spaces such as those corresponding to digitizing an analog signal from a microphone into a set of integers based on some set of voltage thresholds. The mechanism of a *random variable* facilitates addressing all such situations in a simple, rigorous and unified manner.

A random variable is a variable that, from the observer's point of view, takes values non-deterministically, with generally different preferences for different outcomes. Mathematically, however, it is defined as function that maps one sample space into another, with a few technical caveats we will introduce later.

Let us motivate the need for random variables. Consider a probability space (Ω, \mathcal{F}, P) , where Ω is a set of people and let us investigate the probability that a randomly selected person $\omega \in \Omega$ has a cold (we may assume we have a diagnostic method and tools at hand). We start by defining an event A as

$$A = \{\omega \in \Omega : Disease(\omega) = \text{cold}\}$$

and simply calculate the probability of this event. This is a perfectly legitimate approach, but it can be much simplified using the random variable mechanism. We first note that, technically, our diagnostic method corresponds to a function $Disease : \Omega \rightarrow \{\text{cold}, \text{not cold}\}$ that maps the sample space Ω to a new binary sample space $\Omega_{Disease} = \{\text{cold}, \text{not cold}\}$. Even more interestingly, our approach also maps the probability distribution P to a new probability distribution $P_{Disease}$ that is defined on $\Omega_{Disease}$. For example, we can calculate $P_{Disease}(\{\text{cold}\})$ from the probability distribution of the aforementioned event A ; i.e. $P_{Disease}(\{\text{cold}\}) = P(A)$. This is a cluttered notation so we may wish to simplify it by using $P(Disease = \text{cold})$, where $Disease$ is a “random variable”.

We will use capital letters X, Y, \dots to denote random variables (such as $Disease$) and lowercase letters x, y, \dots to indicate elements (such as “cold”) of $\Omega_X, \Omega_Y \dots$. Generally, we write such probabilities as $P(X = x)$, which is a notational relaxation from $P(\{\omega : X(\omega) = x\})$. Before we proceed to formally define random variables, we shall look at two illustrative examples.

Example 3: Consecutive tosses of a fair coin. Consider a process of three coin tosses and two random variables, X and Y , defined on the sample space. We define X as the number of heads in the first toss and Y as the number of heads over all three tosses. Our goal is to find the probability spaces that are created after the transformations.

First, $\Omega = \{HHH, HHT, HTH, HTT, THH, THT, TTH, TTT\}$ and

ω	HHH	HHT	HTH	HTT	THH	THT	TTH	TTT
$X(\omega)$	1	1	1	1	0	0	0	0
$Y(\omega)$	3	2	2	1	2	1	1	0

Let us only focus on variable Y . Clearly, $Y : \Omega \rightarrow \{0, 1, 2, 3\}$ but we also need to find \mathcal{F}_Y and P_Y . To calculate P_Y , a simple approach is to find its pmf p_Y . For example, let us calculate $p_Y(2) = P_Y(\{2\})$ as

$$\begin{aligned}
P_Y(\{2\}) &= P(Y = 2) \\
&= P(\{\omega : Y(\omega) = 2\}) \\
&= P(\{HHT, HTH, THH\}) \\
&= \frac{3}{8},
\end{aligned}$$

because of the uniform distribution in the original space (Ω, \mathcal{F}, P) . In a similar way, we can calculate that $P(Y = 0) = P(Y = 3) = 1/8$, and that $P(Y = 1) = 3/8$. In this example, we took that $\mathcal{F} = \mathcal{P}(\Omega)$ and $\mathcal{F}_Y = \mathcal{P}(\Omega_Y)$. As a final note, we mention that all the randomness is defined in the original probability space (Ω, \mathcal{F}, P) and that the new probability space $(\Omega_Y, \mathcal{F}_Y, P_Y)$ simply inherits it through a deterministic transformation. \square

Example 4: Quantization. Consider (Ω, \mathcal{F}, P) where $\Omega = [0, 1]$, $\mathcal{F} = \mathcal{B}(\Omega)$, and P is induced by a uniform pdf. Define $X : \Omega \rightarrow \{0, 1\}$ as

$$X(\omega) = \begin{cases} 0 & \omega \leq 0.5 \\ 1 & \omega > 0.5 \end{cases}$$

and find the transformed probability space.

Technically, we have changed the sample space to $\Omega_X = \{0, 1\}$. For an event space $\mathcal{F}_X = \mathcal{P}(\Omega_X) = \{\emptyset, \{0\}, \{1\}, \{0, 1\}\}$ we would like to understand the new probability distribution P_X . We have

$$\begin{aligned}
p_X(0) &= P_X(\{0\}) \\
&= P(X = 0) \\
&= P(\{\omega : \omega \in [0, 0.5]\}) \\
&= \frac{1}{2}
\end{aligned}$$

and

$$\begin{aligned}
p_X(1) &= P_X(\{1\}) \\
&= P(X = 1) \\
&= P(\{\omega : \omega \in (0.5, 1]\}) \\
&= \frac{1}{2}
\end{aligned}$$

From here we can easily see that $P_X(\{0, 1\}) = 1$ and $P_X(\emptyset) = 0$, and so P_X is indeed a probability distribution. Again, P_X is naturally defined using P . Thus, we have transformed the probability space (Ω, \mathcal{F}, P) into $(\Omega_X, \mathcal{F}_X, P_X)$.

□

1.2.1 Formal definition of random variable

We now formally define a random variable. Given a probability space (Ω, \mathcal{F}, P) , a random variable X is a function $X : \Omega \rightarrow \Omega_X$ such that for every $A \in \mathcal{B}(\Omega_X)$ it holds that $\{\omega : X(\omega) \in A\} \in \mathcal{F}$. It follows that

$$P_X(A) = P(\{\omega : X(\omega) \in A\}).$$

It is important to mention that, by default, we defined the event space of a random variable to be the Borel field of Ω_X . This is convenient because a Borel field of a countable set Ω is its power set. Thus, we are working with the largest possible event spaces for both discrete and continuous random variables.

Consider now a *discrete random variable* X defined on (Ω, \mathcal{F}, P) . As we can see from the previous examples, the probability distribution for X can be found as

$$\begin{aligned}
p_X(x) &= P_X(\{x\}) \\
&= P(\{\omega : X(\omega) = x\})
\end{aligned}$$

for $\forall x \in \Omega_X$. The probability of an event A can be found as

$$\begin{aligned}
P_X(A) &= P(\{\omega : X(\omega) \in A\}) \\
&= \sum_{x \in A} p_X(x)
\end{aligned}$$

for $\forall A \subseteq \Omega_X$.

The case of *continuous random variables* is more complicated, but reduces to an approach that is similar to that of discrete random variables. Here we first define a *cumulative distribution function* (cdf) as

$$\begin{aligned} F_X(t) &= P_X(\{x : x \leq t\}) \\ &= P_X((-\infty, t]) \\ &= P(\{\omega : X(\omega) \leq t\}) \\ &= P(X \leq t), \end{aligned}$$

where $P(X \leq t)$, as before, presents a minor abuse of notation. If the cumulative distribution function is differentiable, the probability density function of a continuous random variable is defined as

$$p_X(x) = \left. \frac{dF_X(t)}{dt} \right|_{t=x}.$$

Alternatively, if p_X exists, then

$$F_X(t) = \int_{-\infty}^t p_X(x) dx,$$

for each $t \in \mathbb{R}$. Our focus will be exclusively on random variables that have their probability density functions; however, for a more general view, we should always keep in mind “if one exists” when referring to pdfs.

The probability that a random variable will take a value from interval $(a, b]$ can now be calculated as

$$\begin{aligned} P_X((a, b]) &= P_X(a < X \leq b) \\ &= \int_a^b p_X(x) dx \\ &= F_X(b) - F_X(a), \end{aligned}$$

which follows from the properties of integration.

Suppose now that the random variable X transforms a probability space (Ω, \mathcal{F}, P) into $(\Omega_X, \mathcal{F}_X, P_X)$. To describe the resulting probability space, we commonly use probability mass and density functions inducing P_X . For example, if P_X is induced by a Gaussian distribution with parameters μ and σ^2 , we use

$$X : \mathcal{N}(\mu, \sigma^2) \quad \text{or} \quad X \sim \mathcal{N}(\mu, \sigma^2).$$

Both notations indicate that the probability density function for the random variable X is

$$p_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}.$$

This distribution in turn provides all the necessary information about the probability space; i.e. $\Omega_X = \mathbb{R}$, and $\mathcal{F}_X = \mathcal{B}(\mathbb{R})$.

A group of k random variables $\{X_i\}_{i=1}^k$ defined on the same probability space (Ω, \mathcal{F}, P) is called a *random vector* or a multivariate (multidimensional) random variable. We have already seen an example of a random vector provided by random variables (X, Y) in Example 3. A generalization of a random vector to infinite sets is referred to as a *random process* or *stochastic process*; i.e. $\{X_i : i \in \mathcal{T}\}$, where \mathcal{T} is an index set usually interpreted as a set of time indices. In the case of discrete time indices (e.g. $\mathcal{T} = \mathbb{N}$) the random process is called a discrete-time random process; otherwise (e.g. $\mathcal{T} = \mathbb{R}$) it is called a continuous-time random process. There are many models in machine learning that deal with temporally-connected random variables (e.g. autoregressive models for time series, Markov chains, hidden Markov models). The language of random variables, through stochastic processes, nicely enables formalization of these models. Most of these notes, however, will deal with simpler settings only requiring (i.i.d.) multivariate random variables.

1.2.2 Joint and marginal distributions

Let us first look at two discrete random variables X and Y defined on the same probability space (Ω, \mathcal{F}, P) . We define the *joint probability distribution* $p_{XY}(x, y)$ of X and Y as

$$\begin{aligned} p_{XY}(x, y) &= P(X = x, Y = y) \\ &= P(\{\omega : X(\omega) = x\} \cap \{\omega : Y(\omega) = y\}). \end{aligned}$$

We can extend this to a k -dimensional random variable $\mathbf{X} = (X_1, X_2, \dots, X_k)$ and define a multidimensional probability mass function as $p_{\mathbf{X}}(\mathbf{x})$, where $\mathbf{x} = (x_1, x_2, \dots, x_k)$ is a vector of values, such that each x_i is chosen from some Ω_{X_i} .

A *marginal distribution* is defined for a subset of $\mathbf{X} = (X_1, X_2, \dots, X_k)$ by summing or integrating over the remaining variables. A marginal distribution $p_{X_i}(x_i)$ is defined as

$$p_{X_i}(x_i) = \sum_{x_1} \cdots \sum_{x_{i-1}} \sum_{x_{i+1}} \cdots \sum_{x_k} p_{\mathbf{X}}(x_1, \dots, x_k),$$

where the variable in the j -th sum takes values from Ω_{X_j} . The previous equation directly follows from Eq. (1.1) and is also referred to as *sum rule*.

In the continuous case, we define a multidimensional cdf as

$$\begin{aligned} F_{\mathbf{X}}(\mathbf{t}) &= P_{\mathbf{X}}(\{\mathbf{x} : x_i \leq t_i, i = 1 \dots k\}) \\ &= P(X_1 \leq t_1, X_2 \leq t_2, \dots, X_k \leq t_k) \end{aligned}$$

and the probability density function, if it exists, is defined as

$$p_{\mathbf{X}}(\mathbf{x}) = \frac{\partial^k}{\partial t_1 \dots \partial t_k} F_{\mathbf{X}}(t_1, \dots, t_k) \Big|_{\mathbf{t}=\mathbf{x}}.$$

The marginal density $p_{X_i}(x_i)$ is defined as

$$p_{X_i}(x_i) = \int_{x_1} \dots \int_{x_{i-1}} \int_{x_{i+1}} \dots \int_{x_k} p_{\mathbf{X}}(\mathbf{x}) dx_1 \dots dx_{i-1} dx_{i+1} \dots dx_k$$

If the sample space for each discrete random variable is seen as a countable subset of \mathbb{R} , then the probability space for any k -dimensional random variable \mathbf{X} (discrete or continuous) can be defined as $(\mathbb{R}^k, \mathcal{B}(\mathbb{R})^k, P_{\mathbf{X}})$.

Example 5: Three tosses of a fair coin (again). Consider two random variables from Example 3 and calculate their probability spaces, joint and marginal distributions. Recall X is the number of heads in the first toss and Y is the number of heads over all three tosses.

A joint probability mass function $p(x, y) = P(X = x, Y = y)$ is shown below

		Y			
		0	1	2	3
X	0	1/8	1/4	1/8	0
	1	0	1/8	1/4	1/8

but let us step back for a moment and show how we can calculate it. Let us consider two sets $A = \{\text{HHH}, \text{HHT}, \text{HTH}, \text{HTT}\}$ and $B = \{\text{HHT}, \text{HTH}, \text{THH}\}$, corresponding to the events that the first toss was heads and that there were exactly two heads over the three tosses, respectively. Now, let us look at the probability of the intersection of A and B

$$\begin{aligned} P(A \cap B) &= P(\{\text{HHT}, \text{HTH}\}) \\ &= \frac{1}{4} \end{aligned}$$

We can represent the probability of the logical statement $X = 1 \wedge Y = 2$ as

$$\begin{aligned} p_{XY}(1, 2) &= P(X = 1, Y = 2) \\ &= P(A \cap B) \\ &= P(\{\text{HHT}, \text{HTH}\}) \\ &= \frac{1}{4}. \end{aligned}$$

The marginal probability distribution can be found in a straightforward way as

$$p_X(x) = \sum_{y \in \Omega_Y} p_{XY}(x, y),$$

where $\Omega_Y = \{0, 1, 2, 3\}$. Thus,

$$\begin{aligned} p_X(0) &= \sum_{y \in \Omega_Y} p_{XY}(0, y) \\ &= \frac{1}{2}. \end{aligned}$$

We note for the end that we needed $|\Omega_X| \cdot |\Omega_Y| - 1$ numbers (because the sum must equal 1) to fully describe the joint distribution $p_{XY}(x, y)$. Asymptotically, this corresponds to an exponential growth of the number of entries in the table with the number of random variables (k). For example, if $|\Omega_{X_i}| = 2$ for $\forall X_i$, there are $2^k - 1$ free elements in the joint probability distribution. Estimating such distributions from data is intractable and is one form of the *curse of dimensionality*. □

1.2.3 Conditional distributions

The conditional probability distribution for two random variables X and Y , $p(y|x)$, is defined as

$$\begin{aligned} p_{Y|X}(y|x) &= P(Y = y | X = x) \\ &= P(\{\omega : Y(\omega) = y\} | \{\omega : X(\omega) = x\}) \\ &= \frac{P(\{\omega : X(\omega) = x\} \cap \{\omega : Y(\omega) = y\})}{P(\{\omega : X(\omega) = x\})} \\ &= \frac{P_{XY}(X = x, Y = y)}{P_X(X = x)} \\ &= \frac{p_{XY}(x, y)}{p_X(x)} \end{aligned}$$

or simply $p(x, y) = p(y|x) \cdot p(x)$. This is a direct consequence of the product rule from Eq. (1.2). From this expression, we can calculate the probability of an event A as

$$P_{Y|X}(Y \in A|X = x) = \begin{cases} \sum_{y \in A} p_{Y|X}(y|x) & Y : \text{discrete} \\ \int_{y \in A} p_{Y|X}(y|x) dy & Y : \text{continuous} \end{cases}$$

The extension to more than two variables is straightforward. We can write

$$p(x_k|x_1, \dots, x_{k-1}) = \frac{p(x_1, \dots, x_k)}{p(x_1, \dots, x_{k-1})}.$$

By a recursive application of the product rule, we obtain

$$p(x_1, \dots, x_k) = p(x_1) \prod_{l=2}^k p(x_l|x_1, \dots, x_{l-1}) \quad (1.3)$$

which is referred to as the *chain rule*.

1.2.4 Independence of random variables

Two random variables are independent if their joint probability distribution can be expressed as

$$p_{XY}(x, y) = p_X(x) \cdot p_Y(y).$$

As before, k random variables are (collectively, mutually) independent if a joint probability distribution of any subset of variables can be expressed as a product of individual (marginal) probability distributions of its components.

Another, different, form of independence can be found even more frequently in probabilistic calculations. It represents independence between variables in the presence of some other random variable (evidence), e.g.

$$p_{XY|Z}(x, y|z) = p_{X|Z}(x|z) \cdot p_{Y|Z}(y|z)$$

and is referred to as *conditional independence*. Interestingly, the two forms of independence are unrelated; i.e. neither one implies the other. We show this in two simple examples from Figure 1.5.

A. X and Y are independent, but not conditionally independent given Z

$P(X=1)$
a

$$P(Y=y|X=x) = P(Y=y)$$

for example,

$$P(Y=1|X=x) = b$$

$$P(Y=1) = b$$

X	$P(Y=1 X)$
0	b
1	b

X	Y	$P(Z=1 X, Y)$
0	0	c
0	1	$1-c$
1	0	$1-c$
1	1	c

$$P(Y=y|X=x, Z=z) \neq P(Y=y|Z=z)$$

for example,

$$P(Y=1|X=1, Z=1) = bc/(1-c-b(1-2c))$$

$$P(Y=1|Z=1) = b(1-c-a(1-2c))/d$$

$$\text{where } d = P(Z=1)$$

B. X and Z are conditionally independent given Y , but not independent

$P(X=1)$
a

$$P(Z=z|X=x) \neq P(Z=z)$$

for example,

$$P(Z=1|X=1) = d + ce - cd$$

$$P(Z=1) = d + (e-d)(a(c-b) + b)$$

X	$P(Y=1 X)$
0	b
1	c

X	Y	$P(Z=1 X, Y)$
0	0	d
0	1	e
1	0	d
1	1	e

$$P(Z=z|X=x, Y=y) = P(Z=z|Y=y)$$

for example,

$$P(Z=1|X=x, Y=1) = e$$

$$P(Z=1|Y=1) = e$$

Figure 1.5: Independence vs. conditional independence using probability distributions involving three binary random variables. Probability distributions are presented using factorization $p(x, y, z) = p(x)p(y|x)p(z|x, y)$, where all constants $a, b, c, d, e \in [0, 1]$. (A) Variables X and Y are independent, but not conditionally independent given Z . When $c = 0$, $Z = X \oplus Y$, where \oplus is an “exclusive or” operator. (B) Variables X and Z are conditionally independent given Y , but are not independent.

1.2.5 Expectations and moments

Expectations of functions are defined as sums (or integrals) of function values weighted according to the probability distribution function. Given a probability space $(\Omega_X, \mathcal{B}(\Omega_X), P_X)$, we consider a function $f : \Omega_X \rightarrow \mathbb{C}$ and define its expectation function as

$$E[f(X)] = \begin{cases} \sum_{x \in \Omega_X} f(x)p_X(x) & X : \text{discrete} \\ \int_{\Omega_X} f(x)p_X(x)dx & X : \text{continuous} \end{cases}$$

Note that we use a capital X for the random variable (with $f(X)$ the random variable transformed by f) and lower case x when it is an instance (e.g. $p(x)$ is the probability of a specific outcome). It can happen for continuous distributions that $E[f(X)] = \pm\infty$; in such cases we say that the expectation does not exist or is not well-defined.³ For $f(x) = x$, we have a standard expectation $E[X] = \sum xp_X(x)$, or the mean value of X . Using $f(x) = x^k$ results in the k -th moment, $f(x) = \log \frac{1}{p_X(x)}$ gives the well-known entropy function $H(X)$, or differential entropy for continuous random variables, and $f(x) = (x - E[X])^2$ provides the variance of a random variable X , denoted by $V[X]$. Interestingly, the probability of some event $A \subseteq \Omega_X$ can also be expressed in the form of expectation; i.e.

$$P_X(A) = E[1_A(x)],$$

where

$$1_A(x) = \begin{cases} 1 & x \in A \\ 0 & x \notin A \end{cases}$$

is an indicator function. With this, it is possible to express the cdf as $F_X(t) = E[1_{(-\infty, t]}(x)]$.

Function $f(x)$ inside the expectation can also be complex-valued. For example, $\varphi_X(t) = E[e^{itX}]$, where i is the imaginary unit, defines the characteristic function of X . The characteristic function is closely related to the inverse Fourier transform of $p_X(x)$ and is useful in many forms of statistical inference. Several expectation functions are summarized in Table 1.1.

³There is sometimes disagreement on terminology, and some definitions allow the expected value to be infinite, which, for example, still allows the strong law of large numbers. In that setting, an expectation is not well-defined only if both left and right improper integrals are infinite. For our purposes, this is splitting hairs.

$f(x)$	Symbol	Name
x	$E[X]$	Mean
$(x - E[X])^2$	$V[X]$	Variance
x^k	$E[X^k]$	k-th moment; $k \in \mathbb{N}$
$(x - E[X])^k$	$E[(X - E[X])^k]$	k-th central moment; $k \in \mathbb{N}$
e^{tx}	$M_X(t)$	Moment generating function
e^{itx}	$\varphi_X(t)$	Characteristic function
$\log \frac{1}{p_X(x)}$	$H(X)$	(Differential) entropy
$\log \frac{p_X(x)}{q(x)}$	$D(p_X q)$	Kullback-Leibler divergence
$\left(\frac{\partial}{\partial \theta} \log p_X(x \theta)\right)^2$	$\mathcal{I}(\theta)$	Fisher information

Table 1.1: Some important expectation functions $E[f(X)]$ for a random variable X described by its distribution $p_X(x)$. Function $q(x)$ in the definition of the Kullback-Leibler divergence is non-negative and must sum (integrate) to 1; i.e. it is a probability distribution itself. The Fisher information is defined for a family of probability distributions specified by a parameter θ . Note that the moment generating function may not exist for some distributions and all values of t ; however, the characteristic function always exists, even when the density function does not.

$f(x, y)$	Symbol	Name
$(x - E[X])(y - E[Y])$	$\text{cov}(X, Y)$	Covariance
$\frac{(x - E[X])(y - E[Y])}{\sqrt{V[X]V[Y]}}$	$\text{corr}(X, Y)$	Correlation
$\log \frac{p_{XY}(x, y)}{p_X(x)p_Y(y)}$	$I(X; Y)$	Mutual information
$\log \frac{1}{p_{XY}(x, y)}$	$H(X, Y)$	Joint entropy
$\log \frac{1}{p_{X Y}(x y)}$	$H(X Y)$	Conditional entropy

Table 1.2: Some important expectation functions $E[f(X, Y)]$ for two random variables, X and Y , described by their joint distribution $p_{XY}(x, y)$. Mutual information is sometimes referred to as average mutual information.

Given two random variables X and Y and a specific value x assigned to X , we define the conditional expectation as

$$E[f(Y)|x] = \begin{cases} \sum_{y \in \Omega_Y} f(y) p_{Y|X}(y|x) & Y : \text{discrete} \\ \int_{\Omega_Y} f(y) p_{Y|X}(y|x) dy & Y : \text{continuous} \end{cases}$$

where $f : \Omega_Y \rightarrow \mathbb{C}$ is some function. Again, using $f(y) = y$ results in $E[Y|x] = \sum y p_{Y|X}(y|x)$ or $E[Y|x] = \int y p_{Y|X}(y|x) dy$. We shall see later that under some conditions $E[Y|x]$ is referred to as the *regression function*. These types of integrals are often seen and evaluated in Bayesian statistics.

For two random variables X and Y we also define

$$E[f(X, Y)] = \begin{cases} \sum_{x \in \Omega_X} \sum_{y \in \Omega_Y} f(x, y) p_{XY}(x, y) & X, Y : \text{discrete} \\ \int_{\Omega_X} \int_{\Omega_Y} f(x, y) p_{XY}(x, y) dx dy & X, Y : \text{continuous} \end{cases}$$

Expectations can also be defined over a single variable

$$E[f(X, y)] = \begin{cases} \sum_{x \in \Omega_X} f(x, y) p_X(x) & X : \text{discrete} \\ \int_{\Omega_X} f(x, y) p_X(x) dx & X : \text{continuous} \end{cases}$$

where $E[f(X, y)]$ is now a function of y .

We define the covariance function as

$$\begin{aligned} \text{cov}(X, Y) &= E[(x - E[X])(y - E[Y])] \\ &= E[XY] - E[X]E[Y], \end{aligned}$$

with $\text{cov}(X) = \text{cov}(X, X)$ being the variance of the random variable X . Similarly, we define a correlation function as

$$\text{corr}(X, Y) = \frac{\text{cov}(X, Y)}{\sqrt{\text{cov}(X) \cdot \text{cov}(Y)}},$$

which is simply a covariance function normalized by the product of standard deviations. Both covariance and correlation functions have wide applicability in statistics, machine learning, signal processing and many other disciplines. Several important expectations for two random variables are listed in [Table 1.2](#).

Example 6: Three tosses of a fair coin (yet again). Consider two random variables from Examples 3 and 5, and calculate the expectation and variance for both X and Y . Then calculate $E[Y|X = 0]$.

We start by calculating $E[X] = 0 \cdot p_X(0) + 1 \cdot p_X(1) = \frac{1}{2}$. Similarly,

$$\begin{aligned} E[Y] &= \sum_{y=0}^3 y \cdot p_Y(y) \\ &= p_Y(1) + 2p_Y(2) + 3p_Y(3) \\ &= \frac{3}{2} \end{aligned}$$

The conditional expectation can be found as

$$\begin{aligned} E[Y|X=0] &= \sum_{y=0}^3 y \cdot p_{Y|X}(y|0) \\ &= p_{Y|X}(1|0) + 2p_{Y|X}(2|0) + 3p_{Y|X}(3|0) \\ &= 1 \end{aligned}$$

where $p_{Y|X}(y|x) = p_{XY}(x, y)/p_X(x)$. □

In many situations we need to analyze more than two random variables. A simple two-dimensional summary of all pairwise covariance values involving k random variables X_1, X_2, \dots, X_k is called the covariance matrix. More formally, the covariance matrix is defined as

$$\mathbf{\Sigma} = [\Sigma_{ij}]_{i,j=1}^k$$

where

$$\begin{aligned} \Sigma_{ij} &= \text{cov}(X_i, X_j) \\ &= E[(X_i - E[X_i])(X_j - E[X_j])] \end{aligned}$$

with the full matrix written as

$$\mathbf{\Sigma} = E[(\mathbf{X} - E[\mathbf{X}])(\mathbf{X} - E[\mathbf{X}])^\top].$$

Here, the diagonal elements of a $k \times k$ covariance matrix are individual variance values for each variable X_i and the off-diagonal elements are the covariance values between pairs of variables. The covariance matrix is symmetric. It is sometimes called a variance-covariance matrix.

Properties of expectations

Here we review, without proofs, some useful properties of expectations. For any two random variables X and Y , and constant $c \in \mathbb{R}$, it holds that:

1. $E[cX] = cE[X]$
2. $E[X + Y] = E[X] + E[Y]$
3. $V[c] = 0$
4. $V[X] \geq 0$
5. $V[cX] = c^2V[X]$.

In addition, if X and Y are independent random variables, it holds that:

1. $E[XY] = E[X] \cdot E[Y]$
2. $V[X + Y] = V[X] + V[Y]$
3. $\text{cov}(X, Y) = 0$.

1.2.6 Mixtures of distributions

In previous sections we saw that random variables are often described using particular families of probability distributions. This approach can be generalized by considering mixtures of distributions, i.e. linear combinations of other probability distributions. As before, we shall only consider random variables that have their probability mass or density functions.

Given a set of m probability distributions, $\{p_i(x)\}_{i=1}^m$, a finite mixture distribution function, or *mixture model*, $p(x)$ is defined as

$$p(x) = \sum_{i=1}^m w_i p_i(x), \quad (1.4)$$

where $\mathbf{w} = (w_1, w_2, \dots, w_m)$ is a set of non-negative real numbers such that $\sum_{i=1}^m w_i = 1$. We refer to \mathbf{w} as mixing coefficients or, sometimes, as mixing probabilities. A linear combination with such coefficients is called a convex combination. It is straightforward to verify that a function defined in this manner is indeed a probability distribution.

Here we will briefly look into the basic expectation functions of the mixture distribution. Suppose $\{X_i\}_{i=1}^m$ is a set of m random variables described by their respective probability distribution functions $\{p_{X_i}(x)\}_{i=1}^m$. Suppose

also that a random variable X is described by a mixture distribution with coefficients \mathbf{w} and probability distributions $\{p_{X_i}(x)\}_{i=1}^m$. Then, assuming continuous random variables defined on \mathbb{R} , the expectation function is given as

$$\begin{aligned} E[f(X)] &= \int_{-\infty}^{+\infty} f(x)p_X(x)dx \\ &= \int_{-\infty}^{+\infty} f(x) \sum_{i=1}^m w_i p_{X_i}(x)dx \\ &= \sum_{i=1}^m w_i \int_{-\infty}^{+\infty} f(x)p_{X_i}(x)dx \\ &= \sum_{i=1}^m w_i E[f(X_i)]. \end{aligned}$$

We can now apply this formula to obtain the mean, when $f(x) = x$ and the variance, when $f(x) = (x - E[X])^2$, of the random variable X as

$$E[X] = \sum_{i=1}^m w_i E[X_i],$$


and

$$V[X] = \sum_{i=1}^m w_i V[X_i] + \sum_{i=1}^m w_i (E[X_i] - E[X])^2,$$

respectively. A mixture distribution can also be defined for countably and uncountably infinite numbers of components. Such distributions, however, are rare in practice.

Example 7: Signal communications. Consider transmission of a single binary digital signal (bit) over a noisy communication channel shown in Figure 1.6. The magnitude of the signal X emitted by the source is equally likely to be 0 or 1 Volt. The signal is sent over a transmission line (e.g. radio communication, optical fiber, magnetic tape) in which a zero-mean normally distributed noise component Y is added to X . Derive the probability distribution of the signal $Z = X + Y$ that enters the receiver.

We will consider a slightly more general situation where $X : \text{Bernoulli}(\alpha)$ and $Y : \text{Gaussian}(\mu, \sigma^2)$. To find $p_Z(z)$ we will use characteristic functions of random variables X , Y and Z , written as $\varphi_X(t) = E[e^{itX}]$, $\varphi_Y(t) = E[e^{itY}]$



X : Bernoulli(α)

$$Z = X + Y$$

Y : Gaussian(μ, σ^2)

Figure 1.6: A digital signal communication system with additive noise.

and $\varphi_Z(t) = E[e^{itZ}]$. Without derivation we write

$$\begin{aligned}\varphi_X(t) &= 1 - \alpha + \alpha e^{it} \\ \varphi_Y(t) &= e^{it\mu - \frac{\sigma^2 t^2}{2}}\end{aligned}$$

and subsequently

$$\begin{aligned}\varphi_Z(t) &= \varphi_{X+Y}(t) \\ &= \varphi_X(t) \cdot \varphi_Y(t) \\ &= (1 - \alpha + \alpha e^{it}) \cdot e^{it\mu - \frac{\sigma^2 t^2}{2}} \\ &= \alpha e^{it(\mu+1) - \frac{\sigma^2 t^2}{2}} + (1 - \alpha) e^{it\mu - \frac{\sigma^2 t^2}{2}}.\end{aligned}$$

By performing integration on $\varphi_Z(t)$ we can easily verify that

$$p_Z(z) = \alpha \cdot \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(z-\mu-1)^2} + (1 - \alpha) \cdot \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(z-\mu)^2},$$

which is a mixture of two normal distributions $\mathcal{N}(\mu + 1, \sigma^2)$ and $\mathcal{N}(\mu, \sigma^2)$ with coefficients $w_1 = \alpha$ and $w_2 = 1 - \alpha$, respectively. Observe that a convex combination of random variables $Z = w_1 X + w_2 Y$ does not imply $p_Z(z) = w_1 p_X(x) + w_2 p_Y(y)$. □

1.2.7 Graphical representation of probability distributions

We saw earlier that a joint probability distribution can be *factorized* using the chain rule from Eq. (1.3). Such factorizations can be visualized using a

directed graph representation, where nodes represent random variables and edges depict dependence. For example,

$$p(x, y, z) = p(x)p(y|x)p(z|x, y)$$

is shown in Figure 1.7a. Graphical representations of probability distributions using directed acyclic graphs, together with conditional probability distributions, are called *Bayesian networks* or *belief networks*. They facilitate interpretation as well as effective statistical inference.

Visualizing relationships between variables becomes particularly convenient when we want to understand and analyze conditional independence properties of variables. Figure 1.7b shows the the same factorization of $p(x, y, z)$ where variable Z is independent of X given Y . To carefully determine conditional independence and dependence properties, however, one usually uses the *d-separation* rules for belief networks. Though often relationships are intuitive, sometimes dependence properties can get more complicated due to multiple relationships between nodes. For example, in Figure 1.7c, two nodes do not have an edge, but are conditionally dependent through another node. We will not further examine d-separation rules at this time, but can easily be found in any standard textbook on graphical models.

Belief networks have a simple, formal definition. Given a set of k random variables $\mathbf{X} = (X_1, \dots, X_k)$, belief networks factorize the joint probability distribution of \mathbf{X} as

$$p(\mathbf{x}) = \prod_{i=1}^k p(x_i | \mathbf{x}_{\text{Parents}(X_i)}),$$


where $\text{Parents}(X)$ denotes the immediate ancestors of node X in the graph. In Figure 1.7B, node Y is a parent of Z , but node X is not a parent of Z .

It is important to mention that there are multiple (how many?) ways of factorizing a distribution. For example, by reversing the order of variables $p(x, y, z)$ can be also factorized as


$$p(x, y, z) = p(z)p(y|z)p(x|y, z),$$

which has a different graphical representation and its own conditional probability distributions, yet the same joint probability distribution as the earlier factorization. Selecting a proper factorization and estimating the conditional probability distributions from data will be discussed in detail later.

Undirected graphs can also be used to factorize probability distributions. The main idea here is to decompose graphs into maximal cliques \mathcal{C} (the



$$P(X = x, Y = y, Z = z) = P(X = x)P(Y = y|X = x)P(Z = z|X = x, Y = y)$$



$$P(X = x, Y = y, Z = z) = P(X = x)P(Y = y|X = x)P(Z = z|Y = y)$$




Figure 1.7: Bayesian network: graphical representation of three joint probability distributions for three discrete (binary) random variables (X, Y, Z) using directed acyclic graphs. The probability mass function $p(x, y, z)$ is defined over $\{0, 1\}^3$. (a) Full factorization; (b) Factorization that shows and ensures conditional independence between Z and X , given Y . Each node is associated with a conditional probability distribution. In discrete cases, these conditional distributions are referred to as conditional probability tables. (c) An example of a model where the lack of an edge between nodes does not indicate independence. Given information about Z , X and Y are actually dependent (i.e., they are conditionally dependent through Z).




Figure 1.8: Markov network: graphical representation of a probability distribution using maximum clique decomposition. Shown is a set of eight random variables with their interdependency structure and maximum clique decomposition (a clique is fully connected subgraph of a given graph). A decomposition into maximum cliques covers all vertices and edges in a graph with the minimum number of cliques. Here, the set of variables is decomposed into four maximal cliques $\mathcal{C} = \{C_1, C_2, C_3, C_4\}$.

smallest set of cliques that covers the graph) and express the distribution in the following form

$$p(\mathbf{x}) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(\mathbf{x}_C),$$

where each $\psi_C(\mathbf{x}_C) \geq 0$ is called the clique potential function and

$$Z = \int_{\mathbf{x}} \prod_{C \in \mathcal{C}} \psi_C(\mathbf{x}_C) d\mathbf{x},$$

is called the partition function, used strictly for normalization purposes. In contrast to conditional probability distributions in directed acyclic graphs, the clique potentials usually do not have conditional probability interpretations and, thus, normalization is necessary. One example of a maximum clique decomposition is shown in Figure 1.8.

The potential functions are typically taken to be strictly positive, $\psi_C(\mathbf{x}_C) > 0$, and expressed as

$$\psi_C(\mathbf{x}_C) = \exp(-E(\mathbf{x}_C)),$$

where $E(\mathbf{x}_C)$ is a user-specified energy function on the clique of random variables \mathbf{X}_C . This leads to the probability distribution of the following form

$$p(\mathbf{x}) = \frac{1}{Z} \exp \left(\sum_{C \in \mathcal{C}} \log \psi_C(\mathbf{x}_C) \right).$$

As formulated, this probability distribution is called the Boltzmann distribution or the Gibbs distribution.

The energy function $E(\mathbf{x})$ must be lower for values of \mathbf{x} that are more likely. It also may involve parameters that are then estimated from the available training data. Of course, in a prediction problem, an undirected graph must be created to also involve the target variables, which were here considered to be a subset of \mathbf{X} .

Consider now any probability distribution over all possible configurations of the random vector \mathbf{X} with its underlying graphical representation. If the following property

$$p(x_i | \mathbf{x}_{-X_i}) = p(x_i | \mathbf{x}_{N(X_i)}) \quad (1.5)$$

is satisfied, the probability distribution is referred to as *Markov network* or a *Markov random field*. In the equation above

$$\mathbf{X}_{-X_i} = (X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_k)$$

and $N(X)$ is a set of random variables neighboring X in the graph; i.e. there exists an edge between X and every node in $N(X)$. The set of random variables in $N(X)$ is also called the Markov blanket of X .

It can be shown that every Gibbs distribution satisfies the property from Eq. (1.5) and, conversely, that for every probability distribution for which Eq. (1.5) holds can be represented as a Gibbs distribution with some choice of parameters. This equivalence of Gibbs distributions and Markov networks was established by the Hammersley-Clifford theorem.

Chapter 2

Basic Principles of Parameter Estimation

In probabilistic modeling, we are typically presented with a set of observations and the objective is to find a model, or function, \hat{M} that shows good agreement with the data and respects certain additional requirements. We shall roughly categorize these requirements into three groups: (i) the ability to generalize well, (ii) the ability to incorporate prior knowledge and assumptions into modeling, and (iii) scalability. First, the model should be able to stand the test of time; that is, its performance on the previously unseen data should not deteriorate once this new data is presented. Models with such performance are said to generalize well. Second, \hat{M} must be able to incorporate information about the model space \mathcal{M} from which it is selected and the process of selecting a model should be able to accept training “advice” from an analyst. Finally, when large amounts of data are available, learning algorithms must be able to provide solutions in reasonable time given the resources such as memory or CPU power. In summary, the choice of a model ultimately depends on the observations at hand, our experience with modeling real-life phenomena, and the ability of algorithms to find good solutions given limited resources.

An easy way to think about finding the “best” model is through learning parameters of a distribution. Suppose we are given a set of observations $\mathcal{D} = \{x_i\}_{i=1}^n$, where $x_i \in \mathbb{R}$ and have knowledge that \mathcal{M} is a family of all univariate Gaussian distributions, e.g. $\mathcal{M} = \text{Gaussian}(\mu, \sigma^2)$, with $\mu \in \mathbb{R}$ and $\sigma \in \mathbb{R}^+$. In this case, the problem of finding the best model (by which we mean function) can be seen as finding the best parameters μ^* and σ^* , i.e. the problem can be seen as *parameter estimation*. We call this process estimation because the typical assumption is that the data was generated by an unknown model from \mathcal{M} whose parameters we are trying to recover from data.

We will formalize parameter estimation using probabilistic techniques and will subsequently find solutions through optimization, occasionally with constraints in the parameter space. The main assumption throughout this

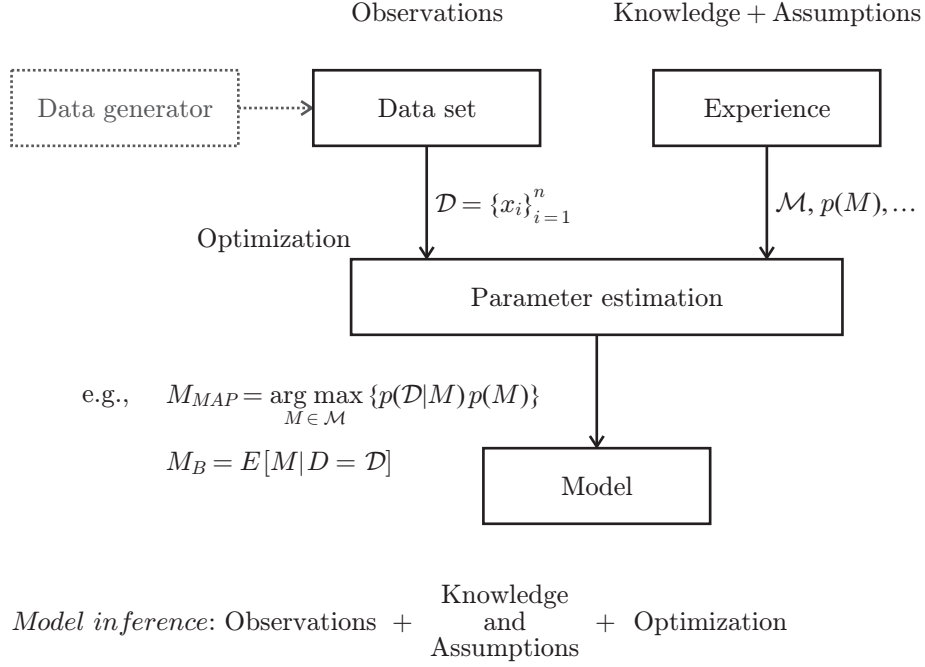


Figure 2.1: Statistical framework for model inference. The estimates of the parameters are made using a set of observations \mathcal{D} as well as experience in the form of model space \mathcal{M} , prior distribution $p(M)$, or specific starting solutions in the optimization step.

part will be that the set of observations \mathcal{D} was generated (or collected) independently and according to the same distribution $p_X(x)$. The statistical framework for model inference is shown in Figure 2.1.

2.1 Maximum a posteriori and maximum likelihood estimation

The idea behind *maximum a posteriori* (MAP) estimation is to find the most probable model for the observed data. Given the data set \mathcal{D} , we formalize the MAP solution as

$$M_{MAP} = \arg \max_{M \in \mathcal{M}} \{p(M|\mathcal{D})\},$$

where $p(M|\mathcal{D})$ is called the *posterior distribution* of the model given the data. In discrete model spaces, $p(M|\mathcal{D})$ is the probability mass function and

the MAP estimate is exactly the most probable model. Its counterpart in continuous spaces is the model with the largest value of the posterior density function. Note that we use words *model*, which is a function, and its *parameters*, which are the coefficients of that function, somewhat interchangeably. However, we should keep in mind the difference, even if only for pedantic reasons.

To calculate the posterior distribution we start by applying the Bayes rule as

$$p(M|\mathcal{D}) = \frac{p(\mathcal{D}|M) \cdot p(M)}{p(\mathcal{D})}, \quad (2.1)$$

where $p(\mathcal{D}|M)$ is called the *likelihood* function, $p(M)$ is the *prior* distribution of the model, and $p(\mathcal{D})$ is the *marginal* distribution of the data. Notice that we use \mathcal{D} for the observed data set, but that we usually think of it as a realization of a multidimensional random variable D drawn according to some distribution $p(\mathcal{D})$. Using the formula of total probability, we can express $p(\mathcal{D})$ as

$$p(\mathcal{D}) = \begin{cases} \sum_{M \in \mathcal{M}} p(\mathcal{D}|M)p(M) & M : \text{discrete} \\ \int_{\mathcal{M}} p(\mathcal{D}|M)p(M)dM & M : \text{continuous} \end{cases}$$

Therefore, the posterior distribution can be fully described using the likelihood and the prior. The field of research and practice involving ways to determine this distribution and optimal models is referred to as *inferential statistics*. The posterior distribution is sometimes referred to as inverse probability.

Finding M_{MAP} can be greatly simplified because $p(\mathcal{D})$ in the denominator does not affect the solution. We shall re-write Eq. (2.1) as

$$\begin{aligned} p(M|\mathcal{D}) &= \frac{p(\mathcal{D}|M) \cdot p(M)}{p(\mathcal{D})} \\ &\propto p(\mathcal{D}|M) \cdot p(M), \end{aligned}$$

where \propto is the proportionality symbol. Thus, we can find the MAP solution by solving the following optimization problem

$$M_{MAP} = \arg \max_{M \in \mathcal{M}} \{p(\mathcal{D}|M)p(M)\}.$$

In some situations we may not have a reason to prefer one model over another and can think of $p(M)$ as a constant over the model space \mathcal{M} . Then, maximum a posteriori estimation reduces to the maximization of the likelihood function; i.e.

$$M_{ML} = \arg \max_{M \in \mathcal{M}} \{p(\mathcal{D}|M)\}.$$

We will refer to this solution as the *maximum likelihood* solution. Formally speaking, the assumption that $p(M)$ is constant is problematic because a uniform distribution cannot be always defined (say, over \mathbb{R}), though there are some solutions to this issue using improper priors. Nonetheless, it may be useful to think of the maximum likelihood approach as a separate technique, rather than a special case of MAP estimation, but keep this connection in mind.

Observe that MAP and ML approaches report solutions corresponding to the mode of the posterior distribution and the likelihood function, respectively. We shall later contrast this estimation technique with the view of the Bayesian statistics in which the goal is to minimize the posterior risk. Such estimation typically results in calculating conditional expectations, which can be complex integration problems. From a different point of view, MAP and ML estimates are called *point estimates*, as opposed to estimated that report confidence intervals for particular group of parameters.

Example 8: Suppose data set $\mathcal{D} = \{2, 5, 9, 5, 4, 8\}$ is an i.i.d. sample from a Poisson distribution with an unknown parameter λ_t . Find the maximum likelihood estimate of λ_t .

The probability density function of a Poisson distribution is expressed as $p(x|\lambda) = \lambda^x e^{-\lambda}/x!$, with some parameter $\lambda \in \mathbb{R}^+$. We will estimate this parameter as

$$\lambda_{ML} = \arg \max_{\lambda \in (0, \infty)} \{p(\mathcal{D}|\lambda)\}. \quad (2.2)$$

We can write the likelihood function as

$$\begin{aligned} p(\mathcal{D}|\lambda) &= p(\{x_i\}_{i=1}^n | \lambda) \\ &= \prod_{i=1}^n p(x_i | \lambda) \\ &= \frac{\lambda^{\sum_{i=1}^n x_i} \cdot e^{-n\lambda}}{\prod_{i=1}^n x_i!}. \end{aligned}$$

To find λ that maximizes the likelihood, we will first take a logarithm (a monotonic function) to simplify the calculation, then find its first derivative with respect to λ , and finally equate it with zero to find the maximum. Specifically, we express the log-likelihood $ll(\mathcal{D}, \lambda) = \ln p(\mathcal{D}|\lambda)$ as

$$ll(\mathcal{D}, \lambda) = \ln \lambda \sum_{i=1}^n x_i - n\lambda - \sum_{i=1}^n \ln(x_i!)$$

and proceed with the first derivative as

$$\begin{aligned} \frac{\partial ll(\mathcal{D}, \lambda)}{\partial \lambda} &= \frac{1}{\lambda} \sum_{i=1}^n x_i - n \\ &= 0. \end{aligned}$$

By substituting $n = 6$ and values from \mathcal{D} , we can compute the solution as

$$\begin{aligned} \lambda_{ML} &= \frac{1}{n} \sum_{i=1}^n x_i \\ &= 5.5, \end{aligned}$$

which is simply a sample mean. The second derivative of the likelihood function is always negative because λ must be positive; thus, the previous expression indeed maximizes the likelihood. Note that to properly maximize this loss, we also need to ensure the constraint $\lambda \in (0, \infty)$ is enforced. Because the solution above is in the constraint set, we know we have the correct solution to Equation (2.2); however, in other situations, we will have to explicitly enforce constraints in the optimization, as we will discuss later. \square

Example 9: Let $\mathcal{D} = \{2, 5, 9, 5, 4, 8\}$ again be an i.i.d. sample from $\text{Poisson}(\lambda_t)$, but now we are also given additional information. Suppose the prior knowledge about λ_t can be expressed using a gamma distribution $\Gamma(x|k, \theta)$ with parameters $k = 3$ and $\theta = 1$. Find the maximum a posteriori estimate of λ_t .

First, we write the probability density function of the gamma family as

$$\Gamma(x|k, \theta) = \frac{x^{k-1} e^{-\frac{x}{\theta}}}{\theta^k \Gamma(k)},$$

where $x > 0$, $k > 0$, and $\theta > 0$. $\Gamma(k)$ is the gamma function that generalizes the factorial function; when k is an integer, we have $\Gamma(k) = (k-1)!$. The

MAP estimate of the parameters can be found as

$$\lambda_{MAP} = \arg \max_{\lambda \in (0, \infty)} \{p(\mathcal{D}|\lambda)p(\lambda)\}.$$

As before, we can write the likelihood function as

$$p(\mathcal{D}|\lambda) = \frac{\lambda^{\sum_{i=1}^n x_i} \cdot e^{-n\lambda}}{\prod_{i=1}^n x_i!}$$

and the prior distribution as

$$p(\lambda) = \frac{\lambda^{k-1} e^{-\frac{\lambda}{\theta}}}{\theta^k \Gamma(k)}.$$

Now, we can maximize the logarithm of the posterior distribution $p(\lambda|\mathcal{D})$ using

$$\begin{aligned} \ln p(\lambda|\mathcal{D}) &\propto \ln p(\mathcal{D}|\lambda) + \ln p(\lambda) \\ &= \ln \lambda(k-1 + \sum_{i=1}^n x_i) - \lambda(n + \frac{1}{\theta}) - \sum_{i=1}^n \ln x_i! - k \ln \theta - \ln \Gamma(k) \end{aligned}$$

to obtain

$$\begin{aligned} \lambda_{MAP} &= \frac{k-1 + \sum_{i=1}^n x_i}{n + \frac{1}{\theta}} \\ &= 5 \end{aligned}$$

after incorporating all data.

A quick look at λ_{MAP} and λ_{ML} suggests that as n grows, both numerators and denominators in the expressions above become increasingly more similar. In fact, it is a well-known result that, in the limit of infinite samples, both the MAP and ML converge to the same model, M , as long as the prior does not have zero probability on M . This result shows that the MAP estimate approaches the ML solution for large data sets. In other words, large data diminishes the importance of prior knowledge. This is an important conclusion because it simplifies mathematical apparatus necessary for practical inference.

To get some intuition for this result, we will show that the MAP and ML estimates converge to the same solution for the above example with a Poisson distribution. Let $s_n = \sum_{i=1}^n x_i$, which is a sample from the random

variable $S_n = \sum_{i=1}^n X_i$. If $\lim_{n \rightarrow \infty} s_n/n^2 = 0$ (i.e., s_n does not grow faster than n^2), then

$$\begin{aligned} |\lambda_{MAP} - \lambda_{ML}| &= \left| \frac{k-1+s_n}{n+1/\theta} - \frac{s_n}{n} \right| \\ &= \left| \frac{k-1}{n+1/\theta} - \frac{s_n}{n(n+1/\theta)} \right| \\ &\leq \frac{|k-1|}{n+1/\theta} + \frac{s_n}{n(n+1/\theta)} \xrightarrow{n \rightarrow \infty} 0 \end{aligned}$$

Note that if $\lim_{n \rightarrow \infty} s_n/n^2 \neq 0$, then both estimators go to ∞ ; however, such a sequence of values has an essentially non-zero probability of occurring. Consistency theorems for ML and MAP estimation state that convergence to the true parameters occurs "almost surely" or "with probability 1" to indicate that these unbounded sequences constitute a set of measure-zero, under certain reasonable conditions (for more, see [11, Theorem 9.13]).

□

Example 10: Let $\mathcal{D} = \{x_i\}_{i=1}^n$ be an i.i.d. sample from a univariate Gaussian distribution. Compare maximum likelihood (ML) and maximum a posteriori (MAP) solutions for inferring the parameters of a Gaussian distribution. For the MAP solution, assume the following prior distribution on the parameters. Use multidimensional Gaussian.

We start by forming the log-likelihood function

$$\begin{aligned} \ln p(D|\mu, \sigma) &= \ln \prod_{i=1}^n p(x_i|\mu, \sigma) \\ &= n \ln \frac{1}{\sqrt{2\pi}} + n \ln \frac{1}{\sigma} - \frac{\sum_{i=1}^n (x_i - \mu)^2}{2\sigma^2}. \end{aligned}$$

We now compute the partial derivatives of the log-likelihood with respect to all parameters as

$$\frac{\partial}{\partial \mu} \ln p(D|\mu, \sigma) = \frac{\sum_{i=1}^n (x_i - \mu)}{\sigma^2}$$

and

$$\frac{\partial}{\partial \sigma} \ln p(D|\mu, \sigma) = -\frac{n}{\sigma} + \frac{\sum_{i=1}^n (x_i - \mu)^2}{\sigma^3}.$$

From here, we can proceed to derive that

$$\mu_{ML} = \frac{1}{n} \sum_{i=1}^n x_i$$

and

$$\sigma_{ML}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu_{ML})^2.$$

□

The relationship with Kullback-Leibler divergence

We now investigate the relationship between maximum likelihood estimation and Kullback-Leibler divergence. Kullback-Leibler divergence between two probability distributions $p(x)$ and $q(x)$ is defined as

$$D_{KL}(p||q) = \int_{-\infty}^{\infty} p(x) \log \frac{p(x)}{q(x)} dx.$$

In information theory, Kullback-Leibler divergence has a natural interpretation of the inefficiency of signal compression when the code is constructed using a suboptimal distribution $q(x)$ instead of the correct (but unknown) distribution $p(x)$ according to which the data has been generated. However, more often than not, Kullback-Leibler divergence is simply considered to be a measure of divergence between two probability distributions. Although this divergence is not a metric (it is not symmetric and does not satisfy the triangle inequality) it has important theoretical properties in that (i) it is always non-negative and (ii) it is equal to zero if and only if $p(x) = q(x)$.

Consider now a divergence between an estimated probability distribution $p(x|\theta)$ and an underlying (true) distribution $p(x|\theta_t)$ according to which the data set $\mathcal{D} = \{x_i\}_{i=1}^n$ was generated. The Kullback-Leibler divergence between $p(x|\theta)$ and $p(x|\theta_t)$ is

$$\begin{aligned} D_{KL}(p(x|\theta_t)||p(x|\theta)) &= \int_{-\infty}^{\infty} p(x|\theta_t) \log \frac{p(x|\theta_t)}{p(x|\theta)} dx \\ &= \int_{-\infty}^{\infty} p(x|\theta_t) \log \frac{1}{p(x|\theta)} dx - \int_{-\infty}^{\infty} p(x|\theta_t) \log \frac{1}{p(x|\theta_t)} dx. \end{aligned}$$

The second term in the above equation is simply the entropy of the true distribution and is not influenced by our choice of the model θ . The first term, on the other hand, can be expressed as

$$\int_{-\infty}^{\infty} p(x|\theta_t) \log \frac{1}{p(x|\theta)} dx = -E[\log p(x|\theta)]$$

Therefore, maximizing $E[\log p(x|\theta)]$ minimizes the Kullback-Leibler divergence between $p(x|\theta)$ and $p(x|\theta_t)$. Using the strong law of large numbers, we know that

$$\frac{1}{n} \sum_{i=1}^n \log p(x_i|\theta) \xrightarrow{a.s.} E[\log p(x|\theta)]$$

when $n \rightarrow \infty$. Thus, when the data set is sufficiently large, maximizing the likelihood function minimizes the Kullback-Leibler divergence and leads to the conclusion that $p(x|\theta_{\text{ML}}) = p(x|\theta_t)$, if the underlying assumptions are satisfied. Under reasonable conditions, we can infer from it that $\theta_{\text{ML}} = \theta_t$. This will hold for families of distributions for which a set of parameters uniquely determines the probability distribution; e.g. it will not generally hold for mixtures of distributions but we will discuss this situation later. This result is only one of the many connections between statistics and information theory.

Chapter 3

Introduction to Prediction Problems

3.1 Problem Statement

We start by defining a data set $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$, where $\mathbf{x}_i \in \mathcal{X}$ is the i -th object and $y_i \in \mathcal{Y}$ is the corresponding target designation. We usually assume that $\mathcal{X} = \mathbb{R}^d$, in which case $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ik})$ is a k -dimensional vector called a *data point* (or *example* or *sample*). Each dimension of \mathbf{x}_i is typically called a *feature* or an *attribute*.

We generally distinguish between two related but different types of prediction problems: classification and regression. Generally speaking, we have a classification problem if \mathcal{Y} is discrete and a regression problem when \mathcal{Y} is continuous. The *classification problem* refers to constructing a function that for a previously unseen data point \mathbf{x} infers (predicts) its class label y . A particular function or algorithm that infers class labels, and is typically optimized to minimize or maximize some objective (or fitness) function, is referred to as a *classifier* or a *classification model*. The cardinality of \mathcal{Y} in classification problems is usually small, e.g. $\mathcal{Y} = \{\text{healthy}, \text{disease}\}$. On the other hand, the *regression problem* refers to constructing a model that for a previously unseen data point \mathbf{x} approximates the target value y as closely as possible, where often times $\mathcal{Y} = \mathbb{R}$.

An example of a data set for classification with $n = 3$ data points and $k = 5$ features is shown in Table 3.1. Problems in which $|\mathcal{Y}| = 2$ are referred to as binary classification problems, whereas problems in which $|\mathcal{Y}| > 2$ are referred to as multi-class classification problems. This can be more complex as, for instance, in classification of text documents into categories such as $\{\text{sports}, \text{medicine}, \text{travel}, \dots\}$. Here, a single document may be related to more than one value in the set; e.g. an article on sports medicine. To account for this, we can certainly say that $\mathcal{Y} = \mathcal{P}(\{\text{sports}, \text{medicine}, \text{travel}, \dots\})$ and treat the problem as multi-class classification, albeit with a very large output space. However, it is often easier to think that more than one value of the output space can be associated with any particular input. We refer to this learn-

	wt [kg]	ht [m]	T [°C]	sbp [mmHg]	dbp [mmHg]	y
\mathbf{x}_1	91	1.85	36.6	121	75	-1
\mathbf{x}_2	75	1.80	37.4	128	85	+1
\mathbf{x}_3	54	1.56	36.6	110	62	-1

Table 3.1: An example of a binary classification problem: prediction of a disease state for a patient. Here, features indicate weight (wt), height (ht), temperature (T), systolic blood pressure (sbp), and diastolic blood pressure (dbp). The class labels indicate presence of a particular disease, e.g. diabetes. This data set contains one positive data point (\mathbf{x}_2) and two negative data points ($\mathbf{x}_1, \mathbf{x}_3$). The class label shows a disease state, i.e. $y_i = +1$ indicates the presence while $y_i = -1$ indicates absence of disease.

ing task as multi-label classification and set $\mathcal{Y} = \{\text{sports, medicine, travel, } \dots\}$. Finally, \mathcal{Y} can be a set of structured outputs, e.g. strings, trees, or graphs. This classification scenario is usually referred to as structured-output learning. The cardinality of the output space in structured-output learning problems is often very high. An example of a regression problem is shown in Table 3.2.

In both prediction scenarios, we assume that the features are easy to collect for each object (e.g. by measuring the height of a person or the square footage of a house), while the target variable is difficult to observe or expensive to collect. Such situations usually benefit from the construction of a computational model that predicts targets from a set of input values. The model is trained using a set of input objects for which target values have already been collected.

Observe that there does not exist a strict distinction between classification and regression. For example, if the output space is $\mathcal{Y} = \{0, 1, 2\}$, we need not treat this problem as classification. This is because there exists a relationship of order among elements of \mathcal{Y} that can simplify model development. For example, we can take $\mathcal{Y} = [0, 2]$ and simply develop a regression model from which we can recover the original discrete values by rounding the raw prediction outputs. The selection of a particular way of modeling depends on the analyst and their knowledge of the domain as well as technical aspects of learning.

	size [sqft]	age [yr]	dist [mi]	inc [\$]	dens [ppl/mi ²]	y
\mathbf{x}_1	1250	5	2.85	56,650	12.5	2.35
\mathbf{x}_2	3200	9	8.21	245,800	3.1	3.95
\mathbf{x}_3	825	12	0.34	61,050	112.5	5.10

Table 3.2: An example of a regression problem: prediction of the price of a house in a particular region. Here, features indicate the size of the house (size) in square feet, the age of the house (age) in years, the distance from the city center (dist) in miles, the average income in a one square mile radius (inc), and the population density in the same area (dens). The target indicates the price a house is sold at, e.g. in hundreds of thousands of dollars.

3.2 Useful notation

In the machine learning literature we use k -tuples $\mathbf{x} = (x_1, x_2, \dots, x_d)$ to denote data points. However, often times we can benefit from the algebraic notation, where each data point \mathbf{x} is a column vector in the d -dimensional Euclidean space: $\mathbf{x} = [x_1 \ x_2 \ \dots \ x_d]^\top \in \mathbb{R}^d$, where \top is the transpose operator. Here, a linear combination of features and some set of coefficients $\mathbf{w} = [w_1 \ w_2 \ \dots \ w_d]^\top \in \mathbb{R}^d$

$$\sum_{i=1}^d w_i x_i = w_1 x_1 + w_2 x_2 + \dots + w_d x_d$$

can be expressed using an inner (dot) product of column vectors $\mathbf{w}^\top \mathbf{x}$. A linear combination $\mathbf{w}^\top \mathbf{x}$ results in a single number. Another useful notation for such linear combinations will be $\langle \mathbf{w}, \mathbf{x} \rangle$.

We will also use an n -by- d matrix $\mathbf{x} = (\mathbf{x}_1^\top, \mathbf{x}_2^\top, \dots, \mathbf{x}_n^\top)$ to represent the entire set of data points and \mathbf{y} to represent a column vector of targets. For example, the i -th row of \mathbf{x} represents data point \mathbf{x}_i^\top . Finally, the j -th column of \mathbf{x} , denoted as \mathbf{f}_j , is an n -by-1 vector which contains the values of feature j for all data points. The notation is further presented in Figure 3.1.

3.3 Optimal classification and regression models

Our goal now is to establish the performance criteria that will be used to evaluate predictors $f : \mathcal{X} \rightarrow \mathcal{Y}$ and subsequently define optimal classification and regression models. We start with classification and consider a situation




Figure 3.1: Data set representation and notation. \mathbf{x} is an n -by- d matrix representing features and data points, whereas y is an n -by-1 vector of targets.

where the joint probability distribution $p(\mathbf{x}, y)$ is either known or can be learned from data. Also, we will consider that we are given a cost function $c : \mathcal{Y} \times \mathcal{Y} \rightarrow [0, \infty)$, where for each prediction \hat{y} and true target value y the classification cost can be expressed as a constant $c(\hat{y}, y)$, regardless of the input $\mathbf{x} \in \mathcal{X}$ given to the classifier. This cost function can simply be stored as a $|\mathcal{Y}| \times |\mathcal{Y}|$ cost matrix.

The criterion for optimality of a classifier will be probabilistic. In particular, we are interested in minimizing the expected cost

$$\begin{aligned} E[C] &= \int_{\mathcal{X}} \sum_y c(\hat{y}, y) p(\mathbf{x}, y) d\mathbf{x} \\ &= \int_{\mathcal{X}} p(\mathbf{x}) \sum_y c(\hat{y}, y) p(y|\mathbf{x}) d\mathbf{x}, \end{aligned}$$

where the integration is over the entire input space $\mathcal{X} = \mathbb{R}^d$. From this equation, we can see that the optimal classifier can be expressed as

$$f_{\text{BR}}(\mathbf{x}) = \arg \min_{\hat{y} \in \mathcal{Y}} \left\{ \sum_y c(\hat{y}, y) p(y|\mathbf{x}) \right\},$$

for any $\mathbf{x} \in \mathcal{X}$. We will refer to this classifier as the *Bayes risk classifier*. One example where the Bayes risk classifier can be useful is the medical domain. Suppose our goal is to decide whether a patient with a particular set of symptoms (\mathbf{x}) should be sent for an additional lab test ($y = 1$ if yes and

$y = -1$ if not), with cost c_{lab} , in order to improve diagnosis. However, if we do not perform a lab test and the patient is later found to have needed the test for proper treatment, we may incur a significant penalty, say c_{lawsuit} . If $c_{\text{lawsuit}} \gg c_{\text{lab}}$, as it is expected to be, then the classifier needs to appropriately adjust its outputs to account for the cost disparity in different forms of incorrect prediction.

In many practical situations, however, it may not be possible to define a meaningful cost matrix and, thus, a reasonable criterion would be to minimize the probability of a classifier's error $P(f(\mathbf{x}) \neq y)$. This corresponds to the situation where the cost function is defined as

$$c(\hat{y}, y) = \begin{cases} 0 & \text{when } y = \hat{y} \\ 1 & \text{when } y \neq \hat{y} \end{cases}$$

After plugging these values in the definition for $f_{\text{BR}}(\mathbf{x})$, the Bayes risk classifier simply becomes the maximum a posteriori (MAP) classifier. That is,

$$f_{\text{MAP}}(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} \{p(y|\mathbf{x})\}.$$

Therefore, if $p(y|\mathbf{x})$ is known or can be accurately learned, we are fully equipped to make the prediction that minimizes the total cost. In other words, we have converted the problem of minimizing the expected classification cost or probability of error, into the problem of learning functions, more specifically learning probability distributions.

The analysis for regression is a natural extension of that for classification. Here too, we are interested in minimizing the expected cost of prediction of the true target y when a predictor $f(\mathbf{x})$ is used. The expected cost can be expressed as

$$E[C] = \int_{\mathcal{X}} \int_{\mathcal{Y}} c(f(\mathbf{x}), y) p(\mathbf{x}, y) dy d\mathbf{x},$$

where $c : \mathbb{R} \times \mathbb{R} \rightarrow [0, \infty)$ is again some cost function between the predicted value $f(\mathbf{x})$ and the true value y . For simplicity, we will consider

$$c(f(\mathbf{x}), y) = (f(\mathbf{x}) - y)^2,$$

which results in

$$\begin{aligned} E[C] &= \int_{\mathcal{X}} \int_{\mathcal{Y}} (f(\mathbf{x}) - y)^2 p(\mathbf{x}, y) dy d\mathbf{x} \\ &= \int_{\mathcal{X}} p(\mathbf{x}) \underbrace{\int_{\mathcal{Y}} (f(\mathbf{x}) - y)^2 p(y|\mathbf{x}) dy}_{g(f(\mathbf{x}))} d\mathbf{x}. \end{aligned}$$

Assuming $f(\mathbf{x})$ is flexible enough to be separately optimized for each unit volume $d\mathbf{x}$, we see that minimizing $E[C]$ leads us to the problem of minimizing

$$g(u) = \int_{\mathcal{Y}} (u - y)^2 p(y|\mathbf{x}) dy,$$

where we used a substitution $u = f(\mathbf{x})$. We can now differentiate g with respect to u as

$$\begin{aligned} \frac{\partial g(u)}{\partial u} &= 2 \int_{\mathcal{Y}} (u - y) p(y|\mathbf{x}) dy = 0 \\ \implies u \underbrace{\int_{\mathcal{Y}} p(y|\mathbf{x}) dy}_{=1} &= \int_{\mathcal{Y}} yp(y|\mathbf{x}) dy \end{aligned}$$

which results in the optimal solution

$$\begin{aligned} f^*(\mathbf{x}) &= \int_{\mathcal{Y}} yp(y|\mathbf{x}) dy \\ &= E[y|\mathbf{x}]. \end{aligned}$$

Therefore, the optimal regression model in the sense of minimizing the square error between the prediction and the true target is the conditional expectation $E[y|\mathbf{x}]$. It may appear that in the above equations, setting $f(\mathbf{x}) = y$ would always lead to $E[C] = 0$. Unfortunately, this would be an invalid operation because for a single input \mathbf{x} there may be multiple possible outputs y and they can certainly appear in the same data set. To be a well-defined function, $f(\mathbf{x})$ must always have the same output for the same input. $E[C] = 0$ can only be achieved if $p(y|\mathbf{x})$ is a delta function for every \mathbf{x} .

Having found the optimal regression model, we can now write the expected cost in the cases of both optimal and suboptimal models $f(\mathbf{x})$. That is, we are interested in expressing $E[C]$ when

1. $f(\mathbf{x}) = E[y|\mathbf{x}]$
2. $f(\mathbf{x}) \neq E[y|\mathbf{x}]$.

We have already found $E[C]$ when $f(\mathbf{x}) = E[y|\mathbf{x}]$. The expected cost can be simply expressed as

$$E[C] = \int_{\mathcal{X}} \int_{\mathcal{Y}} (E[y|\mathbf{x}] - y)^2 p(\mathbf{x}, y) dy d\mathbf{x},$$

which corresponds to the (weighted) squared error between the optimal prediction and the target everywhere in the feature space. This is the best scenario in regression; we cannot achieve a lower cost on average for this squared cost.

The next situation is when $f(\mathbf{x}) \neq E[y|\mathbf{x}]$. Here, we will proceed by decomposing the squared error as

$$\begin{aligned} (f(\mathbf{x}) - y)^2 &= (f(\mathbf{x}) - E[y|\mathbf{x}] + E[y|\mathbf{x}] - y)^2 \\ &= (f(\mathbf{x}) - E[y|\mathbf{x}])^2 + \underbrace{2(f(\mathbf{x}) - E[y|\mathbf{x}])(E[y|\mathbf{x}] - y)}_{g(\mathbf{x}, y)} + (E[y|\mathbf{x}] - y)^2 \end{aligned}$$

We will now take a more detailed look at $g(\mathbf{x}, y)$ when placed under the expectation over $p(\mathbf{x}, y)$

$$\begin{aligned} E_{\mathbf{x}, y}[g(\mathbf{x}, y)] &= \int_{\mathcal{X}} \int_{\mathcal{Y}} (f(\mathbf{x}) - E[y|\mathbf{x}])(E[y|\mathbf{x}] - y) p(\mathbf{x}, y) dy d\mathbf{x} \\ &= \int_{\mathcal{X}} \int_{\mathcal{Y}} f(\mathbf{x}) E[y|\mathbf{x}] p(\mathbf{x}, y) dy d\mathbf{x} - \int_{\mathcal{X}} \int_{\mathcal{Y}} f(\mathbf{x}) y p(\mathbf{x}, y) dy d\mathbf{x} \\ &\quad + \int_{\mathcal{X}} \int_{\mathcal{Y}} E[y|\mathbf{x}] y p(\mathbf{x}, y) dy d\mathbf{x} - \int_{\mathcal{X}} \int_{\mathcal{Y}} E[y|\mathbf{x}] E[y|\mathbf{x}] p(\mathbf{x}, y) dy d\mathbf{x} \\ &= \int_{\mathcal{X}} f(\mathbf{x}) E[y|\mathbf{x}] p(\mathbf{x}) d\mathbf{x} - \int_{\mathcal{X}} f(\mathbf{x}) p(\mathbf{x}) \int_{\mathcal{Y}} y p(y|\mathbf{x}) dy d\mathbf{x} \\ &\quad + \int_{\mathcal{X}} E[y|\mathbf{x}] E[y|\mathbf{x}] p(\mathbf{x}) d\mathbf{x} - \int_{\mathcal{X}} E[y|\mathbf{x}] E[y|\mathbf{x}] p(\mathbf{x}) d\mathbf{x} \\ &= 0. \end{aligned}$$

where the integration over y simplifies to 1 because $E[y|\mathbf{x}]$ is a constant. Therefore, we can now express the expected cost as

$$\begin{aligned} E[C] &= \int_{\mathcal{X}} \int_{\mathcal{Y}} (f(\mathbf{x}) - y)^2 p(\mathbf{x}, y) dy d\mathbf{x} \\ &= \int_{\mathcal{X}} (f(\mathbf{x}) - E[y|\mathbf{x}])^2 p(\mathbf{x}) d\mathbf{x} + \int_{\mathcal{X}} \int_{\mathcal{Y}} (E[y|\mathbf{x}] - y)^2 p(\mathbf{x}, y) dy d\mathbf{x}, \end{aligned}$$

where the first term is the “distance” between the trained model $f(\mathbf{x})$ and the optimal model $E[y|\mathbf{x}]$ and the second term is the “distance” between the optimal model $E[y|\mathbf{x}]$ and the correct target value y .

To sum up, we argued here that optimal classification and regression models critically depend on knowing or accurately learning the posterior distribution $p(y|\mathbf{x})$. This task can be solved in different ways, but a straightforward approach is to assume a functional form for $p(y|\mathbf{x})$, say $p(y|\mathbf{x}, \boldsymbol{\theta})$, where $\boldsymbol{\theta}$ is a set of weights or parameters that are to be learned from the data. Alternatively, we can learn the class-conditional and prior distributions, $p(\mathbf{x}|y)$ and $p(y)$, respectively. Using

$$\begin{aligned} p(y|\mathbf{x}) &= \frac{p(\mathbf{x}, y)}{p(\mathbf{x})} \\ &= \frac{p(\mathbf{x}|y)p(y)}{\sum_y p(\mathbf{x}, y)} \\ &= \frac{p(\mathbf{x}|y)p(y)}{\sum_y p(\mathbf{x}|y)p(y)} \end{aligned}$$

we can see that these two learning approaches are equivalent in theory. The choice depends on our prior knowledge and/or preferences.

Models obtained by directly estimating $p(y|\mathbf{x})$ are called *discriminative models* and models obtained by directly estimating $p(\mathbf{x}|y)$ and $p(y)$ are called *generative models*. Direct estimation of the joint distribution $p(\mathbf{x}, y)$ is relatively rare as it may be more difficult to hypothesize its parametric or non-parametric form from which the parameters are to be found. Finally, in some situations we can train a model without an explicit probabilistic approach in mind. In these cases, we typically aim to show a good performance of an algorithm in practice, say on a large number of data sets, according to a performance measure relevant to the problems at hand.

3.4 Bayes Optimal Models

We saw earlier that optimal prediction models reduce to the learning of the posterior distribution $p(y|\mathbf{x})$ which is then used to minimize the expected cost (risk, loss). However, in practice, the probability distribution $p(y|\mathbf{x})$ must be modeled using a particular functional form and a set of tunable coefficients. When $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{Y} = \{0, 1\}$, one such example is used in

logistic regression, where

$$p(1|\mathbf{x}) = \frac{1}{1 + e^{-(w_0 + \sum_{j=1}^d w_j x_j)}}$$

and $p(0|\mathbf{x}) = 1 - p(1|\mathbf{x})$. Here $(w_0, w_1, \dots, w_d) \in \mathbb{R}^{d+1}$ is a set of weights that are to be inferred from a given data set \mathcal{D} and $\mathbf{x} \in \mathbb{R}^d$ is an input data point. A number of other types of functional relationships can be used as well, providing a vast set of possibilities for modeling distributions. We will adjust our notation to more precisely denote the posterior distribution as

$$p(y|\mathbf{x}) = p(y|\mathbf{x}, f),$$

where f is a particular function from some function (hypothesis) space \mathcal{F} . We can think of \mathcal{F} as a set of all functions from a specified class, say for all $(w_0, w_1, \dots, w_d) \in \mathbb{R}^{k+1}$ in the example above, but we can also extend the functional class beyond simple parameter variation to incorporate non-linear decision surfaces.

In a typical learning problem, we are given a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ and are asked to model $p(y|\mathbf{x})$. For this purpose, we will think of \mathcal{D} as a realization of a random variable D and will assume that \mathcal{D} was drawn according to the true underlying distribution $p(\mathbf{x}, y)$. Thus, our task is to express $p(y|\mathbf{x}, \mathcal{D})$. Using the sum and product rules, will rewrite our original task of estimating $p(y|\mathbf{x})$ as

$$\begin{aligned} p(y|\mathbf{x}, \mathcal{D}) &= \int_{\mathcal{F}} p(y|\mathbf{x}, f, \mathcal{D}) p(f|\mathbf{x}, \mathcal{D}) df \\ &= \int_{\mathcal{F}} p(y|\mathbf{x}, f) p(f|\mathbf{x}, \mathcal{D}) df. \end{aligned}$$

Here we used conditional independence between output Y and the data set D once a particular model f was selected based on \mathcal{D} ; thus $p(y|\mathbf{x}, f, \mathcal{D}) = p(y|\mathbf{x}, f)$. This equation, gives us a sense that the optimal decision can be made through a mixture of distributions $p(y|\mathbf{x}, f)$, where the weights are given as posterior densities $p(f|\mathbf{x}, \mathcal{D})$. In finite hypothesis spaces \mathcal{F} we have that

$$p(y|\mathbf{x}, \mathcal{D}) = \sum_{f \in \mathcal{F}} p(y|\mathbf{x}, f) p(f|\mathbf{x}, \mathcal{D}),$$

and $p(f|\mathbf{x}, \mathcal{D})$ are posterior probabilities. We may further assume that $p(f|\mathbf{x}, \mathcal{D}) = p(f|\mathcal{D})$, in which case the weights can be precomputed based

on the given data set \mathcal{D} . This leads to more efficient calculations of the posterior probabilities.

In classification, we can rewrite our original MAP classifier as

$$f_{\text{MAP}}(\mathbf{x}, \mathcal{D}) = \arg \max_{y \in \mathcal{Y}} \{p(y|\mathbf{x}, \mathcal{D})\},$$

which readily leads to the following formulation

$$\begin{aligned} f_{\text{MAP}}(\mathbf{x}, \mathcal{D}) &= \arg \max_{y \in \mathcal{Y}} \left\{ \int_{\mathcal{F}} p(y|\mathbf{x}, f, \mathcal{D}) p(f|\mathcal{D}) df \right\} \\ &= \arg \max_{y \in \mathcal{Y}} \left\{ \int_{\mathcal{F}} p(y|\mathbf{x}, f) p(\mathcal{D}|f) p(f) df \right\}. \end{aligned}$$

It can be shown that no classifier can outperform the *Bayes optimal classifier*. Interestingly, the Bayes optimal model also hints that a better prediction performance can be achieved by combining multiple models and averaging their outputs. This provides theoretical support for ensemble learning and methods such as bagging and boosting.

One problem in Bayes optimal classification is efficient calculation of $f_{\text{MAP}}(\mathbf{x}, \mathcal{D})$, given that the function (hypothesis) space \mathcal{F} is generally uncountable. One approach to this is sampling of functions from \mathcal{F} according to $p(f)$ and then calculating $p(f|\mathcal{D})$ or $p(\mathcal{D}|f)$. This can be computed until $p(y|\mathbf{x}, \mathcal{D})$ converges.

Chapter 4

Linear Regression and Its Applications

Given a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ the objective is to learn the relationship between features and the target. We usually start by hypothesizing the functional form of this relationship. For example,

$$f(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2$$

where $\mathbf{w} = (w_0, w_1, w_2)$ is a set of parameters that need to be determined (learned) and $\mathbf{x} = (x_1, x_2)$. Alternatively, we may hypothesize that $f(\mathbf{x}) = \alpha + \beta x_1x_2$, where $\boldsymbol{\theta} = (\alpha, \beta)$ is another set of parameters to be learned. In the former case, the target function is modeled as a *linear combination* of features and parameters, i.e.

$$f(\mathbf{x}) = \sum_{j=0}^k w_j x_j,$$

where we extended \mathbf{x} to $(x_0 = 1, x_1, x_2, \dots, x_k)$. Finding the best parameters \mathbf{w} is then referred to as *linear regression problem*, whereas all other types of relationship between the features and the target fall into a category of *non-linear regression*. In either situation, the regression problem can be presented as a probabilistic modeling approach that reduces to parameter estimation; i.e. to an optimization problem with the goal of maximizing or minimizing some performance criterion between target values $\{y_i\}_{i=1}^n$ and predictions $\{f(\mathbf{x}_i)\}_{i=1}^n$. We can think of a particular optimization algorithm as the *learning* or *training algorithm*.

4.1 Maximum likelihood formulation

We now consider a statistical formulation of linear regression. We shall first lay out the assumptions behind this process and subsequently formulate the problem through maximization of the conditional likelihood function. In

following section, we will show how to solve the optimization and analyze the solution and its basic statistical properties.

Let us assume that the observed data set \mathcal{D} is a product of a data generating process in which n data points were drawn independently and according to the same distribution $p(\mathbf{x})$. Assume also that the target variable Y has an underlying linear relationship with features (X_1, X_2, \dots, X_k) , modified by some error term ε that follows a zero-mean Gaussian distribution, i.e. $\varepsilon : N(0, \sigma^2)$. That is, for a given input \mathbf{x} , the target y is a realization of a random variable Y defined as

$$Y = \sum_{j=0}^k \omega_j X_j + \varepsilon,$$

where $\boldsymbol{\omega} = (\omega_0, \omega_1, \dots, \omega_k)$ is a set of unknown coefficients we seek to recover through estimation. Generally, the assumption of normality for the error term is reasonable (recall the central limit theorem!), although the independence between ε and \mathbf{x} may not hold in practice. Using a few simple properties of expectations, we can see that Y also follows a Gaussian distribution, i.e. its conditional density is $p(y|\mathbf{x}, \boldsymbol{\omega}) = N(\boldsymbol{\omega}^\top \mathbf{x}, \sigma^2)$.

In linear regression, we seek to approximate the target as $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$, where weights \mathbf{w} are to be determined. We first write the conditional likelihood function for a single pair (\mathbf{x}, y) as

$$p(y|\mathbf{x}, \mathbf{w}) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y - \sum_{j=0}^k w_j x_j)^2}{2\sigma^2}}$$

and observe that the only change from the conditional density function of Y is that coefficients \mathbf{w} are used instead of $\boldsymbol{\omega}$. Incorporating the entire data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, we can now write the conditional likelihood function as $p(\mathbf{y}|\mathbf{X}, \mathbf{w})$ and find weights as

$$\mathbf{w}_{ML}^* = \arg \max_{\mathbf{w}} \{p(\mathbf{y}|\mathbf{X}, \mathbf{w})\}.$$

Since examples \mathbf{x} are independent and identically distributed (i.i.d.), we have

$$\begin{aligned} p(\mathbf{y}|\mathbf{X}, \mathbf{w}) &= \prod_{i=1}^n p(y_i|\mathbf{x}_i, \mathbf{w}) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y_i - \sum_{j=0}^k w_j x_{ij})^2}{2\sigma^2}}. \end{aligned}$$




Figure 4.1: An example of a linear regression fitting on data set $\mathcal{D} = \{(1, 1.2), (2, 2.3), (3, 2.3), (4, 3.3)\}$. The task of the optimization process is to find the best linear function $f(x) = w_0 + w_1x$ so that the sum of squared errors $e_1^2 + e_2^2 + e_3^2 + e_4^2$ is minimized.

For the reasons of mathematical convenience, we will look at the logarithm (monotonic function) of the likelihood function and express the log-likelihood as

$$\ln(p(\mathbf{y}|\mathbf{X}, \mathbf{w})) = -\sum_{i=1}^n \log\left(\sqrt{2\pi\sigma^2}\right) - \frac{1}{2\sigma^2} \sum_{i=1}^n \left(y_i - \sum_{j=0}^k w_j x_{ij}\right)^2.$$

Given that the first term on the right-hand side is independent of \mathbf{w} , maximizing the likelihood function corresponds exactly to minimizing the sum of squared errors

$$\begin{aligned} E(\mathbf{w}) &= \sum_{i=1}^n (f(\mathbf{x}_i) - y_i)^2 && \triangleright f(\mathbf{x}_i) = \sum_{j=0}^k w_j x_{ij} \\ &= \sum_{i=1}^n e_i^2. \end{aligned}$$

Geometrically, this error is the square of the Euclidean distance between the vector of predictions $\hat{\mathbf{y}} = (f(\mathbf{x}_1), f(\mathbf{x}_2), \dots, f(\mathbf{x}_n))$ and the vector of observed target values $\mathbf{y} = (y_1, y_2, \dots, y_n)$. A simple example illustrating the linear regression problem is shown in Figure 4.1.

To more explicitly see why the maximum likelihood solution corresponds to minimizing $E(\mathbf{w})$, notice that maximizing the likelihood is equivalent to

maximizing the log-likelihood (because log is monotonic) which is equivalent to minimizing the negative log-likelihood. Therefore, the maximum likelihood \mathbf{w}^* corresponds to

$$\begin{aligned}
\mathbf{w}^* &= \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} -\ln(p(\mathbf{y}|\mathbf{X}, \mathbf{w})) \\
&= \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \sum_{i=1}^n \log \left(\sqrt{2\pi\sigma^2} \right) + \frac{1}{2\sigma^2} \sum_{i=1}^n \left(y_i - \sum_{j=0}^k w_j x_{ij} \right)^2 \\
&= \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \sum_{i=1}^n \left(y_i - \sum_{j=0}^k w_j x_{ij} \right)^2 \\
&= \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} E(\mathbf{w})
\end{aligned}$$

In the next sections, we will discuss how to solve this optimization and the properties of the solution.

Note that we could have simply started with some (expert-defined) error function, as was originally done for OLS and using $E(\mathbf{w})$. However, the statistical framework provides insights into the assumptions behind OLS regression. In particular, the assumptions include that the data \mathcal{D} was drawn i.i.d.; there is an underlying linear relationship between features and the target; that the noise (error term) is zero-mean Gaussian and independent of the features; and that there is an absence of noise in the collection of features.

4.2 Ordinary Least-Squares (OLS) Regression

To minimize the sum of squared errors, we shall first re-write $E(\mathbf{w})$ as

$$\begin{aligned}
E(\mathbf{w}) &= \sum_{i=1}^n (f(\mathbf{x}_i) - y_i)^2 \\
&= \sum_{i=1}^n \left(\sum_{j=0}^k w_j x_{ij} - y_i \right)^2,
\end{aligned}$$

where, again, we expanded each data point \mathbf{x}_i by $x_{i0} = 1$ to simplify the expression.

We now calculate the gradient $\nabla E(\mathbf{w})$ and the Hessian matrix $H_{E(\mathbf{w})}$. Finding weights for which $\nabla E(\mathbf{w}) = \mathbf{0}$ will result in an extremum while a

positive semi-definite Hessian will ensure that the extremum is a minimum. Now, we set the partial derivatives to 0 and solve the equations for each weight w_j

$$\begin{aligned}\frac{\partial E}{\partial w_0} &= 2 \sum_{i=1}^n \left(\sum_{j=0}^k w_j x_{ij} - y_i \right) x_{i0} = 0 \\ \frac{\partial E}{\partial w_1} &= 2 \sum_{i=1}^n \left(\sum_{j=0}^k w_j x_{ij} - y_i \right) x_{i1} = 0 \\ &\vdots \\ \frac{\partial E}{\partial w_k} &= 2 \sum_{i=1}^n \left(\sum_{j=0}^k w_j x_{ij} - y_i \right) x_{ik} = 0\end{aligned}$$

This results in a system of $k + 1$ linear equations with $k + 1$ unknowns that can be routinely solved (e.g. by using Gaussian elimination).

While this formulation is useful, it does not allow us to obtain a closed-form solution for \mathbf{w} or discuss the existence or multiplicity of solutions. To address the first point we will exercise some matrix calculus, while the remaining points will be discussed later. We will first write the sum of square errors using the matrix notation as

$$\begin{aligned}E(\mathbf{w}) &= (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2,\end{aligned}$$

where $\|\mathbf{v}\|_2 = \sqrt{\mathbf{v}^\top \mathbf{v}} = \sqrt{v_1^2 + v_2^2 + \dots + v_n^2}$ is the length of vector \mathbf{v} ; it is also called the L_2 norm. We can now formalize the *ordinary least-squares (OLS) linear regression problem* as

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2.$$

We proceed by finding $\nabla E(\mathbf{w})$ and $H_{E(\mathbf{w})}$. The gradient function $\nabla E(\mathbf{w})$ is a derivative of a scalar with respect to a vector. However, the intermediate steps of calculating the gradient require derivatives of vectors with respect to vectors (some of the rules of such derivatives are shown in Table 4.1). Application of the rules from Table 4.1 results in

$$\nabla E(\mathbf{w}) = 2\mathbf{X}^\top \mathbf{X}\mathbf{w} - 2\mathbf{X}^\top \mathbf{y}$$

\mathbf{y}	$\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$
$\mathbf{A}\mathbf{x}$	\mathbf{A}^\top
$\mathbf{x}^\top \mathbf{A}$	\mathbf{A}
$\mathbf{x}^\top \mathbf{x}$	$2\mathbf{x}$
$\mathbf{x}^\top \mathbf{A}\mathbf{x}$	$\mathbf{A}\mathbf{x} + \mathbf{A}^\top \mathbf{x}$

Table 4.1: Useful derivative formulas of vectors with respect to vectors. The derivative of vector \mathbf{y} (say m -dimensional) with respect to vector \mathbf{x} (say n -dimensional) is an $n \times m$ matrix \mathbf{M} with components $M_{ij} = \partial y_j / \partial x_i$, $i \in \{1, 2, \dots, n\}$ and $j \in \{1, 2, \dots, m\}$. A derivative of scalar with respect to a vector, e.g. the gradient, is a special case of this situation that results in an $n \times 1$ column vector.

and, therefore, from $\nabla E(\mathbf{w}) = \mathbf{0}$ we find that

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}. \quad (4.1)$$

The next step is to find the second derivative in order to ensure that we have not found a maximum (or saddle point). This results in

$$H_{E(\mathbf{w})} = 2\mathbf{X}^\top \mathbf{X},$$

which is a positive semi-definite matrix (why? Consider that for any vector $\mathbf{x} \neq \mathbf{0}$, $\mathbf{x}^\top \mathbf{A}^\top \mathbf{A} \mathbf{x} = (\mathbf{A}\mathbf{x})^\top \mathbf{A}\mathbf{x} = \|\mathbf{A}\mathbf{x}\|_2^2 \geq 0$, with equality only if the columns of \mathbf{A} are linearly dependent). Thus, we indeed have found a minimum. This is the global minimum because positive semi-definite Hessian implies convexity of $E(\mathbf{w})$. Furthermore, if the columns of \mathbf{X} are linearly independent, the Hessian is positive definite, which implies that the global minimum is unique. We can now express the predicted target values as

$$\begin{aligned} \hat{\mathbf{y}} &= \mathbf{X}\mathbf{w}^* \\ &= \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}. \end{aligned}$$

The matrix $\mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ is called the *projection matrix*; we will see later that it projects \mathbf{y} to the column space of \mathbf{X} .

Example 11: Consider again data set $\mathcal{D} = \{(1, 1.2), (2, 2.3), (3, 2.3), (4, 3.3)\}$ from Figure 4.1. We want to find the optimal coefficients of the least-squares

fit for $f(x) = w_0 + w_1x$ and then calculate the sum of squared errors on \mathcal{D} after the fit.

The OLS fitting can now be performed using

$$\mathbf{X} = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix}, \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} 1.2 \\ 2.3 \\ 2.3 \\ 3.3 \end{bmatrix},$$

where a column of ones was added to \mathbf{X} to allow for a non-zero intercept ($y = w_0$ when $x = 0$). Substituting \mathbf{X} and \mathbf{y} into Eq. (4.1) results in $\mathbf{w}^* = (0.7, 0.63)$ and the sum of square errors is $E(\mathbf{w}^*) = 0.223$. \square

As seen in the example above, it is a standard practice to add a column of ones to the data matrix \mathbf{X} in order to ensure that the fitted line, or generally a hyperplane, does not have to pass through the origin of the coordinate system. This effect, however, can be achieved in other ways. Consider the first component of the gradient vector

$$\frac{\partial E}{\partial w_0} = 2 \sum_{i=1}^n \left(\sum_{j=0}^k w_j x_{ij} - y_i \right) x_{i0} = 0$$

from which we obtain that

$$\sum_{i=1}^n w_0 = \sum_{i=1}^n y_i - \sum_{j=1}^k w_j \sum_{i=1}^n x_{ij}.$$

When all features (columns of \mathbf{X}) are normalized to have zero mean, i.e. when $\sum_{i=1}^n x_{ij} = 0$ for any column j , it follows that

$$w_0 = \frac{1}{n} \sum_{i=1}^n y_i.$$

We see now that if the target variable is normalized to the zero mean as well, it follows that $w_0 = 0$ and that the column of ones is not needed.

4.2.1 Weighted error function

In some applications it is useful to consider minimizing the weighted error function

$$E(\mathbf{w}) = \sum_{i=1}^n c_i \left(\sum_{j=0}^k w_j x_{ij} - y_i \right)^2,$$

where $c_i > 0$ is a cost for data point i . Expressing this in a matrix form, the goal is to minimize $(\mathbf{X}\mathbf{w} - \mathbf{y})^\top \mathbf{C} (\mathbf{X}\mathbf{w} - \mathbf{y})$, where $\mathbf{C} = \text{diag}(c_1, c_2, \dots, c_n)$. Using a similar approach as above, it can be shown that the weighted least-squares solution \mathbf{w}_C can be expressed as

$$\mathbf{w}_C = \left(\mathbf{X}^\top \mathbf{C} \mathbf{X} \right)^{-1} \mathbf{X}^\top \mathbf{C} \mathbf{y}.$$

In addition, it can be derived that

$$\mathbf{w}_C = \mathbf{w}_{OLS} + \left(\mathbf{X}^\top \mathbf{C} \mathbf{X} \right)^{-1} \mathbf{X}^\top (\mathbf{I} - \mathbf{C}) (\mathbf{X} \mathbf{w}_{OLS} - \mathbf{y}),$$

where \mathbf{w}_{OLS} is provided by Eq. (4.1). We can see that the solutions are identical when $\mathbf{C} = \mathbf{I}$, but also when $\mathbf{X} \mathbf{w}_{OLS} = \mathbf{y}$.

4.2.2 Predicting multiple outputs simultaneously

The extension to multiple outputs is straightforward, where now the target is an k -dimensional vector, $\mathbf{y} \in \mathbb{R}^k$, rather than a scalar, giving target matrix $\mathbf{Y} \in \mathbb{R}^{n \times k}$. Correspondingly, the weights $\mathbf{W} \in \mathbb{R}^{d \times k}$ to give $\mathbf{xW} \in \mathbb{R}^k$, with error

$$\begin{aligned} E(\mathbf{W}) &= \|\mathbf{XW} - \mathbf{Y}\|_F^2 = \sum_{i=1}^n \|\mathbf{X}_{i,:} \mathbf{W} - \mathbf{Y}_{i,:}\|_2^2 &> \text{Frobenius norm} \\ &= \text{trace} \left((\mathbf{XW} - \mathbf{Y})^\top (\mathbf{XW} - \mathbf{Y}) \right) \end{aligned}$$

and solution

$$\mathbf{W}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}.$$

Exercise: Derive this solution, by taking partial derivatives or, preferably, by using gradient rules for matrix variables. A good resource for matrix gradients is the matrix cookbook.

4.2.3 Expectation and variance for the solution vector

Under the data generating model presented above, we shall now treat the solution vector (estimator) \mathbf{w}^* as a random variable and investigate its statistical properties. For each of the data points, we have $Y_i = \sum_{j=0}^k \omega_j X_{ij} + \varepsilon_i$, and use $\boldsymbol{\varepsilon} = (\varepsilon_1, \varepsilon_2, \dots, \varepsilon_n)$ to denote a vector of i.i.d. random variables each

drawn according to $N(0, \sigma^2)$. Let us now look at the expected value (with respect to training data set D) for the weight vector \mathbf{w}^* :

$$\begin{aligned} E_D[\mathbf{w}^*] &= E \left[\left(\mathbf{X}^\top \mathbf{X} \right)^{-1} \mathbf{X}^\top (\mathbf{X} \boldsymbol{\omega} + \boldsymbol{\varepsilon}) \right] \\ &= \boldsymbol{\omega}, \end{aligned}$$

since $E[\boldsymbol{\varepsilon}] = \mathbf{0}$. An estimator whose expected value is the true value of the parameter is called an *unbiased estimator*. The covariance matrix for the optimal set of parameters can be expressed as

$$\begin{aligned} \Sigma[\mathbf{w}^*] &= E \left[(\mathbf{w}^* - \boldsymbol{\omega}) (\mathbf{w}^* - \boldsymbol{\omega})^\top \right] \\ &= E \left[\mathbf{w}^* \mathbf{w}^{*T} \right] - \boldsymbol{\omega} \boldsymbol{\omega}^\top \end{aligned}$$

Taking $\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$, we have

$$\begin{aligned} \Sigma[\mathbf{w}^*] &= E \left[\left(\boldsymbol{\omega} + \mathbf{X}^\dagger \boldsymbol{\varepsilon} \right) \left(\boldsymbol{\omega} + \mathbf{X}^\dagger \boldsymbol{\varepsilon} \right)^\top \right] - \boldsymbol{\omega} \boldsymbol{\omega}^\top \\ &= \boldsymbol{\omega} \boldsymbol{\omega}^\top + E \left[\mathbf{X}^\dagger \boldsymbol{\varepsilon} \boldsymbol{\varepsilon}^\top \mathbf{X}^{\dagger T} \right] - \boldsymbol{\omega} \boldsymbol{\omega}^\top \\ &= \mathbf{X}^\dagger E \left[\boldsymbol{\varepsilon} \boldsymbol{\varepsilon}^\top \right] \mathbf{X}^{\dagger T} \end{aligned}$$

because $\mathbf{X}^\dagger \mathbf{X} = \mathbf{I}$ and $E[\boldsymbol{\varepsilon}] = \mathbf{0}$. Thus, we have

$$\Sigma[\mathbf{w}^*] = \left(\mathbf{X}^\top \mathbf{X} \right)^{-1} \sigma^2,$$

where we exploited the facts that $E[\boldsymbol{\varepsilon} \boldsymbol{\varepsilon}^\top] = \sigma^2 \mathbf{I}$ and that an inverse of a symmetric matrix is also a symmetric matrix. It can be shown that estimator $\mathbf{w}^* = \mathbf{X}^\dagger \mathbf{y}$ is the one with the smallest variance among all unbiased estimators (Gauss-Markov theorem). These results are important not only because the estimated coefficients are expected to match the true coefficients of the underlying linear relationship, but also because the matrix inversion in the covariance matrix suggests stability problems in cases of singular and nearly singular matrices. Numerical stability and sensitivity to perturbations will be discussed later.

4.3 An Algebraic Perspective

Another powerful tool for analyzing and understanding linear regression comes from linear and applied linear algebra. In this section we take a

detour to address fundamentals of linear algebra and then apply these concepts to deepen our understanding of regression. In linear algebra, we are frequently interested in solving the following set of equations, given below in a matrix form

$$\mathbf{Ax} = \mathbf{b}. \quad (4.2)$$

Here, \mathbf{A} is an $m \times n$ matrix, \mathbf{b} is an $m \times 1$ vector, and \mathbf{x} is an $n \times 1$ vector that is to be found. All elements of \mathbf{A} , \mathbf{x} , and \mathbf{b} are considered to be real numbers. We shall start with a simple scenario and assume \mathbf{A} is a square, 2×2 matrix. This set of equations can be expressed as

$$\begin{aligned} a_{11}x_1 + a_{12}x_2 &= b_1 \\ a_{21}x_1 + a_{22}x_2 &= b_2 \end{aligned}$$

For example, we may be interested in solving

$$\begin{aligned} x_1 + 2x_2 &= 3 \\ x_1 + 3x_2 &= 5 \end{aligned}$$

This is a convenient formulation when we want to solve the system, e.g. by Gaussian elimination. However, it is not a suitable formulation to understand the question of the existence of solutions. In order for us to do this, we briefly review the basic concepts in linear algebra.

4.3.1 The four fundamental subspaces

The objective of this section is to briefly review the *four fundamental subspaces* in linear algebra (column space, row space, nullspace, left nullspace) and their mutual relationship. We shall start with our example from above and write the system of linear equations as

$$\begin{bmatrix} 1 \\ 1 \end{bmatrix} x_1 + \begin{bmatrix} 2 \\ 3 \end{bmatrix} x_2 = \begin{bmatrix} 3 \\ 5 \end{bmatrix}.$$

We can see now that by solving $\mathbf{Ax} = \mathbf{b}$ we are looking for the right amounts of vectors $(1, 1)$ and $(2, 3)$ so that their linear combination produces $(3, 5)$; these amounts are $x_1 = -1$ and $x_2 = 2$. Let us define $\mathbf{a}_1 = (1, 1)$ and $\mathbf{a}_2 = (2, 3)$ to be the column vectors of \mathbf{A} ; i.e. $\mathbf{A} = [\mathbf{a}_1 \ \mathbf{a}_2]$. Thus, $\mathbf{Ax} = \mathbf{b}$ will be solvable whenever \mathbf{b} can be expressed as a linear combination of the column vectors \mathbf{a}_1 and \mathbf{a}_2 .

All linear combinations of the columns of matrix \mathbf{A} constitute the *column space* of \mathbf{A} , $C(\mathbf{A})$, with vectors $\mathbf{a}_1 \dots \mathbf{a}_n$ being a basis of this space. Both \mathbf{b} and $C(\mathbf{A})$ lie in the m -dimensional space \mathbb{R}^m . Therefore, what $\mathbf{Ax} = \mathbf{b}$ is saying is that \mathbf{b} must lie in the column space of \mathbf{A} for the equation to have solutions. In the example above, if columns of \mathbf{A} are linearly independent¹, the solution is unique, i.e. there exists only one linear combination of the column vectors that will give \mathbf{b} . Otherwise, because \mathbf{A} is a square matrix, the system has no solutions. An example of such a situation is

$$\begin{bmatrix} 1 & 2 \\ 1 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 3 \\ 5 \end{bmatrix},$$

where $\mathbf{a}_1 = (1,1)$ and $\mathbf{a}_2 = (2,2)$. Here, \mathbf{a}_1 and \mathbf{a}_2 are (linearly) dependent because $2\mathbf{a}_1 - \mathbf{a}_2 = \mathbf{0}$. There is a deep connection between the spaces generated by a set of vectors and the properties of the matrix \mathbf{A} . For now, using the example above, it suffices to say that if \mathbf{a}_1 and \mathbf{a}_2 are independent the matrix \mathbf{A} is non-singular (singularity can be discussed only for square matrices), that is of full rank.

In an equivalent manner to the column space, all linear combinations of the rows of \mathbf{A} constitute the *row space*, denoted by $C(\mathbf{A}^\top)$, where both \mathbf{x} and $C(\mathbf{A}^\top)$ are in \mathbb{R}^n . All solutions to $\mathbf{Ax} = \mathbf{0}$ constitute the *nullspace* of the matrix, $N(\mathbf{A})$, while all solutions of $\mathbf{A}^\top \mathbf{y} = \mathbf{0}$ constitute the so-called *left nullspace* of \mathbf{A} , $N(\mathbf{A}^\top)$. Clearly, $C(\mathbf{A})$ and $N(\mathbf{A}^\top)$ are embedded in \mathbb{R}^m , whereas $C(\mathbf{A}^\top)$ and $N(\mathbf{A})$ are in \mathbb{R}^n . However, the pairs of subspaces are orthogonal (vectors \mathbf{u} and \mathbf{v} are orthogonal if $\mathbf{u}^\top \mathbf{v} = 0$); that is, any vector in $C(\mathbf{A})$ is orthogonal to all vectors from $N(\mathbf{A}^\top)$ and any vector in $C(\mathbf{A}^\top)$ is orthogonal to all vectors from $N(\mathbf{A})$. This is easy to see: if $\mathbf{x} \in N(\mathbf{A})$, then by definition $\mathbf{Ax} = \mathbf{0}$, and thus each row of \mathbf{A} is orthogonal to \mathbf{x} . If each row is orthogonal to \mathbf{x} , then so are all linear combinations of rows.

Orthogonality is a key property of the four subspaces, as it provides useful decomposition of vectors \mathbf{x} and \mathbf{b} from Eq. (4.2) with respect to \mathbf{A} (we will exploit this in the next Section). For example, any $\mathbf{x} \in \mathbb{R}^n$ can be decomposed as

$$\mathbf{x} = \mathbf{x}_r + \mathbf{x}_n,$$

where $\mathbf{x}_r \in C(\mathbf{A}^\top)$ and $\mathbf{x}_n \in N(\mathbf{A})$, such that $\|\mathbf{x}\|_2^2 = \|\mathbf{x}_r\|_2^2 + \|\mathbf{x}_n\|_2^2$. Similarly, every $\mathbf{b} \in \mathbb{R}^m$ can be decomposed as

$$\mathbf{b} = \mathbf{b}_c + \mathbf{b}_l,$$

¹As a reminder, two vectors are independent if their linear combination is only zero when both x_1 and x_2 are zero.

where $\mathbf{b}_c \in C(\mathbf{A})$, $\mathbf{b}_l \in N(\mathbf{A}^\top)$, and $\|\mathbf{b}\|_2^2 = \|\mathbf{b}_c\|_2^2 + \|\mathbf{b}_l\|_2^2$.

We mentioned above that the properties of fundamental spaces are tightly connected with the properties of matrix \mathbf{A} . To conclude this section, let us briefly discuss the *rank* of a matrix and its relationship with the dimensions of the fundamental subspaces. The *basis* of the space is the smallest set of vectors that span the space (this set of vectors is not unique). The size of the basis is also called the dimension of the space. In the example at the beginning of this subsection, we had a two dimensional column space with basis vectors $\mathbf{a}_1 = (1, 1)$ and $\mathbf{a}_2 = (2, 3)$. On the other hand, for $\mathbf{a}_1 = (1, 1)$ and $\mathbf{a}_2 = (2, 2)$ we had a one dimensional column space, i.e. a line, fully determined by any of the basis vectors. Unsurprisingly, the dimension of the space spanned by column vectors equals the *rank* of matrix \mathbf{A} . One of the fundamental results in linear algebra is that the rank of \mathbf{A} is identical to the dimension of $C(\mathbf{A})$, which in turn is identical to the dimension of $C(\mathbf{A}^\top)$.

4.3.2 Minimizing $\|\mathbf{Ax} - \mathbf{b}\|_2$

Let us now look again at the solutions to $\mathbf{Ax} = \mathbf{b}$. In general, there are three different outcomes:

1. there are no solutions to the system
2. there is a unique solution to the system, and
3. there are infinitely many solutions.

These outcomes depend on the relationship between the rank (r) of \mathbf{A} and dimensions m and n . We already know that when $r = m = n$ (square, invertible, full rank matrix \mathbf{A}) there is a unique solution to the system, but let us investigate other situations. Generally, when $r = n < m$ (full column rank), the system has either one solution or no solutions, as we will see momentarily. When $r = m < n$ (full row rank), the system has infinitely many solutions. Finally, in cases when $r < m$ and $r < n$, there are either no solutions or there are infinitely many solutions. Because $\mathbf{Ax} = \mathbf{b}$ may not be solvable, we generalize solving $\mathbf{Ax} = \mathbf{b}$ to minimizing $\|\mathbf{Ax} - \mathbf{b}\|_2$. In such a way, all situations can be considered in a unified framework.

Let us consider the following example

$$\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix}, \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix},$$




Figure 4.2: Illustration of the projection of vector \mathbf{b} to the column space of matrix \mathbf{A} . Vectors \mathbf{p} (\mathbf{b}_c) and \mathbf{e} (\mathbf{b}_l) represent the projection point and the error, respectively.

which illustrates an instance where we are unlikely to have a solution to $\mathbf{Ax} = \mathbf{b}$, unless there is some constraint on b_1 , b_2 , and b_3 ; here, the constraint is $b_3 = 2b_2 - b_1$. In this situation, $C(\mathbf{A})$ is a 2D plane in \mathbb{R}^3 spanned by the column vectors $\mathbf{a}_1 = (1, 1, 1)$ and $\mathbf{a}_2 = (2, 3, 4)$. If the constraint on the elements of \mathbf{b} is not satisfied, our goal is to try to find a point in $C(\mathbf{A})$ that is closest to \mathbf{b} . This happens to be the point where \mathbf{b} is projected to $C(\mathbf{A})$, as shown in Figure 4.2. We will refer to the projection of \mathbf{b} to $C(\mathbf{A})$ as \mathbf{p} . Now, using the standard algebraic notation, we have the following equations

$$\begin{aligned}\mathbf{b} &= \mathbf{p} + \mathbf{e} \\ \mathbf{p} &= \mathbf{Ax}\end{aligned}$$

Since \mathbf{p} and \mathbf{e} are orthogonal, we know that $\mathbf{p}^\top \mathbf{e} = 0$. Let us now solve for \mathbf{x}

$$\begin{aligned}(\mathbf{Ax})^\top (\mathbf{b} - \mathbf{Ax}) &= 0 \\ \mathbf{x}^\top \mathbf{A}^\top \mathbf{b} - \mathbf{x}^\top \mathbf{A}^\top \mathbf{Ax} &= 0 \\ \mathbf{x}^\top (\mathbf{A}^\top \mathbf{b} - \mathbf{A}^\top \mathbf{Ax}) &= 0\end{aligned}$$

and thus

$$\mathbf{x}^* = \left(\mathbf{A}^\top \mathbf{A} \right)^{-1} \mathbf{A}^\top \mathbf{b}.$$

This is exactly the same solution as one that minimized the sum of squared errors and maximized the likelihood. The matrix

$$\mathbf{A}^\dagger = \left(\mathbf{A}^\top \mathbf{A} \right)^{-1} \mathbf{A}^\top$$

is called the Moore-Penrose pseudo-inverse or simply a pseudo-inverse. This is an important matrix because it always exists and is unique, even in situations when the inverse of $\mathbf{A}^\top \mathbf{A}$ does not exist. This happens when \mathbf{A} has dependent columns (technically, \mathbf{A} and $\mathbf{A}^\top \mathbf{A}$ will have the same nullspace that contains more than just the origin of the coordinate system; thus the rank of $\mathbf{A}^\top \mathbf{A}$ is less than n). Let us for a moment look at the projection vector \mathbf{p} . We have

$$\begin{aligned}\mathbf{p} &= \mathbf{A}\mathbf{x} \\ &= \mathbf{A} \left(\mathbf{A}^\top \mathbf{A} \right)^{-1} \mathbf{A}^\top \mathbf{b},\end{aligned}$$

where $\mathbf{A} \left(\mathbf{A}^\top \mathbf{A} \right)^{-1} \mathbf{A}^\top$ is the matrix that projects \mathbf{b} to the column space of \mathbf{A} .

While we arrived at the same result as in previous sections, the tools of linear algebra allow us to discuss OLS regression at a deeper level. Let us examine for a moment the existence and multiplicity of solutions to

$$\arg \min_{\mathbf{x}} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2. \quad (4.3)$$

Clearly, the solution to this problem always exists. However, we shall now see that the solution to this problem is generally not unique and that it depends on the rank of \mathbf{A} . Consider \mathbf{x} to be one solution to Eq. (4.3). Recall that $\mathbf{x} = \mathbf{x}_r + \mathbf{x}_n$ and that it is multiplied by \mathbf{A} ; thus any vector $\mathbf{x} = \mathbf{x}_r + \alpha \mathbf{x}_n$, where $\alpha \in \mathbb{R}$, is also a solution. Observe that \mathbf{x}_r is common to all such solutions; if you cannot see it, assume there exists another vector from the row space and show that it is not possible. If the columns of \mathbf{A} are independent, the solution is unique because the nullspace contains only the origin. Otherwise, there are infinitely many solutions. In such cases, what exactly is the solution found by projecting \mathbf{b} to $C(\mathbf{A})$? Let us look at it:

$$\begin{aligned}\mathbf{x}^* &= \mathbf{A}^\dagger \mathbf{b} \\ &= \left(\mathbf{A}^\top \mathbf{A} \right)^{-1} \mathbf{A}^\top (\mathbf{p} + \mathbf{e}) \\ &= \left(\mathbf{A}^\top \mathbf{A} \right)^{-1} \mathbf{A}^\top \mathbf{p} \\ &= \mathbf{x}_r,\end{aligned}$$

as $\mathbf{p} = \mathbf{A}\mathbf{x}_r$. Given that \mathbf{x}_r is unique, the solution found by the least squares optimization is the one that simultaneously minimizes $\|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2$ and $\|\mathbf{x}\|_2$ (observe that $\|\mathbf{x}\|_2$ is minimized because the solution ignores any component

from the nullspace). Thus, the OLS regression problem is sometimes referred to as the minimum-norm least-squares problem.

Let us now consider situations where $\mathbf{Ax} = \mathbf{b}$ has infinitely many solutions, i.e. when $\mathbf{b} \in C(\mathbf{A})$. This usually arises when $r \leq m < n$. Here, because \mathbf{b} is already in the column space of \mathbf{A} , the only question is what particular solution \mathbf{x} will be found by the minimization procedure. As we have seen above, the outcome of the minimization process is the solution with the minimum L_2 norm $\|\mathbf{x}\|_2$.

To summarize, the goal of the OLS regression problem is to solve $\mathbf{Xw} = \mathbf{y}$, if it is solvable. When $k < n$ this is not a realistic scenario in practice. Thus, we relaxed the requirement and tried to find the point in the column space $C(\mathbf{X})$ that is closest to \mathbf{y} . This turned out to be equivalent to minimizing the sum of square errors (or Euclidean distance) between n -dimensional vectors \mathbf{Xw} and \mathbf{y} . It also turned out to be equivalent to the maximum likelihood solution presented in Section 4.1. When $n < k$, a usual situation in practice is that there are infinitely many solutions. In these situations, our optimization algorithm will find the one with the minimum L_2 norm.

4.4 Linear regression for non-linear problems

At first, it might seem that the applicability of linear regression and classification to real-life problems is greatly limited. After all, it is not clear whether it is realistic (most of the time) to assume that the target variable is a linear combination of features. Fortunately, the applicability of linear regression is broader than originally thought. The main idea is to apply a non-linear transformation to the data matrix \mathbf{X} prior to the fitting step, which then enables a non-linear fit. Obtaining such a useful feature representation is a central problem in machine learning; we will discuss this in detail in Chapter 7. Here, we will first examine a simpler expanded representation that enables non-linear learning: polynomial curve fitting.

4.4.1 Polynomial curve fitting

We start with one-dimensional data. In OLS regression, we would look for the fit in the following form

$$f(x) = w_0 + w_1x,$$




Figure 4.3: Transformation of an $n \times 1$ data matrix \mathbf{X} into an $n \times (p+1)$ matrix Φ using a set of basis functions ϕ_j , $j = 0, 1, \dots, p$.

where x is the data point and $\mathbf{w} = (w_0, w_1)$ is the weight vector. To achieve a polynomial fit of degree p , we will modify the previous expression into

$$f(x) = \sum_{j=0}^p w_j x^j,$$

where p is the degree of the polynomial. We will rewrite this expression using a set of basis functions as

$$\begin{aligned} f(x) &= \sum_{j=0}^p w_j \phi_j(x) \\ &= \mathbf{w}^\top \boldsymbol{\phi}, \end{aligned}$$

where $\phi_j(x) = x^j$ and $\boldsymbol{\phi} = (\phi_0(x), \phi_1(x), \dots, \phi_p(x))$. Applying this transformation to every data point in \mathbf{X} results in a new data matrix Φ , as shown in Figure 4.3.

Following the discussion from Section 4.2, the optimal set of weights is now calculated as

$$\mathbf{w}^* = \left(\Phi^\top \Phi \right)^{-1} \Phi^\top \mathbf{y}.$$

Example 12: In Figure 4.1 we presented an example of a data set with four data points. What we did not mention was that, given a set $\{x_1, x_2, x_3, x_4\}$, the targets were generated by using function $1 + \frac{x}{2}$ and then adding a measurement error $\mathbf{e} = (-0.3, 0.3, -0.2, 0.3)$. It turned out that the optimal coefficients $\mathbf{w}^* = (0.7, 0.63)$ were close to the true coefficients $\boldsymbol{\omega} = (1, 0.5)$, even though the error terms were relatively significant. We will now attempt to estimate the coefficients of a polynomial fit with degrees $p = 2$ and $p = 3$. We will also calculate the sum of squared errors on \mathcal{D} after the fit as well




Figure 4.4: Example of a linear vs. polynomial fit on a data set shown in Figure 4.1. The linear fit, $f_1(x)$, is shown as a solid green line, whereas the cubic polynomial fit, $f_3(x)$, is shown as a solid blue line. The dotted red line indicates the target linear concept.

as on a large discrete set of values $x \in \{0, 0.1, 0.2, \dots, 10\}$ where the target values will be generated using the true function $1 + \frac{x}{2}$.

Using a polynomial fit with degrees $p = 2$ and $p = 3$ results in $\mathbf{w}_2^* = (0.575, 0.755, -0.025)$ and $\mathbf{w}_3^* = (-3.1, 6.6, -2.65, 0.35)$, respectively. The sum of squared errors on \mathcal{D} equals $E(\mathbf{w}_2^*) = 0.221$ and $E(\mathbf{w}_3^*) \approx 0$. Thus, the best fit is achieved with the cubic polynomial. However, the sum of squared errors on the outside data set reveal a poor generalization ability of the cubic model because we obtain $E(\mathbf{w}^*) = 26.9$, $E(\mathbf{w}_2^*) = 3.9$, and $E(\mathbf{w}_3^*) = 22018.5$. This effect is called *overfitting*. Broadly speaking, overfitting is indicated by a significant difference in fit between the data set on which the model was trained and the outside data set on which the model is expected to be applied (Figure 4.4). In this case, the overfitting occurred because the complexity of the model was increased considerably, whereas the size of the data set remained small.

One signature of overfitting is an increase in the magnitude of the coefficients. For example, while the absolute values of all coefficients in \mathbf{w}^* and \mathbf{w}_2^* were less than one, the values of the coefficients in \mathbf{w}_3^* became significantly larger with alternating signs (suggesting overcompensation). We will discuss *regularization* in Section 4.5.2 as an approach to prevent this effect. \square

Polynomial curve fitting is only one way of non-linear fitting because the choice of basis functions need not be limited to powers of x . Among others, non-linear basis functions that are commonly used are the sigmoid function

$$\phi_j(x) = \frac{1}{1 + e^{-\frac{x - \mu_j}{s_j}}}$$

or a Gaussian-style exponential function

$$\phi_j(x) = e^{-\frac{(x - \mu_j)^2}{2\sigma_j^2}},$$

where μ_j , s_j , and σ_j are constants to be determined. However, this approach works only for a one-dimensional input \mathbf{X} . For higher dimensions, this approach needs to be generalized using *radial basis functions*; see Section 7.1 for more details.

4.5 Practical considerations

In this section, we will discuss a few important considerations when implementing and using regression algorithms in practice.

4.5.1 Stability and sensitivity of the solutions

In practice, data sets include large numbers of features, which are sometimes identical, similar, or nearly linearly dependent. This causes $\mathbf{X}^\top \mathbf{X}$ to no longer be invertible and produces solutions that are sensitive to perturbations in \mathbf{y} and \mathbf{X} . To see why, we will look at the *singular value decomposition* of \mathbf{X} . As with the previous linear algebra constructs, it allows us to easily examine properties of \mathbf{X} .

The singular value decomposition of $\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$ for orthonormal matrices² $\mathbf{U} \in \mathbb{R}^{n \times n}$, $\mathbf{V} \in \mathbb{R}^{d \times d}$ and non-negative diagonal matrix $\mathbf{\Sigma} \in \mathbb{R}^{n \times d}$. The diagonal entries in $\mathbf{\Sigma}$ are the singular values; the number of non-zero singular values constitutes the rank of \mathbf{X} . Any matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$ can be decomposed into its singular value decomposition, because any linear transformation can be decomposed into a rotation (multiplication by \mathbf{V}^\top), followed by a scaling (multiplication by $\mathbf{\Sigma}$), followed again by a rotation (multiplication by \mathbf{U}).

Now we can discuss the pseudo-inverse of \mathbf{X} in terms of the singular value decomposition. Notice that

$$\mathbf{X}^\top \mathbf{X} = \mathbf{V}\mathbf{\Sigma}\mathbf{U}^\top \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top = \mathbf{V}\mathbf{\Sigma}_d^2 \mathbf{V}^\top \quad \triangleright \quad \mathbf{\Sigma}_d = \mathbf{\Sigma}(1:d, 1:d)$$

²An orthonormal matrix \mathbf{U} is square matrix that satisfies $\mathbf{U}^\top \mathbf{U} = \mathbf{I}$ and $\mathbf{U}\mathbf{U}^\top = \mathbf{I}$

because \mathbf{U} is orthonormal. The inverse of $\mathbf{X}^\top \mathbf{X}$ exists if \mathbf{X} is full rank, i.e., Σ_d has no zeros on the diagonal, because $(\mathbf{X}^\top \mathbf{X})^{-1} = \mathbf{V} \Sigma_d^{-2} \mathbf{V}^\top$. In the general case, however, the pseudo-inverse is still defined as the inverse of the non-zero singular values, Σ^\dagger . This means that, even for \mathbf{X} that is not full rank, we obtain the solution

$$\mathbf{w} = \mathbf{X}^\dagger \mathbf{y} = \mathbf{V} \Sigma^\dagger \mathbf{U}^\top \mathbf{y} = \sum_{i=1}^{\text{rank of } \mathbf{X}} \frac{\mathbf{v}_i \mathbf{u}_i^\top}{\sigma_i} \mathbf{y}. \quad (4.4)$$

Notice that for full-rank \mathbf{X} ,

$$(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top = \mathbf{V} \Sigma_d^{-2} \mathbf{V}^\top \mathbf{V} \Sigma \mathbf{U}^\top = \mathbf{V} \Sigma_d^{-2} \Sigma \mathbf{U}^\top = \mathbf{V} \Sigma_d^{-1} \mathbf{U}^\top$$

which matches the definition of the pseudo-inverse using the singular value decomposition for non-full rank matrices.

The solution in Equation (4.4) makes it clear why the solution can be sensitive to perturbations. For small singular values, due to being nearly linearly dependent, σ_i^{-1} is large and amplifies any changes in \mathbf{y} . Moreover, for different samples in \mathbf{X} , and for different numbers of samples, we will see different small singular values, which will again strongly affect the solution.

A common strategy to deal with this instability is to drop or truncate small singular values. This is a form of regularization, which we discuss next.

4.5.2 Regularization

So far, we have discussed linear regression in terms of maximum likelihood. But, as before, we can also propose a MAP objective. Instead of specifying no prior over \mathbf{w} , we can select a prior to help *regularize* overfitting to the observed data. We will discuss two common priors (regularizers): the Gaussian prior (ℓ_2 norm) and the Laplace prior (ℓ_1 norm), shown in Figure 4.5.

Taking the log of the zero-mean Gaussian prior, $\mathcal{N}(\mathbf{0}, \lambda^{-1} I)$, we get

$$-\ln p(\mathbf{w}) = \ln(2\pi d \lambda^{-1}) + \frac{\mathbf{w}^\top \mathbf{w}}{2\lambda^{-1}} = \ln(2\pi d \lambda^{-1}) + \frac{\lambda}{2} \mathbf{w}^\top \mathbf{w}.$$

As before, we can drop the first constant which does not affect the selection of \mathbf{w} , and now combine the negative log-likelihood and the negative log prior to get the following ridge regression problem:

$$E(\mathbf{w}) = (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^\top \mathbf{w} \quad \triangleright \|\mathbf{w}\|_2^2 = \mathbf{w}^\top \mathbf{w}$$




Figure 4.5: A comparison between Gaussian and Laplace priors. The Gaussian prior prefers the values to be near zero, whereas the Laplace prior more strongly prefers the values to equal zero.

where λ is a user-selected parameter that is called the *regularization parameter*. The idea is to penalize weight coefficients that are too large; the larger the λ , the more large weights are penalized. If we solve this equation in a similar manner as before, we obtain

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}.$$

This has the nice effect of shifting the squared singular values in Σ_d^2 by λ , removing stability issues with dividing by small singular values, as long as λ is itself large enough.

Similarly, if we choose a Laplace distribution, we get an ℓ_1 penalized loss

$$E(\mathbf{w}) = (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \|\mathbf{w}\|_1.$$

As with the ℓ_2 regularizer for ridge regression, this regularizer penalized large values in \mathbf{w} . However, it also produces more sparse solutions, where entries in \mathbf{w} are zero. This has the effect of feature selection, because zeroing entries in \mathbf{w} is equivalent to removing the corresponding feature.

4.5.3 Handling big data sets

One common approach to handling big datasets is to use *stochastic approximation*, where samples are processed incrementally. To see how this would be done, let us revisit the gradient of the error function, $\nabla E(\mathbf{w})$. We obtained a closed form solution for $\nabla E(\mathbf{w}) = \mathbf{0}$; however, for many other error

functions, solving for $\nabla E(\mathbf{w}) = \mathbf{0}$ in a closed form way is not possible. Instead, we start at some initial \mathbf{w}_0 (typically random), and then step in the direction of the negative of the gradient until we reach a local minimum. This approach is called *gradient descent* and is summarized in Algorithm 1.

Algorithm 1: Batch Gradient Descent($E, \mathbf{X}, \mathbf{y}$)

```

1: // A non-optimized, basic implementation of batch gradient descent
2:  $\mathbf{w} \leftarrow$  random vector in  $\mathbb{R}^d$ 
3:  $\text{err} \leftarrow \infty$ 
4:  $\text{tolerance} \leftarrow 10e^{-4}$ 
5:  $\alpha \leftarrow 0.1$ 
6: while  $|E(\mathbf{w}) - \text{err}| > \text{tolerance}$  do
7:    $\text{err} \leftarrow E(\mathbf{w})$ 
8:   // The step-size  $\alpha$  should be chosen by line-search
9:    $\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla E(\mathbf{w}) = \mathbf{w} - \alpha \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$ 
10: end while
11: return  $\mathbf{w}$ 

```

Algorithm 2: Stochastic Gradient Descent($E, \mathbf{X}, \mathbf{y}$)

```

1:  $\mathbf{w} \leftarrow$  random vector in  $\mathbb{R}^d$ 
2: for  $t = 1, \dots, n$  do
3:   // For many settings, we need step-size  $\alpha_t$  to decrease with time
4:   // For example, a common choice is  $\alpha_t = \alpha_0/t^{-1}$  or
5:   //  $\alpha_t = \alpha_0/t^{-1/2}$  for some initial  $\alpha_0$ , such as  $\alpha_0 = 1.0$ .
6:    $\mathbf{w} \leftarrow \mathbf{w} - \alpha_t \nabla E_t(\mathbf{w}) = \mathbf{w} - \alpha_t (\mathbf{x}_t^\top \mathbf{w} - y_t) \mathbf{x}_t$ 
7: end for
8: return  $\mathbf{w}$ 

```

For a large number of samples n , however, computing the gradient across all samples can be expensive or infeasible. An alternative is to approximate the gradient less accurately with fewer samples. In stochastic approximation, we typically approximate the gradient with one sample³, as in Algorithm 2. Though this approach may appear to be too much of an approximation, there is a long theoretical and empirical history indicating its effectiveness (see for example [4, 3]). With ever increasing data-set size for many scenarios, the generality of stochastic approximation makes it arguably the modern approach to dealing with big data. For specialized scenarios, there are of course other approaches. For one example, see [9].

³Mini-batches are a way to obtain a better approximation but remain efficient.

Chapter 5

Generalized Linear Models

In previous sections, we saw that the statistical framework provided valuable insights into linear regression, especially with respect to explicitly stating most of the assumptions in the system (we will see the full picture only when Bayesian formulation is used). These assumptions were necessary to rigorously estimate parameters of the model, which could then be subsequently used for prediction on previously unseen data points. In this section, we introduce generalized linear models (GLMs) which extend ordinary least-squares regression beyond Gaussian probability distributions and linear dependencies between the features and the target.

We shall first revisit the main points of the ordinary least-squares regression. There, we assumed that a set of i.i.d. data points with their targets $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ were drawn according to some distribution $p(\mathbf{x}, y)$. We also assumed that an underlying relationship between the features and the target was linear, i.e.

$$Y = \sum_{j=0}^k \omega_j X_j + \varepsilon,$$

where $\boldsymbol{\omega}$ was a set of unknown weights and ε was a zero-mean normally distributed random variable with variance σ^2 . In order to simplify generalization, we will slightly reformulate this model. In particular, it will be useful to separate the underlying linear relationship between the features and the target from the fact that Y was normally distributed. That is, we will write that

1. $E[y|\mathbf{x}] = \boldsymbol{\omega}^T \mathbf{x}$
2. $p(y|\mathbf{x}) = N(\mu, \sigma^2)$

with $\mu = \boldsymbol{\omega}^T \mathbf{x}$ connecting the two expressions. This way of formulating linear regression will allow us (i) to generalize the framework to non-linear relationships between the features and the target as well as (ii) to use the error distributions other than Gaussian.

5.1 Loglinear link and Poisson distribution

Let us start with an example. Assume that data points correspond to cities in the world (described by some numerical features) and that the target variable is the number of sunny days observed in a particular year. To establish the GLM model, we will assume (1) a loglinear link between the expectation of the target and linear combination of features, and (2) the Poisson distribution for the target variable. We summarize these assumptions as follows

1. $\log(E[y|\mathbf{x}]) = \boldsymbol{\omega}^T \mathbf{x}$
2. $p(y|\mathbf{x}) = \text{Poisson}(\lambda)$

where $\lambda > 0$ is the parameter (mean and variance) of the Poisson distribution. Exploiting the fact that $E[y|\mathbf{x}] = \lambda$, we connect the two formulas using $\lambda = e^{\boldsymbol{\omega}^T \mathbf{x}}$. In fact, because $\lambda \in \mathbb{R}^+$ and $\boldsymbol{\omega}^T \mathbf{x} \in \mathbb{R}$, it is not appropriate to use a linear link between $E[y|\mathbf{x}]$ and $\boldsymbol{\omega}^T \mathbf{x}$ (i.e. $E[y|\mathbf{x}] = \boldsymbol{\omega}^T \mathbf{x}$). The link function adjusts the range of the linear combination of features (so-called systematic component) to the domain of the parameters (here, mean) of the probability distribution.

We provide a compact summary of the above assumptions via a probability distribution for the target; i.e. $p(y|\mathbf{x}) = \text{Poisson}(e^{\boldsymbol{\omega}^T \mathbf{x}})$. We express this as

$$p(y|\mathbf{x}) = \frac{e^{\boldsymbol{\omega}^T \mathbf{x} y} \cdot e^{-e^{\boldsymbol{\omega}^T \mathbf{x}}}}{y!}$$

for any $y \in \mathbb{N}$. We will now use the maximum likelihood estimation to find the parameters of the regression model. As in previous sections, the likelihood function has the form of the probability distribution, where the data set is observed and the parameters are unknown. Hence, the log-likelihood function has the form

$$l(\mathbf{w}) = \sum_{i=1}^n \mathbf{w}^T \mathbf{x}_i y_i - \sum_{i=1}^n e^{\mathbf{w}^T \mathbf{x}_i} - \sum_{i=1}^n y_i!$$

It is easy to show that $\nabla l(\mathbf{w}) = \mathbf{0}$ does not have a closed-form solution. Therefore, we will use the Newton-Raphson method in which we must first analytically find the gradient vector $\nabla l(\mathbf{w})$ and the Hessian matrix $H_{l(\mathbf{w})}$. We start by deriving the j -th element of the gradient

$$\begin{aligned}
\frac{\partial ll(\mathbf{w})}{\partial w_j} &= \sum_{i=1}^n x_{ij} y_i - \sum_{i=1}^n e^{\mathbf{w}^T \mathbf{x}_i} x_{ij} \\
&= \sum_{i=1}^n x_{ij} \cdot (y_i - e^{\mathbf{w}^T \mathbf{x}_i}) \\
&= \mathbf{f}_j^T \cdot (\mathbf{y} - \mathbf{c}),
\end{aligned}$$

where \mathbf{f}_j^T is the j -th column of \mathbf{X} and \mathbf{c} is a vector with elements $c_i = e^{\mathbf{w}^T \mathbf{x}_i}$. The gradient of the likelihood function can now be expressed as

$$\nabla ll(\mathbf{w}) = \mathbf{X}^T \cdot (\mathbf{y} - \mathbf{c}). \quad (5.1)$$

Note that \mathbf{c} stores a set of predictions for each of the data points, and thus, $\mathbf{y} - \mathbf{c}$ is an error vector. The second partial derivative of the likelihood function can be derived as

$$\begin{aligned}
\frac{\partial^2 ll(\mathbf{w})}{\partial w_j \partial w_k} &= - \sum_{i=1}^n x_{ij} \cdot e^{\mathbf{w}^T \mathbf{x}_i} \cdot x_{ik} \\
&= -\mathbf{f}_j^T \cdot \mathbf{C} \cdot \mathbf{f}_k,
\end{aligned}$$

where \mathbf{C} is an n -by- n diagonal matrix with $c_{ii} = e^{\mathbf{w}^T \mathbf{x}_i}$. The Hessian matrix can now be calculated as

$$H_{ll(\mathbf{w})} = -\mathbf{X}^T \cdot \mathbf{C} \cdot \mathbf{X}. \quad (5.2)$$

which is a negative semi-definite matrix. Substituting Eq. (5.1) and Eq. (5.2) into the Newton-Raphson formula results in the following weight update rule

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \left(\mathbf{X}^T \cdot \mathbf{C}^{(t)} \cdot \mathbf{X} \right)^{-1} \cdot \mathbf{X}^T \cdot (\mathbf{y} - \mathbf{c}^{(t)}),$$

where $\mathbf{c}^{(t)}$ and $\mathbf{C}^{(t)}$ are calculated using the weight vector $\mathbf{w}^{(t)}$. The initial set of weights $\mathbf{w}^{(0)}$ can be set randomly.

5.2 Exponential family of distributions

Exponential family is a class of probability distributions with the following form

$$p(x|\boldsymbol{\theta}) = c(\boldsymbol{\theta})h(x) \exp \left(\sum_{i=1}^m q_i(\boldsymbol{\theta})t_i(x) \right),$$

where $\boldsymbol{\theta} = (\theta_1, \theta_2, \dots, \theta_m)$ is a set of parameters. When $q_i(\boldsymbol{\theta}) = \theta_i$ for $\forall i$, parameters $\theta_1, \theta_2, \dots, \theta_m$ are called natural parameters. This leads to the following (canonical) form of the probability distribution

$$\begin{aligned} p(x|\boldsymbol{\theta}) &= \exp \left(\sum_{i=1}^m \theta_i t_i(x) - a(\boldsymbol{\theta}) + b(x) \right) \\ &= \exp (\boldsymbol{\theta}^T \mathbf{t}(x) - a(\boldsymbol{\theta}) + b(x)), \end{aligned}$$

where $\mathbf{t}(x) = (t_1(x), t_2(x), \dots, t_m(x))$. Many of the often encountered (families of) distributions are members of the exponential family; e.g. exponential, Gaussian, Gamma, Poisson, or the binomial distributions. Therefore, it is useful to generically study the exponential family to better understand commonalities and differences between individual member functions.

Let us look at a couple of examples. The Poisson distribution can be expressed as

$$p(x|\lambda) = \exp (x \log \lambda - \lambda - \log x!),$$

where $\lambda \in \mathbb{R}^+$ and $\mathcal{X} = \mathbb{N}_0$. Thus, $\theta = \log \lambda$, $t(x) = x$, $a(\theta) = e^\theta$, and $b(x) = -\log x!$. Similarly, the Gaussian distribution with mean μ and variance σ^2 can be written as

$$p(x|\mu, \sigma) = \exp \left(-\frac{x^2}{2\sigma^2} + x \frac{\mu}{\sigma^2} - \frac{\mu^2}{2\sigma^2} - \frac{1}{2} \log(2\pi\sigma^2) \right),$$

where $\mu \in \mathbb{R}$, $\sigma \in \mathbb{R}^+$ and $\mathcal{X} = \mathbb{R}$. Therefore, we see that

$$\begin{aligned} \boldsymbol{\theta} &= \left(\frac{\mu}{\sigma^2}, -\frac{1}{2\sigma^2} \right) \\ \mathbf{t}(x) &= (x, x^2) \\ a(\boldsymbol{\theta}) &= \frac{\theta_1^2}{4\theta_2} + \frac{1}{2} \log\left(\frac{\theta_2}{\pi}\right) \\ b(x) &= 0. \end{aligned}$$

From $\int_{\mathcal{X}} p(x)dx = 1$, we can easily derive that

$$a(\boldsymbol{\theta}) = \log \int_{\mathcal{X}} \exp(\boldsymbol{\theta}^T \mathbf{t}(x) + b(x)) dx.$$

Function $a(\boldsymbol{\theta})$ is called a log partitioning function or simply a log-normalizer. It can be derived that

$$\begin{aligned} \frac{\partial a(\boldsymbol{\theta})}{\partial \boldsymbol{\theta}} &= E[\mathbf{t}(x)] \\ \frac{\partial^2 a(\boldsymbol{\theta})}{\partial \boldsymbol{\theta} \partial \boldsymbol{\theta}^T} &= \text{cov}[\mathbf{t}(x)] \end{aligned}$$

These properties are very useful for estimating parameters of the distribution. For example, let us consider a data set of observations $\mathcal{D} = \{x_i\}_{i=1}^n$ and look at the log-likelihood function

$$\begin{aligned} ll(\boldsymbol{\theta}) &= \log \prod_{i=1}^n e^{\boldsymbol{\theta}^T \mathbf{t}(x_i) - a(\boldsymbol{\theta}) + b(x_i)} \\ &= \sum_{i=1}^n \boldsymbol{\theta}^T \mathbf{t}(x_i) - n \cdot a(\boldsymbol{\theta}) + \sum_{i=1}^n b(x_i). \end{aligned}$$

Maximizing the likelihood involves calculating the gradient function

$$\nabla ll(\boldsymbol{\theta}) = \sum_{i=1}^n \mathbf{t}(x_i) - n \nabla a(\boldsymbol{\theta}).$$

Setting the gradient to zero results in

$$\nabla a(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \mathbf{t}(x_i).$$

By combining the previous expressions we see that the likelihood is maximized when the gradient function of the log-normalizer equals the sample mean of $\mathbf{t}(x)$. This result is important because it provides a general expression for estimating the parameters of all distributions in the exponential family. Function $\mathbf{t}(x)$ is called a sufficient statistic (a statistic is simply a function of the data), because all information about the parameters $\boldsymbol{\theta}$ can be inferred by calculating $\mathbf{t}(x)$.

5.3 Formalizing generalized linear models

We shall now formalize the generalized linear models. The two key components of GLMs can be expressed as

1. $f(E[y|\mathbf{x}]) = \boldsymbol{\omega}^T \mathbf{x}$
2. $p(y|\mathbf{x}) \in \text{Exponential Family}$

Here, $f(\cdot)$ is called the link function between the linear combination of the features and parameters of the distribution. On the one hand, the link function adjusts the range of $\boldsymbol{\omega}^T \mathbf{x}$ to the domain of Y (because of this relationship, link functions are usually not selected independently of the distribution for Y). On the other hand, it also provides a mechanism for a non-linear relationship between the features and the target. While the nature of this non-linear relationship is limited (the features enter the system via a linear combination with the parameters) there is still an important flexibility they provide for modeling. Similarly, the generalization to the exponential family from the Gaussian distribution used in ordinary least-squares regression, allows us to model a much wider range of target functions. The choice of the link function and the probability distribution is data dependent.

Generally, there is no guarantee of a closed-form solution for \mathbf{w} . Therefore, GLM formulations usually resort to iterative techniques derived from the Taylor approximation of the log-likelihood. Hence, a single mechanism can be used for a wide range of link functions and probability distributions. Let us write the log-likelihood

$$\begin{aligned} ll(\mathbf{w}) &= \log \prod_{i=1}^n e^{\boldsymbol{\theta}^T \mathbf{t}(x_i) - a(\boldsymbol{\theta}) + b(x_i)} \\ &= \sum_i \sum_m \theta_m t_m(x_i) - n \cdot a(\boldsymbol{\theta}) + \sum_i b(x_i) \\ &= \sum_i ll_i(\mathbf{w}) \end{aligned}$$

and also find the elements of its gradient

$$\frac{\partial ll_i(\mathbf{w})}{\partial w_j} = \sum_m \frac{\partial \theta_m}{\partial w_j} t_m(x_i) - \frac{\partial a(\boldsymbol{\theta})}{\partial w_j}$$

which can be used to easily calculate the update rules of the optimization. Interestingly, the standard versions of the GLM from the literature do not use

the full version of the Newton-Raphson algorithm with the Hessian matrix. Instead, Gauss-Newton and other types of solutions are considered and are generally called iteratively reweighted least-squares (IRLS) algorithms in the statistical literature.

5.4 Logistic regression

At the end, we mention that GLMs extend to classification. One of the most popular uses of GLMs is a combination of a Bernoulli distribution with a logit link function. This framework is frequently encountered and is called logistic regression. We summarize the logistic regression model as follows

1. $\text{logit}(E[y|\mathbf{x}]) = \boldsymbol{\omega}^T \mathbf{x}$
2. $p(y|\mathbf{x}) = \text{Bernoulli}(\alpha)$

where $\text{logit}(x) = \ln \frac{x}{1-x}$, $y \in \{0, 1\}$, and $\alpha \in (0, 1)$ is the parameter (mean) of the Bernoulli distribution. It follows that

$$E[y|\mathbf{x}] = \frac{1}{1 + e^{-\boldsymbol{\omega}^T \mathbf{x}}}$$

and

$$p(y|\mathbf{x}) = \left(\frac{1}{1 + e^{-\boldsymbol{\omega}^T \mathbf{x}}} \right)^y \left(1 - \frac{1}{1 + e^{-\boldsymbol{\omega}^T \mathbf{x}}} \right)^{1-y}.$$

Given a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, where $\mathbf{x}_i \in \mathbb{R}^k$ and $y_i \in \{0, 1\}$, the parameters of the model \mathbf{w} can be found by maximizing the likelihood function.

Chapter 6

Linear Classifiers

Suppose we are interested in building a linear classifier $f : \mathbb{R}^d \rightarrow \{-1, +1\}$. Linear classifiers try to find the relationship between inputs and outputs by constructing a linear function (a point, a line, a plane or a hyperplane) that splits \mathbb{R}^d into two half-spaces. The two half-spaces act as decision regions for the positive and negative examples, respectively. Given a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ consisting of positive and negative examples, there are many ways in which linear classifiers can be constructed. For example, a training algorithm may explicitly work to position the decision surface in order to separate positive and negative examples according to some problem-relevant criteria; e.g. it may try to minimize the fraction of examples on the incorrect side of the decision surface. Alternatively, the goal of the training algorithm may be to directly estimate the posterior distribution $p(y|\mathbf{x})$, in which case the algorithm is more likely to rely on the formal parameter estimation principles; e.g. it may maximize the likelihood. An example of a classifier with a linear decision surface is shown in Figure 6.1.

To simplify the formalism in the following sections, we will add a component $x_0 = 1$ to each input (x_1, \dots, x_d) . This extends the input space to $\mathcal{X} = \mathbb{R}^{d+1}$ but, fortunately, it also leads us to a simplified notation in which the decision boundary in \mathbb{R}^d can be written as $\mathbf{w}^\top \mathbf{x} = 0$, where $\mathbf{w} = (w_0, w_1, \dots, w_d)$ is a set of weights and $\mathbf{x} = (x_0 = 1, x_1, \dots, x_d)$ is any element of the input space. Nevertheless, we should remember that the actual inputs are d -dimensional.

Earlier in the introductory remarks, we presented a classifier as a function $f : \mathcal{X} \rightarrow \mathcal{Y}$ and have transformed the learning problem into approximating $p(y|\mathbf{x})$. In the case of linear classifiers, our flexibility is restricted because our method must learn the posterior probabilities $p(y|\mathbf{x})$ and at the same time have a linear decision surface in \mathbb{R}^d . This, however, can be achieved if $p(y|\mathbf{x})$ is modeled as a monotonic function of $\mathbf{w}^\top \mathbf{x}$; e.g. $\tanh(\mathbf{w}^\top \mathbf{x})$ or $(1 + e^{-\mathbf{w}^\top \mathbf{x}})^{-1}$. Of course, a model trained to learn posterior probabilities $p(y|\mathbf{x})$ can be seen as a function $g : \mathcal{X} \rightarrow [0, 1]$. Then, the conversion from g




Figure 6.1: A data set in \mathbb{R}^2 consisting of nine positive and nine negative examples. The gray line represents a linear decision surface in \mathbb{R}^2 . The decision surface does not perfectly separate positives from negatives.

to f is a straightforward application of the maximum a posteriori principle: the predicted output is positive if $g(\mathbf{x}) \geq 0.5$ and negative if $g(\mathbf{x}) < 0.5$.

6.1 Logistic regression

Let us consider binary classification in \mathbb{R}^d , where $\mathcal{X} = \mathbb{R}^{d+1}$ and $\mathcal{Y} = \{0, 1\}$. The basic idea for many classification approaches is to hypothesize a closed-form representation for the posterior probability that the class label is positive and learn parameters \mathbf{w} from data. In logistic regression, this relationship can be expressed as

$$P(Y = 1|\mathbf{x}, \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}}}, \quad (6.1)$$

which is a monotonic function of $\mathbf{w}^\top \mathbf{x}$. Function $\sigma(t) = (1 + e^{-t})^{-1}$ is called the sigmoid function or the logistic function and is plotted in Figure 6.2.

6.1.1 Predicting class labels

For a previously unseen data point \mathbf{x} and a set of coefficients \mathbf{w}^* found from Eq. (6.4) or Eq. (6.11), we simply calculate the posterior probability as

$$P(Y = 1|\mathbf{x}, \mathbf{w}^*) = \frac{1}{1 + e^{-\mathbf{w}^{*T} \cdot \mathbf{x}}}.$$

If $P(Y = 1|\mathbf{x}, \mathbf{w}^*) \geq 0.5$ we conclude that data point \mathbf{x} should be labeled as positive ($\hat{y} = 1$). Otherwise, if $P(Y = 1|\mathbf{x}, \mathbf{w}^*) < 0.5$, we label the data point




Figure 6.2: Sigmoid function in $[-5, 5]$ interval.

as negative ($\hat{y} = 0$). Thus, the predictor maps a $(d + 1)$ -dimensional vector $\mathbf{x} = (x_0 = 1, x_1, \dots, x_d)$ into a zero or one. Note that $P(Y = 1 | \mathbf{x}, \mathbf{w}^*) \geq 0.5$ only when $\mathbf{w}^\top \mathbf{x} \geq 0$. The expression $\mathbf{w}^\top \mathbf{x} = 0$ represents equation of a hyperplane that separates positive and negative examples. Thus, the logistic regression model is a linear classifier.

6.1.2 Maximum conditional likelihood estimation

To frame the learning problem as parameter estimation, we will assume that the data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ is an i.i.d. sample from a fixed but unknown probability distribution $p(\mathbf{x}, y)$. Even more specifically, we will assume that the data generating process randomly draws a data point \mathbf{x} , a realization of the random vector $(X_0 = 1, X_1, \dots, X_d)$, according to $p(\mathbf{x})$ and then sets its class label Y according to the Bernoulli distribution

$$p(y|\mathbf{x}) = \begin{cases} \left(\frac{1}{1+e^{-\boldsymbol{\omega}^\top \mathbf{x}}} \right)^y & \text{for } y = 1 \\ \left(1 - \frac{1}{1+e^{-\boldsymbol{\omega}^\top \mathbf{x}}} \right)^{1-y} & \text{for } y = 0 \end{cases} \quad (6.2)$$

$$= \sigma(\mathbf{x}^\top \boldsymbol{\omega})^y (1 - \sigma(\mathbf{x}^\top \boldsymbol{\omega}))^{1-y}$$

where $\boldsymbol{\omega} = (\omega_0, \omega_1, \dots, \omega_d)$ is a set of unknown coefficients we want to recover (or learn) from the observed data \mathcal{D} . Based on the principles of parameter estimation, we can estimate $\boldsymbol{\omega}$ by maximizing the conditional

likelihood of the observed class labels $\mathbf{y} = (y_1, y_2, \dots, y_n)$ given the inputs $\mathbf{X} = (\mathbf{x}_1^\top, \mathbf{x}_2^\top, \dots, \mathbf{x}_n^\top)$.

As with the previous maximum likelihood estimation, we shall first write the conditional likelihood function $p(\mathbf{y}|\mathbf{X}, \mathbf{w})$, or simply $l(\mathbf{w})$, as

$$l(\mathbf{w}) = \prod_{i=1}^n p(y_i|\mathbf{x}_i, \mathbf{w}). \quad (6.3)$$

This function can be thought of as the probability of observing a set of labels \mathbf{y} given the set of data points \mathbf{X} and the particular set of coefficients \mathbf{w} . The parameter vector that maximizes the likelihood is

$$\begin{aligned} \mathbf{w}_{\text{ML}} &= \arg \max_{\mathbf{w}} \{l(\mathbf{w})\} \\ &= \arg \max_{\mathbf{w}} \left\{ \prod_{i=1}^n p(y_i|\mathbf{x}_i, \mathbf{w}) \right\} \end{aligned} \quad (6.4)$$

$$= \prod_{i=1}^n \left(\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right)^{y_i} \cdot \left(1 - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right)^{1-y_i}. \quad (6.5)$$

As before, because \log is a monotonic function, maximizing Eq. (6.5) is equivalent to maximizing the log-likelihood function $ll(\mathbf{w}) = \log(l(\mathbf{w}))$

$$ll(\mathbf{w}) = \sum_{i=1}^n \left(y_i \cdot \log \left(\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) + (1 - y_i) \cdot \log \left(1 - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \right). \quad (6.6)$$

The negative of the log-likelihood from Eq. (6.6) is sometimes referred to as cross-entropy; thus, cross-entropy minimization is equivalent to the maximum likelihood method. To make everything more suitable for further steps, we will slightly rearrange Eq. (6.6). Notice first that

$$\left(1 - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) = \frac{e^{-\mathbf{w}^\top \mathbf{x}_i}}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}}$$

giving

$$\begin{aligned} ll(\mathbf{w}) &= \sum_{i=1}^n \left(-y_i \cdot \log \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i} \right) + (1 - y_i) \cdot \log \left(e^{-\mathbf{w}^\top \mathbf{x}_i} \right) \right. \\ &\quad \left. - (1 - y_i) \cdot \log \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i} \right) \right) \\ &= \sum_{i=1}^n \left((y_i - 1) \mathbf{w}^\top \mathbf{x}_i + \log \left(\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \right). \end{aligned} \quad (6.7)$$

It does not take much effort to realize that there is no closed-form solution to $\nabla ll(\mathbf{w}) = \mathbf{0}$ (we did have this luxury in linear regression, but not here). Thus, we have to proceed with iterative optimization methods. We will calculate $\nabla ll(\mathbf{w})$ and $H_{ll(\mathbf{w})}$ to enable either method (first-order with only the gradient or second-order with the gradient and Hessian), both written as a function of inputs \mathbf{X} , class labels \mathbf{y} , and the current parameter vector. We can calculate the first and second partial derivatives of $ll(\mathbf{w})$ as follows

$$\begin{aligned}
\frac{\partial ll(\mathbf{w})}{\partial w_j} &= \sum_{i=1}^n \left((y_i - 1) \cdot x_{ij} - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \cdot e^{-\mathbf{w}^\top \mathbf{x}_i} \cdot (-x_{ij}) \right) \\
&= \sum_{i=1}^n x_{ij} \cdot \left(y_i - 1 + \frac{e^{-\mathbf{w}^\top \mathbf{x}_i}}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \\
&= \sum_{i=1}^n x_{ij} \cdot \left(y_i - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \\
&= \mathbf{f}_j^\top (\mathbf{y} - \mathbf{p}),
\end{aligned}$$

where \mathbf{f}_j is the j -th column (feature) of data matrix \mathbf{X} , \mathbf{y} is an n -dimensional column vector of class labels and \mathbf{p} is an n -dimensional column vector of (estimated) posterior probabilities $p_i = P(Y_i = 1 | \mathbf{x}_i, \mathbf{w})$, for $i = 1, \dots, n$. Considering partial derivatives for every component of \mathbf{w} , we have

$$\nabla ll(\mathbf{w}) = \mathbf{X}^\top (\mathbf{y} - \mathbf{p}). \quad (6.8)$$

The second partial derivative of the log-likelihood function can be found as

$$\begin{aligned}
\frac{\partial^2 ll(\mathbf{w})}{\partial w_j \partial w_d} &= \sum_{i=1}^n x_{ij} \cdot \frac{e^{-\mathbf{w}^\top \mathbf{x}_i}}{(1 + e^{-\mathbf{w}^\top \mathbf{x}_i})^2} \cdot (-x_{id}) \\
&= - \sum_{i=1}^n x_{ij} \cdot \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \cdot \left(1 - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \cdot x_{id} \\
&= -\mathbf{f}_j^\top \mathbf{P} (\mathbf{I} - \mathbf{P}) \mathbf{f}_d,
\end{aligned}$$

where \mathbf{P} is an $n \times n$ diagonal matrix with $P_{ii} = p_i = P(Y_i = 1 | \mathbf{x}_i, \mathbf{w})$ and \mathbf{I} is an $n \times n$ identity matrix. The Hessian matrix $H_{ll(\mathbf{w})}$ can now be calculated as

$$H_{ll(\mathbf{w})} = -\mathbf{X}^\top \mathbf{P} (\mathbf{I} - \mathbf{P}) \mathbf{X}. \quad (6.9)$$

This is a negative semi-definite matrix, which guarantees that the optimum we find will be a global maximum. Negative semi-definite Hessian corresponds to a concave likelihood function and has a global maximum (unique

if negative definite). Substituting Eqs. (6.8-6.9) into Newton-Raphson's method results in the following weight update rule

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \left(\mathbf{X}^\top \mathbf{P}^{(t)} (\mathbf{I} - \mathbf{P}^{(t)}) \mathbf{X} \right)^{-1} \mathbf{X}^\top (\mathbf{y} - \mathbf{p}^{(t)}), \quad (6.10)$$

where the initial weights $\mathbf{w}^{(0)}$ can be calculated using the ordinary least squares regression as $\mathbf{w}^{(0)} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$. Note that the second term on the right-hand side of Eq. (6.10) is calculated in each iteration t ; thus we wrote \mathbf{P} and \mathbf{p} as $\mathbf{P}^{(t)}$ and $\mathbf{p}^{(t)}$, respectively, to indicate that $\mathbf{w}^{(t)}$ was used to calculate them.

The computational complexity of this procedure is $O(d^3 + d^2n)$ in each iteration, assuming $O(d^3)$ time for finding matrix inverses. If this is infeasible for the given dataset, such as if the number of features d is large, we can use a first-order gradient descent update (which only requires $O(dn)$ to compute the matrix vector product) or quasi-Newton methods that approximate the Hessian (e.g., LBFGS).

Weighted conditional likelihood function

In certain situations, it may be justified to allow for unequal importance of each data point. This modifies the conditional likelihood function from Eq. (6.5) to

$$l(\mathbf{w}) = \prod_{i=1}^n p_i^{c_i y_i} \cdot (1 - p_i)^{c_i (1 - y_i)},$$

where $0 \leq c_i \leq 1$ is a cost for data point i . Taking that $\mathbf{C} = \text{diag}(c_1, c_2, \dots, c_n)$ we can now express the gradient of the log-likelihood as

$$\nabla l(\mathbf{w}) = \mathbf{X}^\top \mathbf{C} (\mathbf{y} - \mathbf{p})$$

and the Hessian as

$$H_{ll(\mathbf{w})} = -\mathbf{X}^\top \mathbf{C} \mathbf{P} (\mathbf{I} - \mathbf{P}) \mathbf{X}.$$

It is interesting to observe that the Hessian remains negative semi-definite. Thus, the update rule is expected to converge to a global maximum.

6.1.3 Stochastic optimization for logistic regression

We have derived that the weight update rule depends on the data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ and the predictions on all training examples using the weight

vector from the current step. As in Algorithm 2, we can also use stochastic gradient descent to optimize the logistic regression model. We will first rewrite the update rule of the (first-order) gradient descent method as

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \Delta \mathbf{w}^{(t)},$$

where

$$\begin{aligned} \Delta \mathbf{w}^{(t)} &= \eta \mathbf{X}^\top (\mathbf{y} - \mathbf{p}^{(t)}) \\ &= \eta \sum_{i=1}^n \mathbf{x}_i (y_i - p_i). \end{aligned}$$

We can see that the weight update is simply a linear combination of training data points. This batch gradient step is an unbiased estimate of the gradient of the true function $\int_{\mathcal{X}} \sum_{y \in \mathcal{Y}} p(\mathbf{x}, y) \mathbf{x} (y - \sigma(\mathbf{x}^\top \mathbf{w})) d\mathbf{x} dy$, because we have i.i.d. samples and so this sum converges to this integral over all possible values. Similarly, using one sample to compute the gradient is also an unbiased estimate and has been shown to converge under certain step-size conditions. When the training data is large, it may be beneficial to update the weight vector after each data point is presented to the learning algorithm, rather than computing the batch update for all samples. That is, if data point \mathbf{x}_t with its class label y_t is presented to the learner, the stochastic gradient descent weight update is

$$\Delta \mathbf{w}^{(t)} = \eta \mathbf{x}_t (y_t - p_t).$$

This stochastic gradient descent is a form of stochastic approximation, which comes from a wider class of methods in the area of stochastic optimization. The batch method is often also called sample average approximation.

The training algorithm can now be revised to randomly draw one data point at a time from \mathcal{D} and then update the current weights using the previous equation. Typically, in practice, this entails iterating one or more times over the dataset in order (assuming it is random, with i.i.d. samples). The conditions for convergence typically include conditions on the step-sizes, requiring them to decrease over time. As with batch gradient descent, these stochastic gradient descent updates will converge, though with more oscillation around the true weight vector, with the decreasing step-size progressively smoothing out these oscillations.

6.1.4 Issues with minimizing Euclidean distance

A natural question is why we went down this route for linear classification. Instead of explicitly assuming $P(Y = 1 | \mathbf{x}, \mathbf{w})$ is a Bernoulli distribution and

computing the maximum likelihood solution for $\sigma(\mathbf{x}^\top \mathbf{w}) = E[Y|\mathbf{x}] = P(Y = 1|\mathbf{x}, \mathbf{w})$, we could have simply decided to use $\sigma(\mathbf{x}^\top \mathbf{w})$ to predict targets $y \in \{0, 1\}$ and then tried to minimize their difference, using our favorite loss (the squared loss). Unfortunately, this more haphazard problem specification results in a non-convex optimization. In fact, there is a result that using the Euclidean error for the sigmoid transfer gives exponentially many local minima in the number of features [1]. For interest about this alternate route, we will show that this direction leads to a non-convex optimization.

Let the error function with Euclidean distance now be written as

$$\text{Err}(\mathbf{w}) = \sum_{i=1}^n (y_i - p_i)^2 \quad \triangleright p_i = \sigma(\mathbf{x}_i^\top \mathbf{w}), \quad e_i = y_i - p_i$$

The minimization of $\text{Err}(\mathbf{w})$ is formally expressed as

$$\begin{aligned} \mathbf{w}^* &= \arg \min_{\mathbf{w}} \{\text{Err}(\mathbf{w})\} \\ &= \arg \min_{\mathbf{w}} \left\{ \sum_{i=1}^n (y_i - p_i)^2 \right\}. \end{aligned} \quad (6.11)$$

Similar to the maximum likelihood process, our goal will be to calculate the gradient vector and the Hessian of the error function. The partial derivatives of the error function can be calculated as follows

$$\begin{aligned} \frac{\partial \text{Err}(\mathbf{w})}{\partial w_j} &= \frac{\partial}{\partial w_j} \sum_{i=1}^n e_i^2 \\ &= \sum_{i=1}^n 2 \cdot e_i \cdot \frac{\partial e_i}{\partial w_j} \\ &= 2 \cdot \sum_{i=1}^n \left(y_i - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \cdot \frac{1}{(1 + e^{-\mathbf{w}^\top \mathbf{x}_i})^2} \cdot e^{-\mathbf{w}^\top \mathbf{x}_i} \cdot (-x_{ij}) \\ &= -2 \cdot \sum_{i=1}^n x_{ij} \cdot \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \cdot \left(1 - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \cdot \left(y_i - \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} \right) \\ &= -2 \mathbf{f}_j^\top \mathbf{P} (\mathbf{I} - \mathbf{P}) (\mathbf{y} - \mathbf{p}). \end{aligned}$$

This provides the gradient vector in the following form

$$\nabla \text{Err}(\mathbf{w}) = -2 \mathbf{X}^\top \mathbf{P} (\mathbf{I} - \mathbf{P}) (\mathbf{y} - \mathbf{p}).$$

Matrix $\mathbf{J} = \mathbf{P}(\mathbf{I} - \mathbf{P})\mathbf{X}$ is referred to as Jacobian. In general, Jacobian is an $n \times d$ matrix calculated as

$$J_{\text{Err}(\mathbf{w})} = \begin{bmatrix} \frac{\partial e_1}{\partial w_1} & \frac{\partial e_1}{\partial w_2} & \cdots & \frac{\partial e_1}{\partial w_d} \\ \vdots & \ddots & & \\ \frac{\partial e_n}{\partial w_1} & & & \frac{\partial e_n}{\partial w_d} \end{bmatrix}.$$

The second partial derivative of the error function can be found as

$$\begin{aligned} \frac{\partial^2 \text{Err}(\mathbf{w})}{\partial w_j \partial w_d} &= 2 \cdot \sum_{i=1}^n \frac{\partial e_i}{\partial w_d} \cdot \frac{\partial e_i}{\partial w_j} + e_i \cdot \frac{\partial^2 e_i}{\partial w_j \partial w_d} \\ &= 2 \cdot \sum_{i=1}^n x_{ij} \cdot (p_i^2(1 - p_i)^2 + p_i \cdot (1 - p_i) \cdot (2p_i - 1) \cdot (y_i - p_i)) \cdot x_{ik}. \end{aligned}$$

Thus, the Hessian can be computed as

$$\begin{aligned} H_{\text{Err}(\mathbf{w})} &= 2\mathbf{X}^\top (\mathbf{I} - \mathbf{P})^\top \mathbf{P}^\top \mathbf{P} (\mathbf{I} - \mathbf{P}) \mathbf{X} + 2\mathbf{X}^\top (\mathbf{I} - \mathbf{P})^\top \mathbf{P}^\top \mathbf{E} (2\mathbf{P} - \mathbf{I}) \mathbf{X} \\ &= 2\mathbf{J}^\top \mathbf{J} + 2\mathbf{J}^\top \mathbf{E} (2\mathbf{P} - \mathbf{I}) \mathbf{X}, \end{aligned}$$

where $\mathbf{P} = \text{diag}\{\mathbf{p}\}$, $\mathbf{E} = \text{diag}\{\mathbf{e}\}$ is a diagonal matrix containing elements $E_{ii} = e_i = y_i - p_i$ and \mathbf{I} is an identity matrix.

We can now see that the Hessian is not guaranteed to be positive semi-definite. This means that $\text{Err}(\mathbf{w})$ is not convex, i.e. it must have multiple minima with different values of the objective function. Finding a global optimum depends on how favorable the initial solution $\mathbf{w}^{(0)}$ is and how well the weight update step can escape local minima to find better ones. Minimization of this non-convex function, however, will be much more problematic than the convex cross-entropy.

6.2 Naive Bayes Classifier

Naive Bayes classification is a generative approach to prediction. So far, we have discussed discriminative approaches (linear regression, logistic regression), which attempt to learn $p(y|\mathbf{x})$. For a generative setting, we learn $p(\mathbf{x}, y) = p(\mathbf{x}|y)p(y)$. As you can imagine, this can be a more difficult undertaking, as we also need to learn the distribution over the features themselves. For naive Bayes, we significantly simplify learning this joint distribution by making a strong assumption: the features are conditionally independent




Figure 6.3: Naive Bayes graphical model, with three features.

given the label. This assumption is demonstrated by the graphical model in Figure 6.3.

As with logistic regression, even though we learn a generative model, the decision rule for labeling a point as class 1 (i.e., $y = 1$) is

$$p(y = 1|\mathbf{x}) \geq p(y = 0|\mathbf{x})$$

where now $p(y = 1|\mathbf{x}) = p(\mathbf{x}|y = 1)p(y = 1)$. To start, we will assume a simpler setting with binary features, and then address continuous features. Note that naive Bayes is a linear classifier for binary features; more generally, however, it is not necessarily a linear classifier. Note that a linear classifier is one in which the two classes are separated by a linear plane, i.e., the decision boundary is according to some linear combination of features.

6.2.1 Binary features and linear classification

Let $\mathcal{D} = \{(\mathbf{x}_i, y)\}_{i=1}^n$ be an input data set, where $\mathcal{X} = \{0, 1\}^k$ and $\mathcal{Y} = \{0, 1\}$. Under the naive Bayes assumption, the features are independent given the class label. Therefore, we can write

$$p(\mathbf{x}|y) = \prod_{i=1}^d p(x_i|y).$$

A suitable choice for this simpler univariate probabilities is a Bernoulli distribution, since each x_j is binary, giving

$$p(x_j|y = c) = p_{j,c}^{x_j} (1 - p_{j,c})^{1-x_j}$$

The parameters for the Bernoulli distributions are $p_{j,c} = p(x_j = 1|y = c)$, with a different parameter $p_{j,c}$ for each class value c and for each feature.

We can easily learn this parameter from data by calculating

$$p_{j,c} = \frac{\text{number of times } x_j = 1 \text{ for class } c}{\text{number of datapoints labeled as class } c}.$$

Similarly, we can learn the prior $p_c = p(y = c)$ using

$$p_c = p(y = c) = \frac{\text{number of datapoint labeled as class } c}{\text{total number of datapoints}}.$$

Notice that this approach could also be accomplished for more classes than just two. The prediction on a new point \mathbf{x} is then

$$\begin{aligned} \max_{c \in \mathcal{Y}} p(y = c | \mathbf{x}) &= \max_{c \in \mathcal{Y}} p(\mathbf{x} | y = c) p(y = c) \\ &= \max_{c \in \mathcal{Y}} \prod_{j=1}^d p(x_j | y = c) p(y = c) \\ &= \max_{c \in \mathcal{Y}} \prod_{j=1}^d p_{j,c} p_c \end{aligned}$$

Exercise: The solution above is intuitive, and comes from similarly deriving the maximum likelihood solution. Assuming that you have n datapoints and the chosen distribution $p(x_i | y)$ is Bernoulli as described above, derive the maximum likelihood parameters $p_{j,c}, p_c$ for $i = 1, \dots, d, c = 0, 1$. To make things simpler, use the log of the likelihood.

Linear classification boundary with binary features and targets

Interestingly, naive Bayes classifier with binary features and two classes is a linear classifier. This is somewhat surprising, as this generative approach looks very different from what we did before. To see why this is the case, notice that the classifier will make a positive decision when

$$p(y = 1 | \mathbf{x}) \geq p(y = 0 | \mathbf{x})$$

that is, when

$$p(\mathbf{x} | y = 1) p(y = 1) \geq p(\mathbf{x} | y = 0) p(y = 0)$$

We will shorten this notation using $p(y = 0) = p(0)$, $p(\mathbf{x} | y = 0) = p(\mathbf{x} | 0)$, etc. Using the naive Bayes assumption, we now have

$$p(1) \prod_{j=1}^k p(x_j | 1) \geq p(0) \prod_{j=1}^k p(x_j | 0)$$

which, after applying a logarithm, then becomes

$$\log p(1) + \sum_{j=1}^k \log p(x_j|1) \geq \log p(0) + \sum_{j=1}^k \log p(x_j|0)$$

Let us now investigate class-conditional probabilities $p(x_j|y)$, when $y \in \{0, 1\}$. Recall that each feature is Bernoulli distributed, i.e.

$$p(x_j|1) = p_{j,1}^{x_j} (1 - p_{j,1})^{1-x_j}$$

and

$$p(x_j|0) = p_{j,0}^{x_j} (1 - p_{j,0})^{1-x_j}$$

where parameters $p_{j,c}$ are estimated from the training set. Taking $p(y = c) = p_c$, we have

$$\sum_{j=1}^k x_j \log \frac{p_{j,1}(1 - p_{j,0})}{(1 - p_{j,1})p_{j,0}} + \sum_{j=1}^k \log \frac{1 - p_{j,1}}{1 - p_{j,0}} + \log \frac{p_1}{p_0} \geq 0$$

We can write the previous expression as

$$w_0 + \sum_{j=1}^k w_j x_j \geq 0$$

where

$$w_0 = \log \frac{p_1}{p_0} + \sum_{j=1}^k \log \frac{1 - p_{j,1}}{1 - p_{j,0}}$$

$$w_j = \log \frac{p_{j,1}(1 - p_{j,0})}{(1 - p_{j,1})p_{j,0}} \quad j \in \{1, 2, \dots, k\}$$

Therefore, in the case of binary features, naive Bayes is a linear classifier.

6.2.2 Continuous naive Bayes

For continuous features, a Bernoulli distribution is no longer appropriate for $p(x_j|y)$ and we need to choose a different conditional distribution $p(\mathbf{x}|y)$. A common choice is a Gaussian distribution, now with a different mean and variance for each feature and class, $\mu_{j,c}, \sigma_{j,c}^2$:

$$p(x_j|y = c) = (2\pi\sigma_{j,c}^2)^{-1/2} \exp \left(-\frac{(x_j - \mu_{j,c})^2}{2\sigma_{j,c}^2} \right).$$

Since y is still discrete, we can approximate $p(y)$ using counts as before. The maximum likelihood mean and variance parameters correspond to the sample mean and sample covariance for each given class separately. This involves computing the mean and variance of feature j across the datapoints labeled with class c :

$$\begin{aligned}\mu_{j,c} &= \frac{\sum_{i=1}^n 1(y_i = c) x_j}{\text{number of datapoints labeled as class } c} \\ \sigma_{j,c}^2 &= \frac{\sum_{i=1}^n 1(y_i = c) (x_j - \mu_{j,c})^2}{\text{number of datapoints labeled as class } c} \\ &= \frac{\sum_{i=1}^n 1(y_i = c) (x_j - \mu_{j,c})^2}{\text{number of datapoints labeled as class } c}\end{aligned}$$

Exercise: Derive the maximum likelihood formulation for a Gaussian naive Bayes model, and check that the solution does in fact match the sample mean and variance for each feature and class separately, as above.

6.3 Multinomial logistic regression

Now let us consider discriminative multiclass classification, where $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{Y} = \{1, 2, \dots, k\}$. This setting arises naturally in machine learning, where there is often more than two categories. For example, if we want to predict the blood type (A, B, AB and O) of an individual, then we have four classes. Here we discuss multiclass classification where we only want to label a datapoint with one class out of k . In other settings, one might want to label a datapoint with multiple classes; this is briefly mentioned at the end of this section.

We can nicely generalize to this setting using the idea of multinomials and the corresponding link function, as with the other generalized linear models. The multinomial distribution is a member of the exponential family. We can write

$$p(\mathbf{y}|\mathbf{x}) = \frac{1}{y_1! \dots y_k!} p(y_1 = 1|\mathbf{x})^{y_1} \dots p(y_k|\mathbf{x})^{y_k} \quad (6.12)$$

where the usual numerator $n!$ is 1 because $n = \sum_{j=1}^k y_j = 1$ since we can only have one class value. As with logistic regression, we can parametrize $p(y_j = 1|\mathbf{x}) = \sigma(\mathbf{x}^\top \mathbf{w}_j)$. However, we must also ensure that $\sum_{j=1}^k p(y_j = 1|\mathbf{x}) = 1$. To do so, we “pivot” around the final class, $p(y_k = 1|\mathbf{x}) = 1 - \sum_{j=1}^{k-1} p(y_j = 1|\mathbf{x})$ and only explicitly learn $\mathbf{w}_1, \dots, \mathbf{w}_{k-1}$. Note that these models are not

learned independently, because they are tied by the probability for the last class. The parameters can be represented as a matrix $\mathbf{W} \in \mathbb{R}^{d \times k}$ where $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_k]$ is composed of k weight vectors with $\mathbf{w}_k = \mathbf{0}$. We will see why we fix $\mathbf{w}_k = \mathbf{0}$.

The transfer (inverse of the link) for this setting is the softmax transfer

$$\begin{aligned} \text{softmax}(\mathbf{x}^\top \mathbf{W}) &= \left[\frac{\exp(\mathbf{x}^\top \mathbf{w}_1)}{\sum_{j=1}^k \exp(\mathbf{x}^\top \mathbf{w}_j)}, \dots, \frac{\exp(\mathbf{x}^\top \mathbf{w}_k)}{\sum_{j=1}^k \exp(\mathbf{x}^\top \mathbf{w}_j)} \right] \\ &= \left[\frac{\exp(\mathbf{x}^\top \mathbf{w}_1)}{\mathbf{1}^\top \exp(\mathbf{x}^\top \mathbf{W})}, \dots, \frac{\exp(\mathbf{x}^\top \mathbf{w}_k)}{\mathbf{1}^\top \exp(\mathbf{x}^\top \mathbf{W})} \right] \end{aligned}$$

and the prediction is $\text{softmax}(\mathbf{x}) = \hat{\mathbf{y}} \in [0, 1]^k$, which gives the probability in each entry of being labeled as that class, where $\hat{\mathbf{y}}^\top \mathbf{1} = 1$ signifying that the probabilities sum to 1. Note that this model encompasses the binary setting for logistic regression, because $\sigma(\mathbf{x}^\top \mathbf{w}) = (1 + \exp(-\mathbf{x}^\top \mathbf{w}))^{-1} = \frac{\exp(\mathbf{x}^\top \mathbf{w})}{1 + \exp(\mathbf{x}^\top \mathbf{w})}$. The weights for multinomial logistic regression with two classes are then $\mathbf{W} = [\mathbf{w}, \mathbf{0}]$ giving

$$\begin{aligned} p(y = 1|\mathbf{x}) &= \frac{\exp(\mathbf{x}^\top \mathbf{w})}{\mathbf{1}^\top \exp(\mathbf{x}^\top \mathbf{W})} \\ &= \frac{\exp(\mathbf{x}^\top \mathbf{w})}{\exp(\mathbf{x}^\top \mathbf{W}) + \exp(\mathbf{x}^\top \mathbf{0})} \\ &= \frac{\exp(\mathbf{x}^\top \mathbf{w})}{\exp(\mathbf{x}^\top \mathbf{W}) + 1} \\ &= \sigma(\mathbf{x}^\top \mathbf{w}). \end{aligned}$$

Similarly, for $k > 2$, by fixing $\mathbf{w}_k = \mathbf{0}$, the other weights $\mathbf{w}_1, \dots, \mathbf{w}_{k-1}$ are learned to ensure that $p(y = k|\mathbf{x}) = \frac{\exp(\mathbf{x}^\top \mathbf{w}_k)}{\mathbf{1}^\top \exp(\mathbf{x}^\top \mathbf{W})} = \frac{1}{1 + \sum_{j=1}^{k-1} \exp(\mathbf{x}^\top \mathbf{w}_j)}$ and that $\sum_{j=1}^k p(y = j|\mathbf{x}) = 1$.

With the parameters of the model parameterized by \mathbf{W} and the softmax transfer, we can determine the maximum likelihood formulation. By plugging in the parameterization into Equation (6.12), taking the negative log of that likelihood and dropping constants, we arrive at the following loss

$$\min_{\mathbf{W} \in \mathbb{R}^{d \times k} : \mathbf{W}_{:,k} = \mathbf{0}} \log \left(\mathbf{1}^\top \exp(\mathbf{x}^\top \mathbf{W}) \right) - \mathbf{x}^\top \mathbf{W} \mathbf{y}$$

with gradient

$$\nabla \left(\log \left(\mathbf{1}^\top \exp(\mathbf{x}^\top \mathbf{W}) \right) - \mathbf{x}^\top \mathbf{W} \mathbf{y} \right) = \frac{\exp(\mathbf{x}^\top \mathbf{W})^\top \mathbf{x}^\top}{\mathbf{1}^\top \exp(\mathbf{x}^\top \mathbf{W})} - \mathbf{x} \mathbf{y}^\top.$$

As before, we do not have a closed form solution for this gradient, and will use iterative methods to solve for \mathbf{W} . Note that here, unlike previous methods, we have a constraint on part of the variable. However, this was solely written this way for convenience. We do not optimize $\mathbf{W}_{:k}$, as it is fixed at zero; one can rewrite this minimization and gradient to only apply to the $W_{:(1:k-1)}$. This corresponds to initializing $\mathbf{W}_{:k} = \mathbf{0}$, and then only using the first $k - 1$ columns of the gradient in the update to $W_{:(1:k-1)}$.

The final prediction $\text{softmax}(\mathbf{x}^\top \mathbf{W}) \in [0, 1]$ gives the probabilities of being in a class. As with logistic regression, to pick one class, the highest probability value is chosen. For example, with $k = 4$, we might predict $[0.10.20.60.1]$ and so decide to classify the point into class 3.

Remark about overlapping classes: If you want to predict multiple classes for a datapoint \mathbf{x} , then a common strategy is to learn separate binary predictors for each class. Each predictor is queried separately, and a datapoint will label each class as 0 or 1, with potentially more than one class having a 1. Above, we examined the case where the datapoint was exclusively in one of the provided classes, by setting $n = 1$ in the multinomial.

Chapter 7

Representations for machine learning

At first, it might seem that the applicability of linear regression and classification to real-life problems is greatly limited. After all, it is not clear whether it is realistic (most of the time) to assume that the target variable is a linear combination of features. Fortunately, the applicability of linear regression is broader than originally thought. The main idea is to apply a non-linear transformation to the data matrix \mathbf{x} prior to the fitting step, which then enables a non-linear fit. Obtaining such a useful feature representation is a central problem in machine learning.

We will first examine fixed representations for linear regression: polynomial curve fitting and radial basis function (RBF) networks. Then, we will discuss learning representations.

7.1 Radial basis function networks and kernel representations

The idea of radial basis function (RBF) networks is a natural generalization of the polynomial curve fitting and approaches from the previous Section. Given data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, we start by picking p points to serve as the “centers” in the input space \mathcal{X} . We denote those centers as $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_p$. Usually, these can be selected from \mathcal{D} or computed using some clustering technique (e.g. the EM algorithm, K-means).

When the clusters are determined using a Gaussian mixture model, the basis functions can be selected as

$$\phi_j(\mathbf{x}) = e^{-\frac{1}{2}(\mathbf{x}-\mathbf{c}_j)^T \Sigma_j^{-1}(\mathbf{x}-\mathbf{c}_j)},$$

where the cluster centers and the covariance matrix are found during clustering. When K-means or other clustering is used, we can use

$$\phi_j(\mathbf{x}) = e^{-\frac{\|\mathbf{x}-\mathbf{c}_j\|^2}{2\sigma_j^2}},$$




Figure 7.1: Radial basis function network.

where σ_j 's can be separately optimized; e.g. using a validation set. In the context of multidimensional transformations from \mathbf{x} to Φ , the basis functions can also be referred to as *kernel functions*, i.e. $\phi_j(\mathbf{x}) = k_j(\mathbf{x}, \mathbf{c}_j)$. Matrix

$$\Phi = \begin{bmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_p(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & & \\ \vdots & & \ddots & \\ \phi_0(\mathbf{x}_n) & & & \phi_p(\mathbf{x}_n) \end{bmatrix}$$

is now used as a new data matrix. For a given input \mathbf{x} , the prediction of the target y will be calculated as

$$\begin{aligned} f(\mathbf{x}) &= w_0 + \sum_{j=1}^p w_j \phi_j(\mathbf{x}) \\ &= \sum_{j=0}^p w_j \phi_j(\mathbf{x}) \end{aligned}$$

where $\phi_0(\mathbf{x}) = 1$ and \mathbf{w} is to be found. It can be proved that with a sufficiently large number of radial basis functions we can accurately approximate any function. As seen in Figure 7.1, we can think of RBFs as neural networks.

RBF networks and kernel representations are highly related. The main distinction is that kernel representations use any kernel function for the similarity measure $k(\mathbf{x}, \mathbf{c}_j) = \phi_j(\mathbf{x})$, where radial basis functions are one example

of a kernel. In addition, if an RBF kernel is chosen, for kernel representations typical the centers are selected from the training dataset. For RBF networks, the selection of the centers is left generally as an important step, where they can be selected from the training set but can also be selected in other ways.

7.2 Learning representations

There are many approaches to learning representations. Two dominant approaches are (semi-supervised) matrix factorization techniques and neural networks. Neural networks build on the generalized linear models we have discussed, stacking multiple generalized linear models together. Matrix factorization techniques (e.g., dimensionality reduction, sparse coding) typically factorize the input data into a dictionary and a new representation (a basis). We will first discuss neural networks, and then discuss the many unsupervised and semisupervised learning techniques that are encompassed by matrix factorizations.

7.2.1 Neural networks

Neural networks are a form of supervised representation learning. All the first hidden layers constitute representation layer; the output learning is the supervised prediction part. Neural networks stack multiple generalized linear models. Figure 7.2 shows the graphical model for the generalized linear models we discussed in the previous chapters, where the weights and corresponding transfer can be thought of as being on the arrows (as they are not random variables). Figure 7.3 shows a neural network with one hidden-layer; this is called a two-layer neural network, as there are two layers of weights.

The logistic regression solution specified a linear classifier; by stacking multiple logistic regression models, you obtain a more advance non-linear representation. For now, to compare to logistic regression, assume that your transfer function f is the sigmoid function σ . For logistic regression we learned $\mathbf{W} \in \mathbb{R}^{d \times k}$, with $f(\mathbf{x}\mathbf{W}) = \sigma(\mathbf{x}\mathbf{W}) \approx \mathbf{y}$. When we add a hidden layer, we have two parameter matrices $\mathbf{W}^{(1)} \in \mathbb{R}^{d \times k}$ and $\mathbf{W}^{(2)} \in \mathbb{R}^{k \times k}$,




Figure 7.2: Generalized linear model, such as logistic regression.

Figure 7.3: Standard two-layer neural network.

where k is the dimension of the hidden layer

$$\mathbf{h} = \sigma(\mathbf{W}^{(1)}\mathbf{x}) = \begin{bmatrix} \sigma(\mathbf{W}^{(1)}_{1:\mathbf{x}}) \\ \sigma(\mathbf{W}^{(1)}_{2:\mathbf{x}}) \\ \vdots \\ \sigma(\mathbf{W}^{(1)}_{k:\mathbf{x}}) \end{bmatrix} \in \mathbb{R}^k$$

where the sigmoid function is applied to each entry in $\mathbf{W}^{(1)}\mathbf{x}$ and $\mathbf{W}^{(2)}\mathbf{h}$ (i.e., the division is element-wise). This hidden layer is your new set of features and again you will do the regular logistic regression optimization to learn weights on \mathbf{h} :

$$p(y = 1|\mathbf{x}) = \sigma(\mathbf{W}^{(2)}\mathbf{h}) = \sigma(\mathbf{W}^{(2)}\sigma(\mathbf{W}^{(1)}\mathbf{x})).$$

This composition of transfers seems to complicate matters, but we can still take the gradient w.r.t. our parameters. We simply have more parameters now: $\mathbf{W}^{(1)} \in \mathbb{R}^{k \times d}$, $\mathbf{W}^{(2)} \in \mathbb{R}^{1 \times k}$.

First, we take the gradient of the output loss (the cross-entropy) with respect to $\mathbf{W}^{(2)}$. Then, we take the gradient with respect to $\mathbf{W}^{(1)}$. To do so, we can recall that the chain rule for a function $f : \mathbb{R}^k \rightarrow \mathbb{R}$ and $g : \mathbb{R}^d \rightarrow \mathbb{R}^k$ for parameters $\mathbf{W}^{(1)} \in \mathbb{R}^{k \times d}$ and $\mathbf{h} = g(\mathbf{W}^{(1)})$ is

$$\frac{\partial f(\mathbf{h})}{\partial \mathbf{W}^{(1)}_{ij}} = \text{tr} \left(\left(\frac{\partial f(\mathbf{h})}{\partial \mathbf{h}} \right)^T \frac{\partial \mathbf{h}}{\partial \mathbf{W}^{(1)}_{ij}} \right).$$

(See the matrix cookbook for more information). Due to the size of the network, we will often learn with stochastic gradient descent. Therefore, we will first compute this gradient assuming we only have one sample. Once we have the gradient w.r.t. each parameter matrix, we simply take a step in the direction of the negative of the gradient, as usual. This algorithm is typically called *back propagation*.

Note that in general, we can compute the gradient for any number H of hidden layers. For simplicity of notation, denote each differentiable transfer function f_1, \dots, f_H , where f_H is the sigmoid and produces the final probability:

$$p(y = 1|\mathbf{x}) = f_H(\mathbf{W}^{(H)} f_{H-1}(\mathbf{W}^{(H-1)} \dots f_1(\mathbf{W}^{(1)} \mathbf{x}) \underbrace{\dots}_{H \text{ of these brackets}})$$

where $\mathbf{W}^{(1)} \in \mathbb{R}^{k_1 \times d}, \dots, \mathbf{W}^{(H-1)} \in \mathbb{R}^{k_{H-1} \times k_{H-2}}$ and $\mathbf{W}^{(H)} \in \mathbb{R}^{1 \times k_{H-1}}$.

Backpropagation algorithm

We will start by deriving back propagation for two layers; the extension to multiple layers will be more clear given this derivation. The back-propagation algorithm is simply gradient descent on a non-convex objective, with a careful ordering of computation to avoid repeating computation. In particular, one first propagates forward and computes variable $\mathbf{h} = f_1(\mathbf{W}^{(1)} \mathbf{x}_i)$ and then $\hat{y} = f_2(\mathbf{W}^{(2)} f_1(\mathbf{W}^{(1)} \mathbf{x}_i)) = f_2(\mathbf{W}^{(2)} \mathbf{h})$. We then compute the error between our prediction \hat{y} and the true label. We take the gradient of this error (loss) w.r.t. to our parameters; in this case, for efficient computation, the best ordering is to compute the gradient w.r.t. to the last parameter $\mathbf{W}^{(2)}$ first, and then $\mathbf{W}^{(1)}$. This is the reason for the term back-propagation, since the error is propagated backward from the last layer first.

The choices then involve picking the transfers at each layer, the number of hidden nodes and the loss for the last layer. The matching convex loss $L(\cdot, y)$ depends on the chosen $p(y|\mathbf{x})$ and corresponding transfer function for the last layer of the neural network; we will generically call this the error function $\text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)}) = L(f_2(\mathbf{W}^{(2)} f_1(\mathbf{W}^{(1)} \mathbf{x}_i)), y_i)$, for one sample (\mathbf{x}_i, y_i) . For example, for $p(y = 1|\mathbf{x})$ Gaussian and identity transfer f_2 , we get $\text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)}) = (\mathbf{W}^{(2)} f_1(\mathbf{W}^{(1)} \mathbf{x}_i) - y_i)^2$. If $p(y = 1|\mathbf{x})$ is a Bernoulli distribution, then we would chose the logistic regression loss (the cross entropy).

As before, we will compute gradients of the loss w.r.t. our parameters. First, we take the partial derivative w.r.t. the parameters $\mathbf{W}^{(2)}$ (assuming

$\mathbf{W}^{(1)}$ are fixed).

$$\begin{aligned}
\frac{\partial \text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial \mathbf{W}^{(2)}} &= \frac{\partial L(f_2(\mathbf{W}^{(2)} f_1(\mathbf{W}^{(1)} \mathbf{x}_i)), y_i)}{\partial \mathbf{W}^{(2)}} \\
&= \left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) \frac{\partial \hat{y}}{\partial \mathbf{W}^{(2)}} \quad \triangleright \hat{y} = f_2(\mathbf{W}^{(2)} \mathbf{h}) \\
&= \left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) \frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}} \frac{\partial \boldsymbol{\theta}^{(2)}}{\partial \mathbf{W}^{(2)}} \quad \triangleright \boldsymbol{\theta}^{(2)} = \mathbf{W}^{(2)} \mathbf{h} \\
&= \left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) \frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}} \mathbf{h}^\top
\end{aligned}$$

At this point these equations are abstract; but they are simple to compute for the losses and transfers we have examined. For example, for $L(\hat{y}, y_i) = \frac{1}{2}(\hat{y} - y_i)^2$, and f_2 the identity, we get

$$\begin{aligned}
\left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) &= (\hat{y} - y_i) \\
\frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}} &= 1
\end{aligned}$$

giving

$$\frac{\partial \text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial \mathbf{W}^{(2)}} = \left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) \frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}} \mathbf{h}^\top = (\hat{y} - y_i) \mathbf{h}^\top.$$

The gradient update then looks as usual like $\mathbf{W}^{(2)} = \mathbf{W}^{(2)} - \alpha(\hat{y} - y_i) \mathbf{h}^\top$ for some step-size α . Below we see what this looks like for the sigmoid transfer.

Next, to compute the partial gradient with respect to $\mathbf{W}^{(1)}$, we use the above computed gradient according to the chain rule and then further compute the gradient of \mathbf{h} with respect to $\mathbf{W}^{(1)}$. For $\mathbf{W}^{(2)}$, the $\mathbf{h} = f_1(\mathbf{W}^{(1)} \mathbf{x}_i)$ is a constant that does not depend on $\mathbf{W}^{(2)}$; however, it does depend on $\mathbf{W}^{(1)}$, so we continue applying the chain rule from here for that variable. Note that if we had started computing the partial derivative w.r.t. $\mathbf{W}^{(1)}$ from the loss, we would reach the exact same point as above and then would next to take the gradient of \mathbf{h} w.r.t. $\mathbf{W}^{(1)}$. Therefore, it makes more sense to first compute the gradient w.r.t. $\mathbf{W}^{(2)}$, and use this information to then

compute the gradient for $\mathbf{W}^{(1)}$.

$$\begin{aligned}
\frac{\partial \text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial \mathbf{W}^{(1)}} &= \frac{\partial L(f_2(\mathbf{W}^{(2)} f_1(\mathbf{W}^{(1)} \mathbf{x}_i)), y_i)}{\partial \mathbf{W}^{(1)}} \\
&= \left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) \frac{\partial \hat{y}}{\partial \mathbf{W}^{(1)}} \quad \triangleright \hat{y} = f_2(\mathbf{W}^{(2)} \mathbf{h}) \\
&= \left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) \frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}} \frac{\partial \boldsymbol{\theta}^{(2)}}{\partial \mathbf{W}^{(1)}} \quad \triangleright \boldsymbol{\theta}^{(2)} = \mathbf{W}^{(2)} \mathbf{h} \\
&= \left(\frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \right) \frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}} \frac{\partial \mathbf{W}^{(2)} \mathbf{h}}{\partial \mathbf{W}^{(1)}}.
\end{aligned}$$

Now $\mathbf{W}^{(2)} \mathbf{h}$ is still a scalar, and the partial derivative w.r.t. $\mathbf{W}^{(1)}$ should be in $\mathbb{R}^{d \times k}$. We can compute this gradient as

$$\begin{aligned}
\frac{\partial \mathbf{W}^{(2)} \mathbf{h}}{\partial \mathbf{W}^{(1)}} &= \frac{\partial \sum_{j=1}^k \mathbf{W}_j^{(2)} \mathbf{h}_j}{\partial \mathbf{W}^{(1)}} \\
&= \frac{\partial \sum_{j=1}^k \mathbf{W}_j^{(2)} f_1(\mathbf{W}_{j:}^{(1)} \mathbf{x})}{\partial \mathbf{W}^{(1)}} \\
&= \sum_{j=1}^k \mathbf{W}_j^{(2)} \frac{\partial f_1(\mathbf{W}_{j:}^{(1)} \mathbf{x})}{\partial \mathbf{W}^{(1)}}
\end{aligned}$$

because we always assume that the transfer on the vector is applied element-wise, and so is a vector of values where the transfer is applied to each scalar $\mathbf{W}_{j:}^{(1)} \mathbf{x}$. Now notice that the gradient of this transfer for a specific j satisfies

$$\begin{aligned}
\frac{\partial f_1(\mathbf{W}_{j:}^{(1)} \mathbf{x})}{\partial \mathbf{W}^{(1)}_k} &= 0 \quad \triangleright k \neq j \\
\frac{\partial f_1(\mathbf{W}_{j:}^{(1)} \mathbf{x})}{\partial \mathbf{W}^{(1)}_{j:}} &= \frac{\partial f_1(\boldsymbol{\theta}_j^{(1)})}{\partial \boldsymbol{\theta}_j^{(1)}} \frac{\partial \boldsymbol{\theta}_j^{(1)}}{\partial \mathbf{W}_{j:}^{(1)}} \\
&= \frac{\partial f_1(\boldsymbol{\theta}_j^{(1)})}{\partial \boldsymbol{\theta}_j^{(1)}} \mathbf{x}^\top
\end{aligned}$$

So for each j , $\frac{\partial f_1(\mathbf{W}_{j:}^{(1)} \mathbf{x})}{\partial \mathbf{W}^{(1)}}$ is a $\mathbb{R}^{d \times k}$ matrix with every row equal to zero except

for the j th row. Therefore, we can write

$$\begin{aligned}\frac{\partial \mathbf{W}^{(2)} \mathbf{h}}{\partial \mathbf{W}^{(1)}} &= \sum_{j=1}^k \mathbf{W}_j^{(2)} \frac{\partial f_1(\mathbf{W}_{j:}^{(1)} \mathbf{x})}{\partial \mathbf{W}^{(1)}} \\ &= \left(\mathbf{W}^{(2)} \circ \left[\frac{\partial f_1(\boldsymbol{\theta}_1^{(1)})}{\partial \boldsymbol{\theta}_1^{(1)}}, \dots, \frac{\partial f_1(\boldsymbol{\theta}_k^{(1)})}{\partial \boldsymbol{\theta}_k^{(1)}} \right] \right)^\top \mathbf{x}^\top\end{aligned}$$

where again \circ indicates element-wise multiplication. With an abuse of notation, we will write

$$\nabla f_1(\boldsymbol{\theta}^{(1)}) = \left[\frac{\partial f_1(\boldsymbol{\theta}_1^{(1)})}{\partial \boldsymbol{\theta}_1^{(1)}}, \dots, \frac{\partial f_1(\boldsymbol{\theta}_k^{(1)})}{\partial \boldsymbol{\theta}_k^{(1)}} \right] \in \mathbb{R}^{1 \times k}.$$

Letting $\Delta_i = \frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}}$, we get the two gradients

$$\begin{aligned}\frac{\partial \text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial \mathbf{W}^{(1)}} &= \Delta_i \left(\mathbf{W}^{(2)} \circ \nabla f_1(\boldsymbol{\theta}^{(1)}) \right)^\top \mathbf{x}_i^\top \\ \frac{\partial \text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial \mathbf{W}^{(2)}} &= \Delta_i \mathbf{h}^\top\end{aligned}$$

Example 13: Let $p(y = 1|\mathbf{x})$ be a Bernoulli distribution, with f_1 and f_2 both sigmoid functions. The loss is the cross-entropy. We can derive the two-layer update rule with these settings, by plugging-in above.

$$\begin{aligned}L(\hat{y}, y_i) &= -y_i \log(\hat{y}) - (1 - y_i) \log(1 - \hat{y}) && \triangleright \text{cross-entropy} \\ \frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} &= -\frac{y_i}{\hat{y}} + \frac{1 - y_i}{1 - \hat{y}} \\ f_1(\mathbf{W}^{(1)} \mathbf{x}) &= \sigma(\mathbf{W}^{(1)} \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{W}^{(1)} \mathbf{x})} \\ f_2(\mathbf{W}^{(2)} \mathbf{h}) &= \sigma(\mathbf{W}^{(2)} \mathbf{h}) = \frac{1}{1 + \exp(-\mathbf{W}^{(2)} \mathbf{h})} \\ \nabla \sigma(\boldsymbol{\theta}) &= \sigma(\boldsymbol{\theta}) \circ (\mathbf{1} - \sigma(\boldsymbol{\theta}))\end{aligned}$$

The explicit Hadamard product \circ is given to make it clear that we are using element-wise multiplication between the two vectors $\sigma(\boldsymbol{\theta})$ and $(\mathbf{1} - \sigma(\boldsymbol{\theta}))$, rather than a dot product.

Now we can plug-in these basic gradients to obtain

$$\begin{aligned}
\mathbf{h} &= \sigma(\mathbf{W}^{(1)}\mathbf{x}) \\
\hat{y} &= \sigma(\mathbf{W}^{(2)}\mathbf{h}) \\
\Delta_i &= \frac{\partial L(\hat{y}, y_i)}{\partial \hat{y}} \frac{\partial f_2(\boldsymbol{\theta}^{(2)})}{\partial \boldsymbol{\theta}^{(2)}} = \left(-\frac{y_i}{\hat{y}} + \frac{1-y_i}{1-\hat{y}} \right) \hat{y}(1-\hat{y}) \\
\frac{\partial \text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial \mathbf{W}^{(2)}} &= \Delta_i \mathbf{h}^\top \\
\frac{\partial \text{Err}_i(\mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial \mathbf{W}^{(1)}} &= \Delta_i \left(\mathbf{W}^{(2)} \circ \nabla f_1(\boldsymbol{\theta}^{(1)}) \right)^\top \mathbf{x}_i^\top \\
&= \Delta_i \left(\mathbf{W}^{(2)} \circ \mathbf{h} \circ (\mathbf{1} - \mathbf{h}) \right)^\top \mathbf{x}_i^\top
\end{aligned}$$

Now the update simply consists of stepping in the direction of these gradients, as is usual for gradient descent. We start with some initial $\mathbf{W}^{(1)}$ and $\mathbf{W}^{(2)}$ (say filled with random values), and then apply the gradient descent rules with these gradients. The back-propagation algorithm tells us how to compute the values in this gradient in an efficient order. \square

This algorithm has a clear extension to multiple layers. We simply proceed in the same way as with $\mathbf{W}^{(2)}$, which is to continue computing the gradient with the chain rule.

Exercise: Compute the gradient updates for a three layer neural network, for generic transfers f_1 , f_2 and f_3 and matching loss $L(\cdot, y)$.

Exercise: Modify the above gradient updates with a vector output $\mathbf{y} \in \mathbb{R}^k$, instead of scalar $y \in \mathbb{R}$.

7.2.2 Unsupervised learning and matrix factorization

Another strategy to obtaining a new representation is through matrix factorization. In matrix factorization, the data matrix \mathbf{X} is factorized into a dictionary \mathbf{D} and a basis or new representation \mathbf{H} (see Figure 7.4). In fact, many unsupervised learning algorithms (e.g., dimensionality reduction, sparse coding) and semi-supervised learning algorithms (e.g., supervised dictionary learning) can actually be formulated as matrix factorizations. We will look at principal components analysis as an example. The remaining algorithms are simply summarized in the below table.

Principal components analysis (PCA) is a standard dimensionality reduction technique, where the input data $\mathbf{x} \in \mathbb{R}^d$ is projected into a lower dimensional $\mathbf{h} \in \mathbb{R}^k$ spanned by the space of principal components. These principal components are the directions of maximal variance in the




Figure 7.4: Matrix factorization of data matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$.

data. To obtain these principal components $\mathbf{D} \in \mathbb{R}^{d \times k}$, the common solution technique is to obtain the singular value decomposition of the data matrix $\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top \in \mathbb{R}^{n \times d}$, giving $\mathbf{D} = \mathbf{V}_k^\top \in \mathbb{R}^k$ and $\mathbf{H} = \mathbf{U}_k\mathbf{\Sigma}_k$ where $\mathbf{\Sigma}_k \in \mathbb{R}^{k \times k}$ consists of the top largest k singular values (in descending order) and $\mathbf{U}_k \in \mathbb{R}^{n \times k}$ and $\mathbf{V}_k \in \mathbb{R}^{k \times d}$ are the corresponding singular vectors, i.e., $\mathbf{U}_k = \mathbf{U}_{:,1:k}$ and $\mathbf{V}_k = \mathbf{V}_{:,1:k}$. The new representation for \mathbf{X} (using PCA) is this \mathbf{H} .

Interestingly, this factorization corresponds to solving

$$\min_{\mathbf{D} \in \mathbb{R}^{k \times d}, \mathbf{H} \in \mathbb{R}^{n \times k}} \|\mathbf{X} - \mathbf{H}\mathbf{D}\|_F^2$$

This formulation makes it clear how we can obtain out-of-sample representations. For a new sample \mathbf{x} , the representation can be obtained using

$$\min_{\mathbf{h} \in \mathbb{R}^k} \|\mathbf{x} - \mathbf{h}\mathbf{D}\|_2^2.$$

Moreover, this formulation makes it simpler to incorporate supervised information, as we will see.

There are many variants of unsupervised learning algorithms that actually correspond to factorizing the data matrix, summarized in Table 7.5. In many cases, they simply correspond to defining an interesting kernel on \mathbf{X} , and then factorizing that kernel (i.e., kernel PCA). If the entry is empty, this specifies no regularization and no constraints. For a more complete list, see [10]) and [12]. Properties of the representation that are typically encoded are those that obtain lower-dimensional features, sparse features, and/or local features. Lower-dimensional features are encoded by taking $k < d$ (PCA),

Algorithm	Loss	Constraints & Regularizers
CCA ≡ orthonormal PLS	$\left\ \begin{bmatrix} \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \\ \mathbf{Y}(\mathbf{Y}^\top \mathbf{Y})^{-1} \end{bmatrix} - \mathbf{H}\mathbf{D} \right\ _F^2$	
Isomap	$\ \mathbf{K} - \mathbf{H}\mathbf{D}\ _F^2$ $\mathbf{K} = -\frac{1}{2}(I - \mathbf{e}\mathbf{e}')S(I - \mathbf{e}\mathbf{e}')$ with $S_{i,j} = \ \mathbf{X}_{i:} - \mathbf{X}_{j:}\ $	
K-means clustering	$\ \mathbf{X} - \mathbf{H}\mathbf{D}\ _F^2$	$\mathbf{H} \in \{0,1\}^{n \times k}$ $\mathbf{H}\mathbf{1} = \mathbf{1}$
K-medians clustering	$\ \mathbf{X} - \mathbf{H}\mathbf{D}\ _{1,1}$	$\mathbf{H} \in \{0,1\}^{n \times k}$ $\mathbf{H}\mathbf{1} = \mathbf{1}$
Laplacian eigenmaps ≡ Kernel LPP	$\ \mathbf{K} - \mathbf{H}\mathbf{D}\ _F^2$ for $\mathbf{K} = \mathbf{L}^\dagger$	
Metric multi-dimensional scaling	$\ \mathbf{K} - \mathbf{H}\mathbf{D}\ _F^2$ for isotropic kernel K	
Normalized-cut	$\ (\Lambda^{-1}\mathbf{X} - \mathbf{H}\mathbf{D})\Lambda^{1/2}\ _F^2$	$\mathbf{H} \in \{0,1\}^{n \times k}$ $\mathbf{H}\mathbf{1} = \mathbf{1}$
Partial least squares	$\ \mathbf{X}\mathbf{Y}' - \mathbf{D}\mathbf{H}\ _F^2$	
PCA	$\ \mathbf{X} - \mathbf{H}\mathbf{D}\ _F^2$	
Kernel PCA	$\ \mathbf{K} - \mathbf{H}\mathbf{D}\ _F^2$	
Ratio cut	$\ \mathbf{K} - \mathbf{H}\mathbf{D}\ _F^2$ for $\mathbf{K} = \mathbf{L}^\dagger$	

Figure 7.5: Table of unsupervised learning algorithms that correspond to a matrix factorization.

or to add an ℓ_2 regularizer on each column of $\mathbf{H}_{:i}$ (relaxed rank PCA):

$$\begin{aligned}
& \min_{\mathbf{D} \in \mathbb{R}^{k \times d}, \mathbf{H} \in \mathbb{R}^{n \times k}} \|\mathbf{X} - \mathbf{H}\mathbf{D}\|_F^2 + \alpha \sum_{i=1}^k \|\mathbf{H}_{:i}\|_2^2 + \alpha \sum_{i=1}^k \|\mathbf{D}_{i:}\|_2^2 \\
& = \min_{\mathbf{D} \in \mathbb{R}^{k \times d}, \mathbf{H} \in \mathbb{R}^{n \times k}} \|\mathbf{X} - \mathbf{H}\mathbf{D}\|_F^2 + \alpha \|\mathbf{H}\|_F^2 + \alpha \|\mathbf{D}\|_F^2
\end{aligned}$$

Note that a regularizer is also added to \mathbf{D} , to ensure that \mathbf{D} does not become too large. Sparse features are typically encoded using an ℓ_1 regularizer on each column of $\mathbf{H}_{:i}$ (e.g., sparse coding). The ℓ_1 regularizer promotes zeroed entries, and so prefers columns of $\mathbf{H}_{:i}$ that have many zeros. Having many sparse features typically also corresponds to a local representation, where similar inputs have only a small number of features that are non-negligibly greater than zero. As with the regression setting, we can generalize this Euclidean loss to any convex loss:

$$L_x(\mathbf{H}, \mathbf{D}, \mathbf{X}) = \sum_{i=1}^n L_x(\mathbf{H}_{i:}, \mathbf{D}, \mathbf{X}_{i:}).$$

Above, we have used

$$\begin{aligned} L_x(\mathbf{H}, \mathbf{D}, \mathbf{X}) &= \sum_{i=1}^n \|\mathbf{H}_{i:} \mathbf{D} - \mathbf{X}_{i:}\|_2^2 \\ &= \|\mathbf{H} \mathbf{D} - \mathbf{X}\|_F^2. \end{aligned}$$

For intuition, you can simply imagine this linear loss, but keep in mind that it can be generalized.

Our focus remains on prediction, and so we will next see how this factorization approach can also be used for supervised (or semi-supervised) representation learning. The many other unsupervised objectives (summarized in the table) can be used in conjunction with the more general supervised representation learning setting described next.

Supervised dictionary learning

As with neural networks, we can inform the learning of this hidden representation \mathbf{h} by incorporating a supervised loss. The goal is to use this new representation for learning; therefore, we would like to again represent $E[y|\mathbf{x}]$ with $f(\mathbf{W}\mathbf{h})$ using this new representation imputed from \mathbf{x} . For the chosen transfer exponential family distribution and transfer f , let $L_y(\cdot, y)$ be the corresponding loss (i.e., the negative log-likelihood, with constants removed). Then the supervised dictionary learning problem, which adds a supervised component to the unsupervised matrix factorization, can be quite simply written as

$$\min_{\mathbf{D} \in \mathbb{R}^{k \times d}, \mathbf{W} \in \mathbb{R}^{k \times 1}, \mathbf{H} \in \mathbb{R}^{n \times k}} L_x(\mathbf{H} \mathbf{D}, \mathbf{X}) + \sum_{i=1}^n L_y(\mathbf{H}_{i:} \mathbf{W}, y_i).$$

As before, it is common to add regularization to these parameters, including for subspace and sparsity properties. This supervised dictionary learning model is given in Figure 7.6.

Algorithms

The most common strategy to learn these supervised dictionary models is to do an alternating minimization over the variables. Because \mathbf{D} and \mathbf{H} are multiplied, commonly one optimizes the inner variable first (say \mathbf{H}), and then descends on the \mathbf{D} and \mathbf{W} , the other variables which are only related through \mathbf{H} . This alternating procedure is guaranteed to converge to a local minima. Recent evidence, however, suggests that for many cases with




Figure 7.6: Supervised dictionary learning. Notice that the graphical model looks similar to a neural network, but the arrows are directed out from the hidden layer into the input layer (which is being factorized). An important difference is that the hidden layer \mathbf{h} is now treated explicitly as a variable and is learned.

supervised dictionary learning, this procedure actually returns the global minimum (see e.g. [2]). Therefore, this procedure appears to be a viable strategy forward for these models.

One first computes the solution to \mathbf{H} for the given data \mathbf{X}, \mathbf{Y} , with \mathbf{D} and \mathbf{W} fixed. This involves computing the gradient w.r.t. to \mathbf{H} , and then performing gradient descent. Then, with \mathbf{H} fixed, one computes the solution to \mathbf{D} and \mathbf{W} , again by computing the gradient and descending to a local minimum. This corresponds to a batch gradient descent, with alternating minimization. The stochastic gradient descent for these models is more complicated, and we will not address it here.

Chapter 8

Empirical Evaluation

Empirical evaluation of models and algorithms that construct them lies in the core of machine learning. Although the field is rife with theoretical models and extremely powerful apparatus for their analysis, at the end of the day, machine learning is a field with real-life data and its goal is development of methods that would work well on such data. Thus, evaluation of methods on real-life data cannot be avoided.

Given a data set $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$, where $\mathbf{x}_i \in \mathcal{X}$ is the i -th object and $y_i \in \mathcal{Y}$ is the corresponding target designation, we will distinguish between evaluation of prediction models and the evaluation of algorithms that construct them. A prediction model is a function $f : \mathcal{X} \rightarrow \mathcal{Y}$ constructed from data set \mathcal{D} using an algorithm $a : (\mathcal{X} \times \mathcal{Y})^n \rightarrow \mathcal{F}$, where \mathcal{F} is called a function or hypothesis space and $f \in \mathcal{F}$. We can also say that $f = a(\mathcal{D})$. We can think of a being a particular learning algorithm, say logistic regression or neural networks, where f is a particular function with fixed set of parameters.

8.1 Comparison of Learning Algorithms

To compare learning algorithms, it is necessary to collect a list of data sets such that the performance can be compared on each such data set. Suppose we have a set of learning problems $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_m$ and wish to compare learning algorithms a_1 and a_2 . We can carry out such a comparison using a counting test as follows: for each data set both algorithms are evaluated in terms of relevant performance measure and the algorithm with a higher performance accuracy is awarded a win, while the other one is given a loss (in case of exactly the same performance, we can provide a win/loss randomly).

We are now interested in providing statistical evidence that say algorithm a_1 is better than algorithm a_2 . Suppose a_1 has k wins out of m and algorithm a_2 has $m - k$ wins, as shown in Table 8.1. We would like to evaluate the null hypothesis H_0 that algorithms a_1 and a_2 have the same performance by

	\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\mathcal{D}_4		\mathcal{D}_{m-1}	\mathcal{D}_m
a_1	1	0	1	1	\dots	0	1
a_2	0	1	0	0		1	0

Table 8.1: A counting test where learning algorithms a_1 and a_2 are compared on a set of m independent data sets. An algorithm with a better performance on a particular data set collects a win (1), whereas the other algorithm collects a loss (0).

providing an alternative hypothesis H_1 that algorithm a_1 is better than a_2 . In short,

H_0 : $\text{quality}(a_1) = \text{quality}(a_2)$

H_1 : $\text{quality}(a_1) > \text{quality}(a_2)$

If the null hypothesis is true, the win/loss on each data set will be equally likely and determined by minor variation. Therefore, the probability of a win on any data set will be roughly equal to $p = 1/2$. Now, we can express the probability that algorithm a_1 collected k wins or more under the null hypothesis using binomial distribution

$$P = \sum_{i=k}^m \binom{m}{i} p^i (1-p)^{m-i}$$

and refer to it as the P-value. A typical approach in these cases is to establish a significance values, say, $\alpha = 0.05$ and reject the null hypothesis if $P \leq \alpha$. If the P-value is greater than α we say that there is insufficient evidence for rejecting H_0 . For sufficiently low P-values, we may conclude that there is sufficient evidence that algorithm a_1 is better than algorithm a_2 .

The choice of the significance threshold α is somewhat arbitrary. Typically, 5% is a reasonable value, but lower values indicate that the particular situation of k wins out of m was so unlikely, that we can consider the evidence for rejecting H_0 very strong.

8.2 Performance of Classification Models

Name	Symbol	Definition
Classification error	$error$	$error = \frac{fp+fn}{tp+fp+tn+fn}$
Classification accuracy	$accuracy$	$accuracy = 1 - error$
True positive rate	tpr	$tpr = \frac{tp}{tp+fn}$
False negative rate	fnr	$fnr = \frac{fn}{tp+fn}$
True negative rate	tnr	$tnr = \frac{tn}{tn+fp}$
False positive rate	fpr	$fpr = \frac{fp}{tn+fp}$
Precision	pr	$pr = \frac{tp}{tp+fp}$
Recall	rc	$rc = \frac{tp}{tp+fn}$
F-measure	F	

Table 8.2: Some

Chapter 9

Advanced topics

9.1 Bayesian estimation

Maximum a posteriori and maximum likelihood approaches report the solution that corresponds to the mode of the posterior distribution and the likelihood function, respectively. This approach, however, does not consider the possibility of skewed distributions, multimodal distributions or simply large regions with similar values of $p(M|\mathcal{D})$. Bayesian estimation addresses those concerns.

The main idea in Bayesian statistics is minimization of the *posterior risk*

$$R = \int_{\mathcal{M}} \ell(M, \hat{M}) \cdot p(M|\mathcal{D}) dM$$

where \hat{M} is our estimate and $\ell(M, \hat{M})$ is some loss function between two models. When $\ell(M, \hat{M}) = (M - \hat{M})^2$ (ignore the abuse of notation), we can minimize the posterior risk as follows

$$\begin{aligned} \frac{\partial}{\partial \hat{M}} R &= 2\hat{M} - 2 \int_{\mathcal{M}} M \cdot p(M|\mathcal{D}) dM \\ &= 0 \end{aligned}$$

from which it can be derived that the minimizer of the posterior risk is the posterior mean function, i.e.

$$\begin{aligned} M_B &= \int_{\mathcal{M}} M \cdot p(M|\mathcal{D}) dM \\ &= E_M[M|\mathcal{D}]. \end{aligned}$$

We shall refer to M_B as the Bayes estimator. It is important to mention that computing the posterior mean usually involves solving complex integrals. In

some situations, these integrals can be solved analytically; in others, numerical integration is necessary.

Example 11. Let $\mathcal{D} = \{2, 5, 9, 5, 4, 8\}$ yet again be an i.i.d. sample from $\text{Poisson}(\lambda_t)$. Suppose the prior knowledge about the parameter of the exponential distribution can be expressed using a gamma distribution with parameters $k = 3$ and $\theta = 1$. Find the Bayesian estimate of λ_t .

We want to find $E[\lambda|\mathcal{D}]$. Let us first write the posterior distribution as

$$\begin{aligned} p(\lambda|\mathcal{D}) &= \frac{p(\mathcal{D}|\lambda)p(\lambda)}{p(\mathcal{D})} \\ &= \frac{p(\mathcal{D}|\lambda)p(\lambda)}{\int_0^\infty p(\mathcal{D}|\lambda)p(\lambda)d\lambda}, \end{aligned}$$

where, as shown in previous examples, we have that

$$p(\mathcal{D}|\lambda) = \frac{\lambda^{\sum_{i=1}^n x_i} \cdot e^{-n\lambda}}{\prod_{i=1}^n x_i!}$$

and

$$p(\lambda) = \frac{\lambda^{k-1} e^{-\frac{\lambda}{\theta}}}{\theta^k \Gamma(k)}.$$

Before calculating $p(\mathcal{D})$, let us first note that

$$\int_0^\infty x^{\alpha-1} e^{-\beta x} dx = \frac{\Gamma(\alpha)}{\beta^\alpha}.$$

Now, we can derive that

$$\begin{aligned} p(\mathcal{D}) &= \int_0^\infty p(\mathcal{D}|\lambda)p(\lambda)d\lambda \\ &= \int_0^\infty \frac{\lambda^{\sum_{i=1}^n x_i} \cdot e^{-n\lambda}}{\prod_{i=1}^n x_i!} \cdot \frac{\lambda^{k-1} e^{-\frac{\lambda}{\theta}}}{\theta^k \Gamma(k)} d\lambda \\ &= \frac{\Gamma(k + \sum_{i=1}^n x_i)}{\theta^k \Gamma(k) \prod_{i=1}^n x_i! (n + \frac{1}{\theta})^{\sum_{i=1}^n x_i + k}} \end{aligned}$$

and subsequently that

$$\begin{aligned}
p(\lambda|\mathcal{D}) &= \frac{p(\mathcal{D}|\lambda)p(\lambda)}{p(\mathcal{D})} \\
&= \frac{\lambda^{\sum_{i=1}^n x_i} \cdot e^{-n\lambda}}{\prod_{i=1}^n x_i!} \cdot \frac{\lambda^{k-1} e^{-\frac{\lambda}{\theta}}}{\theta^k \Gamma(k)} \cdot \frac{\theta^k \Gamma(k) \prod_{i=1}^n x_i! (n + \frac{1}{\theta})^{\sum_{i=1}^n x_i + k}}{\Gamma(k + \sum_{i=1}^n x_i)} \\
&= \frac{\lambda^{k-1 + \sum_{i=1}^n x_i} \cdot e^{-\lambda(n + 1/\theta)} \cdot (n + \frac{1}{\theta})^{\sum_{i=1}^n x_i + k}}{\Gamma(k + \sum_{i=1}^n x_i)}.
\end{aligned}$$

Finally,

$$\begin{aligned}
E[\lambda|\mathcal{D}] &= \int_0^\infty \lambda p(\lambda|\mathcal{D}) d\lambda \\
&= \frac{k + \sum_{i=1}^n x_i}{n + \frac{1}{\theta}} \\
&= 5.14
\end{aligned}$$

which is nearly the same solution as the MAP estimate found in Example 9.

□

It is evident from the previous example that selection of the prior distribution has important implications on calculation of the posterior mean. We have not picked the gamma distribution by chance; that is, when the likelihood was multiplied by the prior, the resulting distribution remained in the same class of functions as the likelihood. We shall refer to such prior distributions as *conjugate priors*. Conjugate priors are also simplifying the mathematics; in fact, this is a major reason for their consideration. Interestingly, in addition to the Poisson distribution, the gamma distribution is a conjugate prior to the exponential distribution as well as the gamma distribution itself.

9.2 Parameter estimation for mixtures of distributions

We now investigate parameter estimation for mixture models, which is most commonly carried out using the *expectation-maximization (EM) algorithm*. As before, we are given a set of i.i.d. observations $\mathcal{D} = \{x_i\}_{i=1}^n$, with the goal of estimating the parameters of the mixture distribution

$$p(x|\theta) = \sum_{j=1}^m w_j p(x|\theta_j).$$

In the equation above, we used $\theta = (w_1, w_2, \dots, w_m, \theta_1, \theta_2, \dots, \theta_m)$ to combine all parameters. Just to be more concrete, we shall assume that each $p(x_i|\theta_j)$ is an exponential distribution with parameter λ_j , i.e. $p(x|\theta_j) = \lambda_j e^{-\lambda_j x}$, where $\lambda_j > 0$. Finally, we shall assume that m is given and will address simultaneous estimation of θ and m later.

Let us attempt to find the maximum likelihood solution first. By plugging the formula for $p(x|\theta)$ into the likelihood function we obtain

$$\begin{aligned} p(\mathcal{D}|\theta) &= \prod_{i=1}^n p(x_i|\theta) \\ &= \prod_{i=1}^n \left(\sum_{j=1}^m w_j p(x_i|\theta_j) \right), \end{aligned} \quad (9.1)$$

which, unfortunately, is difficult to maximize using differential calculus (why?). Note that although $p(\mathcal{D}|\theta)$ has $O(m^n)$ terms, it can be calculated in $O(mn)$ time as a log-likelihood.

Before introducing the EM algorithm, let us for a moment present two hypothetical scenarios that will help us to understand the algorithm. First, suppose that information is available as to which mixing component generated which data point. That is, suppose that $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ is an i.i.d. sample from some distribution $p_{XY}(x, y)$, where $y \in \mathcal{Y} = \{1, 2, \dots, m\}$ specifies the mixing component. How would the maximization be performed then? Let us write the likelihood function as

$$\begin{aligned} p(\mathcal{D}|\theta) &= \prod_{i=1}^n p(x_i, y_i|\theta) \\ &= \prod_{i=1}^n p(x_i|y_i, \theta) p(y_i|\theta) \\ &= \prod_{i=1}^n w_{y_i} p(x_i|\theta_{y_i}), \end{aligned} \quad (9.2)$$

where $w_j = p_Y(j) = P(Y = j)$. The log-likelihood is

$$\begin{aligned}\log p(\mathcal{D}|\theta) &= \sum_{i=1}^n (\log w_{y_i} + \log p(x_i|\theta_{y_i})) \\ &= \sum_{j=1}^m n_j \log w_j + \sum_{i=1}^n \log p(x_i|\theta_{y_i}),\end{aligned}$$

where n_j is the number of data points in \mathcal{D} generated by the j -th mixing component.

It is useful to observe here that when $\mathbf{y} = (y_1, y_2, \dots, y_n)$ is known, the internal summation operator in Eq. (9.1) disappears. More importantly, it follows that Eq. (9.2) can be maximized in a relatively straightforward manner. Let us show how. To find $\mathbf{w} = (w_1, w_2, \dots, w_m)$ we need to solve a constrained optimization problem, which we will do by using the method of Lagrange multipliers. We shall first form the Lagrangian function as

$$L(\mathbf{w}, \alpha) = \sum_{j=1}^m n_j \log w_j + \alpha \left(\sum_{j=1}^m w_j - 1 \right)$$

where α is the Lagrange multiplier. Then, by setting $\frac{\partial}{\partial w_k} L(\mathbf{w}, \alpha) = 0$ for every $k \in \mathcal{Y}$ and $\frac{\partial}{\partial \alpha} L(\mathbf{w}, \alpha) = 0$, we derive that $w_k = -\frac{n_k}{\alpha}$ and $\alpha = -n$. Thus,

$$w_k = \frac{1}{n} \sum_{i=1}^n I(y_i = k),$$

where $I(\cdot)$ is the indicator function. To find all θ_j , we recall that we assumed a mixture of exponential distributions. Thus, we proceed by setting

$$\frac{\partial}{\partial \lambda_k} \sum_{i=1}^n \log p(x_i|\lambda_{y_i}) = 0,$$

for each $k \in \mathcal{Y}$. We obtain that

$$\lambda_k = \frac{n_k}{\sum_{i=1}^n I(y_i = k) \cdot x_i},$$

which is simply the inverse mean over those data points generated by the k -th mixture component. In summary, we observe that if the mixing component designations \mathbf{y} are known, the parameter estimation is greatly simplified. This was achieved by decoupling the estimation of mixing proportions and all parameters of the mixing distributions.

In the second hypothetical scenario, suppose that parameters θ are known, and that we would like to estimate the best configuration of the mixture designations \mathbf{y} (one may be tempted to call them “class labels”). This task looks like clustering in which cluster memberships need to be determined based on the known set of mixing distributions and mixing probabilities. To do this we can calculate the posterior distribution of \mathbf{y} as

$$\begin{aligned} p(\mathbf{y}|\mathcal{D}, \theta) &= \prod_{i=1}^n p(y_i|x_i, \theta) \\ &= \prod_{i=1}^n \frac{w_{y_i} p(x_i|\theta_{y_i})}{\sum_{j=1}^m w_j p(x_i|\theta_j)} \end{aligned} \quad (9.3)$$

and subsequently find the best configuration out of m^n possibilities. Obviously, because of the i.i.d. assumption each element y_i can be estimated separately and, thus, this estimation can be completed in $O(mn)$ time. The MAP estimate for y_i can be found as

$$\hat{y}_i = \arg \max_{y_i \in \mathcal{Y}} \left\{ \frac{w_{y_i} p(x_i|\theta_{y_i})}{\sum_{j=1}^m w_j p(x_i|\theta_j)} \right\}$$

for each $i \in \{1, 2, \dots, n\}$.

In reality, neither “class labels” \mathbf{y} nor the parameters θ are known. Fortunately, we have just seen that the optimization step is relatively straightforward if one of them is known. Therefore, the intuition behind the EM algorithm is to form an iterative procedure by *assuming* that either \mathbf{y} or θ is known and calculate the other. For example, we can initially pick some value for θ , say $\theta^{(0)}$, and then estimate \mathbf{y} by computing $p(\mathbf{y}|\mathcal{D}, \theta^{(0)})$ as in Eq. (9.3). We can refer to this estimate as $\mathbf{y}^{(0)}$. Using $\mathbf{y}^{(0)}$ we can now refine the estimate of θ to $\theta^{(1)}$ using Eq. (9.2). We can then iterate these two steps until convergence. In the case of mixture of exponential distributions, we arrive at the following algorithm:

1. Initialize $\lambda_k^{(0)}$ and $w_k^{(0)}$ for $\forall k \in \mathcal{Y}$
2. Calculate $y_i^{(0)} = \arg \max_{k \in \mathcal{Y}} \left\{ \frac{w_k^{(0)} p(x_i|\lambda_k^{(0)})}{\sum_{j=1}^m w_j^{(0)} p(x_i|\lambda_j^{(0)})} \right\}$ for $\forall i \in \{1, 2, \dots, n\}$
3. Set $t = 0$
4. Repeat until convergence

- (a) $w_k^{(t+1)} = \frac{1}{n} \sum_{i=1}^n I(y_i^{(t)} = j)$
- (b) $\lambda_k^{(t+1)} = \frac{\sum_{i=1}^n I(y_i^{(t)} = k)}{\sum_{i=1}^n I(y_i^{(t)} = k) \cdot x_i}$
- (c) $t = t + 1$
- (d) $y_i^{(t+1)} = \arg \max_{k \in \mathcal{Y}} \left\{ \frac{w_k^{(t)} p(x_i | \lambda_k^{(t)})}{\sum_{j=1}^m w_j^{(t)} p(x_i | \lambda_j^{(t)})} \right\}$

5. Report $\lambda_k^{(t)}$ and $w_k^{(t)}$ for $\forall k \in \mathcal{Y}$

This procedure is not quite yet the EM algorithm; rather, it is a version of it referred to as *classification EM algorithm*. In the next section we will introduce the EM algorithm.

The expectation-maximization algorithm

The previous procedure was designed to iteratively estimate class memberships and parameters of the distribution. In reality, it is not necessary to compute \mathbf{y} ; after all, we only need to estimate θ . To accomplish this, at each step t , we can use $p(\mathbf{y}|\mathcal{D}, \theta^{(t)})$ to maximize the *expected log-likelihood* of both \mathcal{D} and \mathbf{y}

$$E_{\mathbf{Y}}[\log p(\mathcal{D}, \mathbf{y}|\theta)|\theta^{(t)}] = \sum_{\mathbf{y}} \log p(\mathcal{D}, \mathbf{y}|\theta) p(\mathbf{y}|\mathcal{D}, \theta^{(t)}), \quad (9.4)$$

which can be carried out by integrating the log-likelihood function of \mathcal{D} and \mathbf{y} over the posterior distribution for \mathbf{y} in which the current values of the parameters $\theta^{(t)}$ are assumed to be known. We can now formulate the expression for the parameters in step $t + 1$ as

$$\theta^{(t+1)} = \arg \max_{\theta} \left\{ E[\log p(\mathcal{D}, \mathbf{y}|\theta)|\theta^{(t)}] \right\}. \quad (9.5)$$

The formula above is all that is necessary to create the update rule for the EM algorithm. Note, however, that inside of it we always have to re-compute $E_{\mathbf{Y}}[\log p(\mathcal{D}, \mathbf{y}|\theta)|\theta^{(t)}]$ function because the parameters $\theta^{(t)}$ have been updated from the previous step. We then can perform maximization. Hence the name “expectation-maximization”, although it is perfectly valid to think of the EM algorithm as an iterative maximization of expectation from Eq. (9.4), i.e. “expectation maximization”.

We now proceed as follows

$$\begin{aligned}
E[\log p(\mathcal{D}, \mathbf{y}|\theta)|\theta^{(t)}] &= \sum_{y_1=1}^m \cdots \sum_{y_n=1}^m \log p(\mathcal{D}, \mathbf{y}|\theta) p(\mathbf{y}|\mathcal{D}, \theta^{(t)}) \\
&= \sum_{y_1=1}^m \cdots \sum_{y_n=1}^m \sum_{i=1}^n \log p(x_i, y_i|\theta) \prod_{l=1}^n p(y_l|x_l, \theta^{(t)}) \\
&= \sum_{y_1=1}^m \cdots \sum_{y_n=1}^m \sum_{i=1}^n \log (w_{y_i} p(x_i|\theta_{y_i})) \prod_{l=1}^n p(y_l|x_l, \theta^{(t)}).
\end{aligned}$$

After several simplification steps, that we omit for space reasons, the expectation of the likelihood can be written as

$$E[\log p(\mathcal{D}, \mathbf{y}|\theta)|\theta^{(t)}] = \sum_{i=1}^n \sum_{j=1}^m \log (w_j p(x_i|\theta_j)) p_{Y_i}(j|x_i, \theta^{(t)}),$$

from which we can see that \mathbf{w} and $\{\theta_j\}_{j=1}^m$ can be separately found. In the final two steps, we will first derive the update rule for the mixing probabilities and then by assuming the mixing distributions are exponential, derive the update rules for their parameters.

To maximize $E[\log p(\mathcal{D}, \mathbf{y}|\theta)|\theta^{(t)}]$ with respect to \mathbf{w} , we observe that this is an instance of constrained optimization because it must hold that $\sum_{i=1}^m w_i = 1$. We will use the method of Lagrange multipliers; thus, for each $k \in \mathcal{Y}$ we need to solve

$$\frac{\partial}{\partial w_k} \left(\sum_{j=1}^m \log w_j \sum_{i=1}^n p_{Y_i}(j|x_i, \theta^{(t)}) + \alpha \left(\sum_{j=1}^m w_j - 1 \right) \right) = 0,$$

where α is the Lagrange multiplier. It is relatively straightforward to show that

$$w_k^{(t+1)} = \frac{1}{n} \sum_{i=1}^n p_{Y_i}(k|x_i, \theta^{(t)}). \quad (9.6)$$

Similarly, to find the solution for the parameters of the mixture distributions, we obtain that

$$\lambda_k^{(t+1)} = \frac{\sum_{i=1}^n p_{Y_i}(k|x_i, \theta^{(t)})}{\sum_{i=1}^n x_i p_{Y_i}(k|x_i, \theta^{(t)})} \quad (9.7)$$

for $k \in \mathcal{Y}$. As previously shown, we have

$$p_{Y_i}(k|x_i, \theta^{(t)}) = \frac{w_k^{(t)} p(x_i|\lambda_k^{(t)})}{\sum_{j=1}^m w_j^{(t)} p(x_i|\lambda_j^{(t)})}, \quad (9.8)$$

which can be computed and stored as an $n \times m$ matrix. In summary, for the mixture of m exponential distributions, we summarize the EM algorithm by combining Eqs. (9.6-9.8) as follows:

1. Initialize $\lambda_k^{(0)}$ and $w_k^{(0)}$ for $\forall k \in \mathcal{Y}$
2. Set $t = 0$
3. Repeat until convergence

- (a) $p_{Y_i}(k|x_i, \theta^{(t)}) = \frac{w_k^{(t)} p(x_i|\lambda_k^{(t)})}{\sum_{j=1}^m w_j^{(t)} p(x_i|\lambda_j^{(t)})}$ for $\forall(i, k)$
- (b) $w_k^{(t+1)} = \frac{1}{n} \sum_{i=1}^n p_{Y_i}(k|x_i, \theta^{(t)})$
- (c) $\lambda_k^{(t+1)} = \frac{\sum_{i=1}^n p_{Y_i}(k|x_i, \theta^{(t)})}{\sum_{i=1}^n x_i p_{Y_i}(k|x_i, \theta^{(t)})}$
- (d) $t = t + 1$

4. Report $\lambda_k^{(t)}$ and $w_k^{(t)}$ for $\forall k \in \mathcal{Y}$

Similar update rules can be obtained for different probability distributions; however, a separate derivatives have to be found.

Notice the difference between the CEM and the EM algorithms.

Identifiability

When estimating the parameters of a mixture, it is possible that for some parametric families one obtains multiple solutions. In other words,

$$p(x; \theta) = \sum_{j=1}^m w_j p(x; \theta_j),$$

but can also be expressed as

$$p(x; \theta') = \sum_{j=1}^{m'} w'_j p(x; \theta'_j),$$

The parameters are identifiable if

$$\sum_{j=1}^m w_j p(x; \theta_j) = \sum_{j=1}^{m'} w'_j p(x; \theta'_j),$$

implies that $m = m'$ for each $j \in \{1, 2, \dots, m\}$ there exists some $l \in \{1, 2, \dots, m\}$ such that $w_j = w'_l$ and $\theta_j = \theta'_l$.

9.3 Boosting

Algorithm 3: AdaBoost algorithm. Typically, $T = 100$.

Input:

$\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, where $\mathbf{x}_i \in \mathcal{X}$ and $y_i \in \mathcal{Y}$.
Weak learning algorithm a that maps \mathcal{X} to \mathcal{Y}
Positive integer T

Initialization:

Initialize sampling distribution $p^{(1)}(i) = \frac{1}{n}$ for $\forall i \in \{1, 2, \dots, n\}$

Loop:

for $t = 1$ to T

Sample data set $\mathcal{D}^{(t)}$ from \mathcal{D} according to $p^{(t)}(i)$

Learn model $f_t(\mathbf{x})$ from $\mathcal{D}^{(t)}$

Calculate error ϵ_t on training data \mathcal{D} as

$$\epsilon_t = \sum_{i: f_t(\mathbf{x}_i) \neq y_i} p^{(t)}(i)$$

Set $\beta_t = \frac{\epsilon_t}{1 - \epsilon_t}$

Set $w_t = \ln \frac{1}{\beta_t}$

Set $p^{(t+1)}(i) = \frac{p^{(t)}(i)}{Z} \cdot \begin{cases} \beta_t & \text{if } f_t(\mathbf{x}_i) = y_i \\ 1 & \text{otherwise} \end{cases}$, where Z is a

normalizer

end

Output:

$$f(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} \left(\sum_{t=1}^T w_t \cdot I(f_t(\mathbf{x}) = y) \right)$$

Bibliography

- [1] P Auer, M Herbster, and M K Warmuth. Exponentially many local minima for single neurons. In *Advances in Neural Information Processing Systems*, 1996.
- [2] Francis Bach, Julien Mairal, and Jean Ponce. Convex sparse matrix factorizations. *arXiv:0812.1869v1 [cs.LG]*, December 2008.
- [3] L on Bottou and Yann Le Cun. On-line learning for very large data sets. *Applied Stochastic Models in Business and Industry*, 21(2):137–151, 2005.
- [4] Léon Bottou. Online Learning and Stochastic Approximations. *Online Learning and Neural Networks*, 1998.
- [5] S Chrétien and A O Hero. On EM algorithms and their proximal generalizations. *ESAIM: Probability and Statistics*, 2008.
- [6] M L Overton and R S Womersley. Optimality conditions and duality theory for minimizing sums of the largest eigenvalues of symmetric matrices. *Mathematical Programming*, 62(1-3):321–357, February 1993.
- [7] R Salakhutdinov, S Roweis, and Z Ghahramani. Optimization with EM and expectation-conjugate-gradient. *ICML*, 2003.
- [8] R Salakhutdinov, S Roweis, and Z Ghahramani. Relationship between gradient and EM steps in latent variable models. Technical report, University of Toronto, 2004.
- [9] Dinah Shender and John Lafferty. Computation-Risk Tradeoffs for Covariance-Thresholded Regression. *International Conference on Machine Learning*, pages 756–764, 2013.
- [10] Ajit P Singh and Geoffrey J Gordon. A unified view of matrix factorization models. In *Machine Learning and Knowledge Discovery in Databases*, pages 358–373. Springer Berlin Heidelberg, Berlin, Heidelberg, 2008.
- [11] Larry Wasserman. *All of Statistics: A Concise Course in Statistical Inference*. Springer, January 2004.
- [12] M White. *Regularized factor models*. PhD thesis, University of Alberta, 2014.
- [13] Lei Xu and Michael I Jordan. On Convergence Properties of the EM Algorithm for Gaussian Mixtures. *Neural Computation*, 1996.

Appendix A

Optimization background

A.1 Second order optimization: Newton-Raphson method

A function $f(x)$ in the neighborhood of point x_0 , can be approximated using the Taylor series as

$$f(x) = \sum_{n=0}^{\infty} \frac{f^{(n)}(x_0)}{n!} (x - x_0)^n,$$

where $f^{(n)}(x_0)$ is the n -th derivative of function $f(x)$ evaluated at point x_0 . Also, $f(x)$ is considered to be infinitely differentiable. For practical reasons, we will approximate this function using the first three terms of the series as

$$f(x) \approx f(x_0) + (x - x_0)f'(x_0) + \frac{1}{2}(x - x_0)^2 f''(x_0).$$

The optimum of this function can be found by finding the first derivative and setting it to zero (technically, one should check the second derivative as well)

$$f'(x) \approx f'(x_0) + (x - x_0)f''(x_0) = 0.$$

Solving this equation for x gives us

$$x = x_0 - \frac{f'(x_0)}{f''(x_0)}.$$

Note that the approach assumes that a good enough solution x_0 already exists. However, this equation, also provides a basis for an iterative process in finding the optimum of function $f(x)$. For example, if $x^{(i)}$ is the value of x in the i -th step, then the value in step $i + 1$ can be obtained as

$$x^{(i+1)} = x^{(i)} - \frac{f'(x^{(i)})}{f''(x^{(i)})}. \quad (\text{A.1})$$

This method is called the Newton-Raphson method of optimization. We can generalize this approach to functions of vector variables $\mathbf{x} = (x_1, x_2, \dots, x_k)$. The Taylor approximation for a vector function can be written as

$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \cdot (\mathbf{x} - \mathbf{x}_0) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_0)^T \cdot H_{f(\mathbf{x}_0)} \cdot (\mathbf{x} - \mathbf{x}_0),$$

where

$$\nabla f(\mathbf{x}) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_k} \right)$$

is the gradient of function f and

$$H_{f(\mathbf{x})} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_k} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & & \\ \vdots & & \ddots & \\ \frac{\partial^2 f}{\partial x_k \partial x_1} & & & \frac{\partial^2 f}{\partial x_k^2} \end{bmatrix}$$

is the Hessian matrix of function f . Here, the gradient of f and its Hessian are evaluated at point \mathbf{x}_0 . Consequently, Eq. A.1 is modified into the following form

$$\mathbf{x}^{(i+1)} = \mathbf{x}^{(i)} - (H_{f(\mathbf{x}^{(i)})})^{-1} \cdot \nabla f(\mathbf{x}^{(i)}), \quad (\text{A.2})$$

In Eq. A.2, both gradient and Hessian are evaluated at point $\mathbf{x}^{(i)}$.